

---

# AWS Serverless Application Model Developer Guide



## **AWS Serverless Application Model: Developer Guide**

Copyright © 2020 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

## Table of Contents

What Is AWS SAM?	1
Benefits of Using AWS SAM	1
Next Step	2
Getting Started	3
Installing the AWS SAM CLI	3
Linux	3
Windows	6
macOS	8
Setting Up AWS Credentials	10
Using the AWS CLI	10
Not Using the AWS CLI	11
Tutorial: Deploying a Hello World Application	11
Prerequisites	12
Step 1: Download a Sample AWS SAM Application	12
Step 2: Build Your Application	13
Step 3: Deploy Your Application to the AWS Cloud	14
Step 4: Testing Your Application Locally (Optional)	16
Troubleshooting	18
Clean Up	20
Conclusion	20
Next Steps	21
AWS SAM Specification	22
Template Anatomy	22
YAML	23
Template Sections	23
Globals	24
Resource and Property Reference	28
AWS::Serverless::Api	28
AWS::Serverless::Application	53
AWS::Serverless::Function	56
AWS::Serverless::HttpApi (Beta)	102
AWS::Serverless::LayerVersion	111
AWS::Serverless::SimpleTable	114
Resource Attributes	117
Intrinsic Functions	118
API Gateway Extensions	118
Authoring	120
Validating AWS SAM Template Files	120
Building Applications with Dependencies	120
Working with Layers	121
Using Nested Applications	122
Defining a Nested Application from the AWS Serverless Application Repository	123
Defining a Nested Application from the Local File System	124
Deploying Nested Applications	124
Controlling Access to APIs	125
Choosing a Mechanism to Control Access	125
Example: Defining Lambda Token Authorizers	126
Example: Defining Lambda Request Authorizers	126
Example: Defining Amazon Cognito User Pools	127
Example: Defining IAM Permissions	128
Example: Defining API Keys	129
Example: Defining Resource Policies	129
Example: Defining Customized Responses	130
Testing and Debugging	132

Invoking Functions Locally .....	132
Environment Variable File .....	133
Layers .....	134
Running API Gateway Locally .....	134
Layers .....	136
Running Automated Tests .....	136
Generating Sample Event Payloads .....	138
Step-Through Debugging Lambda Functions Locally .....	138
Using AWS Toolkits .....	138
Running AWS SAM Locally .....	139
Node.js .....	139
Python .....	141
Golang .....	143
Passing Additional Runtime Debug Arguments .....	144
Deploying .....	145
Packaging and Deploying Using the AWS SAM CLI .....	145
Publishing Serverless Applications .....	145
Automating Deployments .....	145
Deploying Gradually .....	146
Monitoring .....	149
Working with Logs .....	149
Fetching Logs by AWS CloudFormation Stack .....	149
Fetching Logs by Lambda Function Name .....	149
Tailing Logs .....	149
Viewing Logs for a Specific Time Range .....	149
Filtering Logs .....	149
Error Highlighting .....	150
JSON Pretty Printing .....	150
Publishing .....	151
Prerequisites .....	151
Publishing a New Application .....	152
Step 1: Add a Metadata Section to the AWS SAM Template .....	152
Step 2: Package the Application .....	152
Step 3: Publish the Application .....	153
Publishing a New Version of an Existing Application .....	153
Additional Topics .....	154
Metadata Section Properties .....	154
Properties .....	154
Use Cases .....	155
Example .....	156
Example Applications .....	157
Process DynamoDB Events .....	157
Before You Begin .....	157
Step 1: Initialize the Application .....	157
Step 2: Test the Application Locally .....	157
Step 3: Package the Application .....	158
Step 4: Deploy the Application .....	158
Process Amazon S3 Events .....	159
Before You Begin .....	159
Step 1: Initialize the Application .....	159
Step 2: Package the Application .....	159
Step 3: Deploy the Application .....	160
Step 4: Test the Application Locally .....	160
AWS SAM Reference .....	162
AWS SAM Specification .....	162
AWS SAM CLI Command Reference .....	162
AWS SAM Policy Templates .....	162

AWS SAM CLI Command Reference .....	163
sam build .....	163
sam deploy .....	164
sam init .....	166
sam local generate-event .....	168
sam local invoke .....	169
sam local start-api .....	170
sam local start-lambda .....	172
sam logs .....	173
sam package .....	174
sam publish .....	175
sam validate .....	176
AWS SAM CLI Config .....	176
Config Basics .....	177
Example .....	177
Options .....	177
AWS SAM Policy Templates .....	178
Examples .....	178
Policy Template Table .....	179
Policy Template List .....	182
AWS SAM CLI Telemetry .....	210
Disabling Telemetry for a Session .....	211
Disabling Telemetry for Your Profile in All Sessions .....	211
Types of Information Collected .....	211
Learn More .....	212
Document History .....	213

# What Is the AWS Serverless Application Model (AWS SAM)?

The AWS Serverless Application Model (AWS SAM) is an open-source framework that you can use to build [serverless applications](#) on AWS.

A **serverless application** is a combination of Lambda functions, event sources, and other resources that work together to perform tasks. Note that a serverless application is more than just a Lambda function—it can include additional resources such as APIs, databases, and event source mappings.

You can use AWS SAM to define your serverless applications. AWS SAM consists of the following components:

- **AWS SAM template specification.** You use this specification to define your serverless application. It provides you with a simple and clean syntax to describe the functions, APIs, permissions, configurations, and events that make up a serverless application. You use an AWS SAM template file to operate on a single, deployable, versioned entity that's your serverless application. For the full AWS SAM template specification, see [AWS Serverless Application Model \(AWS SAM\) Specification \(p. 22\)](#).
- **AWS SAM command line interface (AWS SAM CLI).** You use this tool to build serverless applications that are defined by AWS SAM templates. The CLI provides commands that enable you to verify that AWS SAM template files are written according to the specification, invoke Lambda functions locally, step-through debug Lambda functions, package and deploy serverless applications to the AWS Cloud, and so on. For details about how to use the AWS SAM CLI, including the full AWS SAM CLI Command Reference, see [AWS SAM CLI Command Reference \(p. 162\)](#).

This guide shows you how to use AWS SAM to define, test, and deploy a simple serverless application. It also provides an [example application \(p. 11\)](#) that you can download, test locally, and deploy to the AWS Cloud. You can use this example application as a starting point for developing your own serverless applications.

## Benefits of Using AWS SAM

Because AWS SAM integrates with other AWS services, creating serverless applications with AWS SAM provides the following benefits:

- **Single-deployment configuration.** AWS SAM makes it easy to organize related components and resources, and operate on a single stack. You can use AWS SAM to share configuration (such as memory and timeouts) between resources, and deploy all related resources together as a single, versioned entity.
- **Extension of AWS CloudFormation.** Because AWS SAM is an extension of AWS CloudFormation, you get the reliable deployment capabilities of AWS CloudFormation. You can define resources by using AWS CloudFormation in your AWS SAM template. Also, you can use the full suite of resources, intrinsic functions, and other template features that are available in AWS CloudFormation.

- **Built-in best practices.** You can use AWS SAM to define and deploy your infrastructure as config. This makes it possible for you to use and enforce best practices such as code reviews. Also, with a few lines of configuration, you can enable safe deployments through CodeDeploy, and can enable tracing by using AWS X-Ray.
- **Local debugging and testing.** The AWS SAM CLI lets you locally build, test, and debug serverless applications that are defined by AWS SAM templates. The CLI provides a Lambda-like execution environment locally. It helps you catch issues upfront by providing parity with the actual Lambda execution environment. To step through and debug your code to understand what the code is doing, you can use AWS SAM with AWS toolkits like the [AWS Toolkit for JetBrains](#), [AWS Toolkit for PyCharm](#), [AWS Toolkit for IntelliJ](#), and [AWS Toolkit for Visual Studio Code](#). This tightens the feedback loop by making it possible for you to find and troubleshoot issues that you might run into in the cloud.
- **Deep integration with development tools.** You can use AWS SAM with a suite of AWS tools for building serverless applications. You can discover new applications in the [AWS Serverless Application Repository](#). For authoring, testing, and debugging AWS SAM-based serverless applications, you can use the [AWS Cloud9 IDE](#). To build a deployment pipeline for your serverless applications, you can use [CodeBuild](#), [CodeDeploy](#), and [CodePipeline](#). You can also use [AWS CodeStar](#) to get started with a project structure, code repository, and a CI/CD pipeline that's automatically configured for you. To deploy your serverless application, you can use the [Jenkins plugin](#). You can use the [Stackery.io toolkit](#) to build production-ready applications.

## Next Step

[Getting Started with AWS SAM \(p. 3\)](#)

# Getting Started with AWS SAM

To get started with AWS SAM, use the AWS SAM CLI to create a serverless application that you can package and deploy in the AWS Cloud. You can run the application both in the AWS Cloud or locally on your development host.

To install the AWS SAM CLI, including everything that needs to be installed or configured to use the AWS SAM CLI, see [Installing the AWS SAM CLI \(p. 3\)](#). After the AWS SAM CLI is installed, you can run through the following tutorial.

## Topics

- [Installing the AWS SAM CLI \(p. 3\)](#)
- [Setting Up AWS Credentials \(p. 10\)](#)
- [Tutorial: Deploying a Hello World Application \(p. 11\)](#)

## Installing the AWS SAM CLI

AWS SAM provides you with a command line tool, the AWS SAM CLI, that makes it easy for you to create and manage serverless applications. You need to install and configure a few things in order to use the AWS SAM CLI.

To install the AWS SAM CLI, see the following instructions for your development host:

## Topics

- [Installing the AWS SAM CLI on Linux \(p. 3\)](#)
- [Installing the AWS SAM CLI on Windows \(p. 6\)](#)
- [Installing the AWS SAM CLI on macOS \(p. 8\)](#)

## Installing the AWS SAM CLI on Linux

The following steps help you to install and configure the required prerequisites for using the AWS SAM CLI on your Linux host:

1. Create an AWS account.
2. Configure IAM permissions.
3. Install Docker. Note: Docker is only a prerequisite for testing your application locally.
4. Install Homebrew.
5. Install the AWS SAM CLI.

### Step 1: Create an AWS Account

If you don't already have an AWS account, see [aws.amazon.com](https://aws.amazon.com) and choose **Create an AWS Account**. For detailed instructions, see [Create and Activate an AWS Account](#).



## Step 2: Create an IAM User with Administrator Permissions

If you don't already have an IAM user with administrator permissions, see [Creating Your First IAM Admin User and Group](#) in the *IAM User Guide*.

In addition, you must set up AWS credentials to enable the AWS SAM CLI to make AWS service calls. For example, the AWS SAM CLI makes calls to Amazon S3 and AWS CloudFormation. For more information about setting up AWS credentials, see [Setting Up AWS Credentials \(p. 10\)](#).

## Step 3: Install Docker

### Note

Docker is only a prerequisite for testing your application locally and to build deployment packages using the `--use-container` flag. You may skip this section or install Docker at a later time if you do not plan to use these features initially.

Docker is an application that runs containers on your Linux machines. AWS SAM provides a local environment that's similar to AWS Lambda to use as a Docker container. You can use this container to build, test, and debug your serverless applications.

You must have Docker installed and working to be able to run serverless projects and functions locally with the AWS SAM CLI. The AWS SAM CLI uses the `DOCKER_HOST` environment variable to contact the Docker daemon. The following steps describe how to install, configure, and verify a Docker installation to work with the AWS SAM CLI.

Docker is available on many different operating systems, including most modern Linux distributions, like CentOS, Debian, Ubuntu, etc. For more information about how to install Docker on your particular operating system, go to the [Docker installation guide](#).

If you are using Amazon Linux 2, follow these steps to install Docker:

1. Update the installed packages and package cache on your instance.

```
sudo yum update -y
```

2. Install the most recent Docker Community Edition package.

```
sudo amazon-linux-extras install docker
```

3. Start the Docker service.

```
sudo service docker start
```

4. Add the `ec2-user` to the `docker` group so you can execute Docker commands without using `sudo`.

```
sudo usermod -a -G docker ec2-user
```

5. Log out and log back in again to pick up the new `docker` group permissions. You can accomplish this by closing your current SSH terminal window and reconnecting to your instance in a new one. Your new SSH session will have the appropriate `docker` group permissions.
6. Verify that the `ec2-user` can run Docker commands without `sudo`.

```
docker ps
```

You should see the following output, showing Docker is installed and running:

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
PORTS	NAMES			

#### Note

In some cases, you may need to reboot your instance to provide permissions for the `ec2-user` to access the Docker daemon. Try rebooting your instance if you see the following error:

```
Cannot connect to the Docker daemon. Is the docker daemon running on this host?
```

If you run into issues installing Docker, see the [Troubleshooting \(p. 6\)](#) section later in this guide, or the [Troubleshooting](#) section of the *Docker installation guide* for additional troubleshooting tips.

## Step 4: Install Homebrew

The recommended approach for installing the AWS SAM CLI on Linux is to use the Homebrew package manager. For more information about Homebrew, see [Homebrew Documentation](#).

To install Homebrew, run the following:

```
sh -c "$(curl -fsSL https://raw.githubusercontent.com/Linuxbrew/install/master/install.sh)"
```

Next, add Homebrew to your PATH by running the following commands. These commands work on all major flavors of Linux by adding either `~/.profile` on Debian/Ubuntu or `~/.bash_profile` on CentOS/Fedora/RedHat:

```
test -d ~/.linuxbrew && eval "$( ~/.linuxbrew/bin/brew shellenv )  
test -d /home/linuxbrew/.linuxbrew && eval "$( /home/linuxbrew/.linuxbrew/bin/brew shellenv )  
test -r ~/.bash_profile && echo "eval \"\${$(brew --prefix)/bin/brew shellenv}\""  
>> ~/.bash_profile  
echo "eval \"\${$(brew --prefix)/bin/brew shellenv}\"" >> ~/.profile
```

Verify that Homebrew is installed:

```
brew --version
```

You should see output like the following on successful installation of Homebrew:

```
Homebrew 2.1.6  
Homebrew/homebrew-core (git revision ef21; last commit 2019-06-19)
```

## Step 5: Install the AWS SAM CLI

Follow these steps to install the AWS SAM CLI using Homebrew:

```
brew tap aws/tap  
brew install aws-sam-cli
```

Verify the installation:

```
sam --version
```

You should see output like the following after successful installation of the AWS SAM CLI:

```
SAM CLI, version 0.33.0
```

You're now ready to start development.

## Troubleshooting

### Docker Error: "Cannot connect to the Docker daemon. Is the docker daemon running on this host?"

In some cases, you may need to reboot your instance to provide permissionst for the `ec2-user` to access the Docker daemon. If you receive this error, try rebooting your instance.

### Shell error: "command not found"

Your shell is not able to locate the AWS SAM CLI executable in the path. If you receive this error, verify the location of directory where the AWS SAM CLI executable was installed, and verify that directory is on your path.

For example, if you used the instructions in this topic to both 1) Install Homebrew, and 2) Use Homebrew to install the AWS SAM CLI, then the AWS SAM CLI executable will be installed to the following location:

```
/home/homebrew/.homebrew/bin/sam
```

## Next Steps

You're now ready to begin building your own serverless applications using AWS SAM! If you want to start with sample serverless applications, choose one of the following links:

- [Tutorial: Deploying a Hello World Application \(p. 11\)](#) – Step-by-step instructions to download, build, and deploy a simple serverless application.
- [AWS SAM example applications in GitHub](#) – Sample applications in the AWS SAM GitHub repository that you can further experiment with.

## Installing the AWS SAM CLI on Windows

The following steps help you to install and configure the required prerequisites for using the AWS SAM CLI on your Windows host:

1. Create an AWS account.
2. Configure IAM permissions.
3. Install Docker. Note: Docker is only a prerequisite for testing your application locally.
4. Install the AWS SAM CLI.

### Step 1: Create an AWS Account

If you don't already have an AWS account, see [aws.amazon.com](https://aws.amazon.com) and choose **Create an AWS Account**. For detailed instructions, see [Create and Activate an AWS Account](#).

## Step 2: Create an IAM User with Administrator Permissions

If you don't already have an IAM user with administrator permissions, see [Creating Your First IAM Admin User and Group](#) in the *IAM User Guide*.

In addition, you must set up AWS credentials to enable the AWS SAM CLI to make AWS service calls. For example, the AWS SAM CLI makes calls to Amazon S3 and AWS CloudFormation. For more information about setting up AWS credentials, see [Setting Up AWS Credentials \(p. 10\)](#).

## Step 3: Install Docker

### Note

Docker is only a prerequisite for testing your application locally and building deployment packages using the `--use-container` flag. You can skip this section or install Docker at a later time if you don't plan to use these features initially.

Docker is an application that runs containers on your Linux machines. AWS SAM provides a local environment that's similar to AWS Lambda to use as a Docker container. You can use this container to build, test, and debug your serverless applications.

You must have Docker installed and working to be able to run serverless projects and functions locally with the AWS SAM CLI. The AWS SAM CLI uses the `DOCKER_HOST` environment variable to contact the Docker daemon. The following steps describe how to install, configure, and verify a Docker installation to work with the AWS SAM CLI.

### 1. Install Docker.

Docker Desktop supports the most recent Windows operating system. For legacy versions of Windows, the Docker Toolbox is available. Choose your version of Windows for the correct Docker installation steps:

- To install Docker for Windows 10, see [Install Docker Desktop for Windows](#).
- To install Docker for older versions of Windows, see [Install Docker Toolbox on Windows](#).

### 2. Configure your shared drives.

The AWS SAM CLI requires that the project directory, or any parent directory, is listed in a shared drive. Choose your version of Windows below for the correct shared drive instructions:

- To share drives on Windows 10, see [Docker Shared Drives](#).
- To share drives on older versions of Windows, see [Add Shared Directories](#).

### 3. Verify the installation.

After Docker is installed, verify that it's working. Also confirm that you can run Docker commands from the AWS SAM CLI (for example, `docker ps`). You don't need to install, fetch, or pull any containers—the AWS SAM CLI does this automatically as required.

If you run into issues installing Docker, see the [Logs and troubleshooting](#) section of the *Docker installation guide* for additional troubleshooting tips.

## Step 4: Install the AWS SAM CLI

Windows Installer (MSI) files are the package installer files for the Windows operating system.

Follow these steps to install the AWS SAM CLI using the MSI file.

### 1. Install the AWS SAM CLI [64-bit](#).

### Note

If you operate on 32-bit machine, execute the following command: `pip install aws-sam-cli`

2. Verify the installation.

After completing the installation, verify it by opening a new command prompt or PowerShell prompt. You should be able to invoke `sam` from the command line.

```
sam --version
```

You should see output like the following after successful installation of the AWS SAM CLI:

```
SAM CLI, version 0.33.0
```

You're now ready to start development.

## Next Steps

You're now ready to begin building your own serverless applications using AWS SAM! If you want to start with sample serverless applications, choose one of the following links:

- [Tutorial: Deploying a Hello World Application \(p. 11\)](#) – Step-by-step instructions to download, build, and deploy a simple serverless application.
- [AWS SAM example applications in GitHub](#) – Sample applications in the AWS SAM GitHub repository that you can further experiment with.

## Installing the AWS SAM CLI on macOS

The following steps help you to install and configure the required prerequisites for using the AWS SAM CLI on your macOS host:

1. Create an AWS account.
2. Configure IAM permissions.
3. Install Docker. Note: Docker is only a prerequisite for testing your application locally.
4. Install Homebrew.
5. Install the AWS SAM CLI.

### Step 1: Create an AWS Account

If you don't already have an AWS account, see [aws.amazon.com](https://aws.amazon.com) and choose **Create an AWS Account**. For detailed instructions, see [Create and Activate an AWS Account](#).

### Step 2: Create an IAM User with Administrator Permissions

If you don't already have an IAM user with administrator permissions, see [Creating Your First IAM Admin User and Group](#) in the *IAM User Guide*.

In addition, you must set up AWS credentials to enable the AWS SAM CLI to make AWS service calls. For example, the AWS SAM CLI makes calls to Amazon S3 and AWS CloudFormation. For more information about setting up AWS credentials, see [Setting Up AWS Credentials \(p. 10\)](#).

## Step 3: Install Docker

### Note

Docker is only a prerequisite for testing your application locally and to build deployment packages using the `--use-container` flag. You may skip this section or install Docker at a later time if you do not plan to use these features initially.

Docker is an application that runs containers on your macOS machines. AWS SAM provides a local environment that's similar to AWS Lambda to use as a Docker container. You can use this container to build, test, and debug your serverless applications.

You must have Docker installed and working to be able to run serverless projects and functions locally with the AWS SAM CLI. The AWS SAM CLI uses the `DOCKER_HOST` environment variable to contact the Docker daemon. The following steps describe how to install, configure, and verify a Docker installation to work with the AWS SAM CLI.

#### 1. Install Docker

The AWS SAM CLI supports Docker running on macOS Sierra 10.12 or above. To install Docker see [Install Docker Desktop for Mac](#).

#### 2. Configure your shared drives

The AWS SAM CLI requires that the project directory, or any parent directory, is listed in a shared drive. To share drives on macOS, see [File sharing](#).

#### 3. Verify the installation

After Docker is installed, verify that it's working. Also confirm that you can run Docker commands from the AWS SAM CLI (for example, `docker ps`). You don't need to install, fetch, or pull any containers—the AWS SAM CLI does this automatically as required.

If you run into issues installing Docker, see the [Logs and troubleshooting](#) section of the *Docker installation guide* for additional troubleshooting tips.

## Step 4: Install Homebrew

The recommended approach for installing the AWS SAM CLI on macOS is to use the Homebrew package manager. For more information about Homebrew, see [Homebrew Documentation](#).

To install Homebrew, run the following and follow the prompts:

```
/usr/bin/ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
```

Verify that Homebrew is installed:

```
brew --version
```

You should see output like the following on successful installation of Homebrew:

```
Homebrew 2.1.6  
Homebrew/homebrew-core (git revision ef21; last commit 2019-06-19)
```

## Step 5: Install the AWS SAM CLI

Follow these steps to install the AWS SAM CLI using Homebrew:

```
brew tap aws/tap
brew install aws-sam-cli
```

Verify the installation:

```
sam --version
```

You should see output like the following after successful installation of the AWS SAM CLI:

```
SAM CLI, version 0.33.0
```

You're now ready to start development.

## Next Steps

You're now ready to begin building your own serverless applications using AWS SAM! If you want to start with sample serverless applications, choose one of the following links:

- [Tutorial: Deploying a Hello World Application \(p. 11\)](#) – Step-by-step instructions to download, build, and deploy a simple serverless application.
- [AWS SAM example applications in GitHub](#) – Sample applications in the AWS SAM GitHub repository that you can further experiment with.

# Setting Up AWS Credentials

The AWS SAM command line interface (CLI) requires you to set AWS credentials so that it can make calls to AWS services on your behalf. For example, the AWS SAM CLI makes calls to Amazon S3 and AWS CloudFormation.

You might have already set AWS credentials to work with AWS tools, like one of the AWS SDKs or the AWS CLI. If you haven't, this topic shows you the recommended approaches for setting AWS credentials.

To set AWS credentials, you must have the *access key ID* and your *secret access key* for the IAM user you want to configure. For information about access key IDs and secret access keys, see [Managing Access Keys for IAM Users](#) in the *IAM User Guide*.

Next, determine whether you have the AWS CLI installed. Then follow the instructions in one of the following sections:

## Using the AWS CLI

If you have the AWS CLI installed, use the **aws configure** command and follow the prompts:

```
$ aws configure
AWS Access Key ID [None]: your_access_key_id
AWS Secret Access Key [None]: your_secret_access_key
Default region name [None]:
Default output format [None]:
```

For information about the **aws configure** command, see [Quickly Configuring the AWS CLI](#) in the *AWS Command Line Interface User Guide*.

## Not Using the AWS CLI

If you don't have the AWS CLI installed, you can either create a credentials file or set environment variables:

- **Credentials file** – You can set credentials in the AWS credentials file on your local system. This file must be located in one of the following locations:
  - `~/.aws/credentials` on Linux or macOS
  - `C:\Users\USERNAME\.aws\credentials` on Windows

This file should contain lines in the following format:

```
[default]
aws_access_key_id = your_access_key_id
aws_secret_access_key = your_secret_access_key
```

- **Environment variables** – You can set the `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` environment variables.

To set these variables on Linux or macOS, use the **export** command:

```
export AWS_ACCESS_KEY_ID=your_access_key_id
export AWS_SECRET_ACCESS_KEY=your_secret_access_key
```

To set these variables on Windows, use the **set** command:

```
set AWS_ACCESS_KEY_ID=your_access_key_id
set AWS_SECRET_ACCESS_KEY=your_secret_access_key
```


## Tutorial: Deploying a Hello World Application

In this guide, you download, build, and deploy a sample Hello World application using AWS SAM. You then test the application in the AWS Cloud, and optionally test it locally on your development host.

This application implements a simple API backend. It consists of an API Gateway endpoint and a Lambda function. When you send a GET request to the API Gateway endpoint, the Lambda function is invoked. This function returns a `hello world` message.

The following diagram shows the components of this application:





The following is a preview of commands that you run to create your Hello World application. For more details about each of these commands, see the sections later in this page

```
#Step 1 - Download a sample application
sam init

#Step 2 - Build your application
cd sam-app
sam build

#Step 3 - Deploy your application
sam deploy --guided
```

## Prerequisites

This guide assumes that you've completed the steps in the [Installing the AWS SAM CLI \(p. 3\)](#) for your OS. It assumes that you've done the following:

1. Created an AWS account.
2. Configured IAM permissions.
3. Installed Docker. Note: Docker is only a prerequisite for testing your application locally.
4. Installed Homebrew. Note: Homebrew is only a prerequisite for Linux and macOS.
5. Installed the AWS SAM CLI. Note: Make sure you have version 0.33.0 or later. You can check which version you have by executing the command `sam --version`.

## Step 1: Download a Sample AWS SAM Application

**Command to run:**

```
sam init
```

Follow the on-screen prompts. For this tutorial we recommend you choose AWS Quick Start Templates, the runtime of your choice, and the Hello World Example.

**Example output:**

```
-----  
Generating application:  
-----  
Name: sam-app  
Runtime: python3.7  
Dependency Manager: pip  
Application Template: hello-world  
Output Directory: .  
  
Next steps can be found in the README file at ./sam-app/README.md
```

### What AWS SAM is doing:

This command creates a directory with the name you provided as the project name. The contents of the project directory are similar to the following (these contents are created when one of the Python runtimes and the Hello World Example are chose):

```
sam-app/  
  ### README.md  
  ### events/  
  #   ### event.json  
  ### hello_world/  
  #   ### __init__.py  
  #   ### app.py           #Contains your AWS Lambda handler logic.  
  #   ### requirements.txt #Contains any Python dependencies the application requires,  
used for sam build  
  ### template.yaml       #Contains the AWS SAM template defining your application's AWS  
resources.  
  ### tests/  
    ### unit/  
      ### __init__.py  
      ### test_handler.py
```

There are three especially important files:

- `template.yaml`: Contains the AWS SAM template that defines your application's AWS resources.
- `hello_world/app.py`: Contains your actual Lambda handler logic.
- `hello_world/requirements.txt`: Contains any Python dependencies that the application requires, and is used for `sam build`.

## Step 2: Build Your Application

### Command to run:

First change into the project directory (that is, the directory where the `template.yaml` file for the sample application is located; by default is `sam-app`), then run this command:

```
sam build
```

### Example output:

```
Build Succeeded
```

```
Built Artifacts   : .aws-sam/build
Built Template    : .aws-sam/build/template.yaml

Commands you can use next
=====
[*] Invoke Function: sam local invoke
[*] Deploy: sam deploy --guided
```

### What AWS SAM is doing:

The AWS SAM CLI comes with abstractions for a number of Lambda runtimes to build your dependencies, and copies the source code into staging folders so that everything is ready to be packaged and deployed. The `sam build` command builds any dependencies that your application has, and copies your application source code to folders under `aws-sam/build` to be zipped and uploaded to Lambda.

You can see the following top-level tree under `.aws-sam`:

```
.aws_sam/
  ### build/
    ### HelloWorldFunction/
      ### template.yaml
```

`HelloWorldFunction` is a directory that contains your `app.py` file, as well as third-party dependencies that your application uses.

## Step 3: Deploy Your Application to the AWS Cloud

### Command to run:

```
sam deploy --guided
```

Follow the on-screen prompts. You can just respond with `Enter` to accept the default options provided in the interactive experience.

### Example output:

```
Deploying with following values
=====
Stack name           : sam-app
Region               : us-east-1
Confirm changeset    : False
Deployment s3 bucket : sam-bucket
Capabilities          : [ "CAPABILITY_IAM" ]
Parameter overrides  : {}

Initiating deployment
=====

Waiting for changeset to be created..

CloudFormation stack changeset
-----
Operation      LogicalResourceId
ResourceType

+ Add
HelloWorldFunctionHelloWorldPermissionProd  AWS::Lambda::Permission
```

+ Add	AWS::ApiGateway::Deployment	ServerlessRestApiDeployment47fc2d5f9d
+ Add	AWS::ApiGateway::Stage	ServerlessRestApiProdStage
+ Add	AWS::ApiGateway::RestApi	ServerlessRestApi
* Modify	AWS::IAM::Role	HelloWorldFunctionRole
* Modify	AWS::Lambda::Function	HelloWorldFunction
-----		
2019-11-21 14:33:24 - Waiting for stack create/update to complete		
CloudFormation events from changeset		
-----		
ResourceStatus	ResourceType	
LogicalResourceId	ResourceStatusReason	
-----		
UPDATE_IN_PROGRESS	AWS::IAM::Role	
HelloWorldFunctionRole	-	
UPDATE_COMPLETE	AWS::IAM::Role	
HelloWorldFunctionRole	-	
UPDATE_IN_PROGRESS	AWS::Lambda::Function	
HelloWorldFunction	-	
UPDATE_COMPLETE	AWS::Lambda::Function	
HelloWorldFunction	-	
CREATE_IN_PROGRESS	AWS::ApiGateway::RestApi	
ServerlessRestApi	-	
CREATE_COMPLETE	AWS::ApiGateway::RestApi	
ServerlessRestApi	-	
CREATE_IN_PROGRESS	AWS::ApiGateway::RestApi	Resource creation Initiated
ServerlessRestApi		
CREATE_IN_PROGRESS	AWS::ApiGateway::Deployment	Resource creation Initiated
ServerlessRestApiDeployment47fc2d5		f9d
CREATE_IN_PROGRESS	AWS::Lambda::Permission	
HelloWorldFunctionHelloWorldPermis	Resource creation Initiated	sionProd
CREATE_IN_PROGRESS	AWS::Lambda::Permission	
HelloWorldFunctionHelloWorldPermis	-	sionProd
CREATE_IN_PROGRESS	AWS::ApiGateway::Deployment	
ServerlessRestApiDeployment47fc2d5	-	f9d
CREATE_COMPLETE	AWS::ApiGateway::Deployment	
ServerlessRestApiDeployment47fc2d5	-	f9d
CREATE_IN_PROGRESS	AWS::ApiGateway::Stage	
ServerlessRestApiProdStage	-	
CREATE_IN_PROGRESS	AWS::ApiGateway::Stage	Resource creation Initiated
ServerlessRestApiProdStage		
CREATE_COMPLETE	AWS::ApiGateway::Stage	
ServerlessRestApiProdStage	-	
CREATE_COMPLETE	AWS::Lambda::Permission	
HelloWorldFunctionHelloWorldPermis	-	sionProd
UPDATE_COMPLETE_CLEANUP_IN_PROGRES	AWS::CloudFormation::Stack	sam-app
-		
S		
UPDATE_COMPLETE	AWS::CloudFormation::Stack	sam-app
-		
-----		
Stack sam-app outputs:		
-----		

OutputKey-Description	OutputValue
----- HelloWorldFunctionIamRole - Implicit IAM Role created for Hello World arn:aws:iam::123456789012:role/sam-app-function HelloWorldFunctionRole-104VTJ0TST7M0 HelloWorldApi - API Gateway endpoint URL for Prod stage for Hello World https://0ks2zue0zh.execute-api.us-east-1.amazonaws.com/Prod/hello/ function HelloWorldFunction - Hello World Lambda Function ARN arn:aws:lambda:us-east-1:123456789012:function:sam-app-  HelloWorldFunction-1TY92MJX0BXU5 -----	
Successfully created/updated stack - sam-app in us-east-1	

### What AWS SAM is doing:

This command deploys your application to the AWS cloud. It takes the deployment artifacts you build with the `sam build` command, packages and uploads them to an Amazon S3 bucket created by AWS SAM CLI, and deploys the application using AWS CloudFormation. In the output of the `deploy` command you can see the changes being made to your AWS CloudFormation stack.

If your application created a HTTP endpoint, the Outputs generated by `sam deploy` also show you the endpoint URL for your test application. You can use `curl` to send a request to your application using that endpoint URL. For example:

```
curl https://<restapi.id>.execute-api.us-east-1.amazonaws.com/Prod/hello/
```

You should see output like the following after successfully deploying your application:

```
{"message": "hello world"}
```

If you see `{"message": "hello world"}` after executing the `curl` command, it means that you've successfully deployed your serverless application to AWS, and are calling your live Lambda function. Otherwise, see the [Troubleshooting \(p. 18\)](#) section later in this tutorial.

## Step 4: Testing Your Application Locally (Optional)

When you're developing your application, you might also find it useful to test locally. The AWS SAM CLI provides the `sam local` command to run your application using Docker containers that simulate the execution environment of Lambda. There are two options to do this:

- Host your API locally
- Invoke your Lambda function directly

This step describes both options.

### Host Your API Locally

#### Command to run:

```
sam local start-api
```

**Example output:**

```
2019-07-12 15:27:58 Mounting HelloWorldFunction at http://127.0.0.1:3000/hello [GET]
2019-07-12 15:27:58 You can now browse to the above endpoints to invoke your functions.
You do not need to restart/reload SAM CLI while working on your functions, changes will be
reflected instantly/automatically. You only need to restart SAM CLI if you update your AWS
SAM template
2019-07-12 15:27:58 * Running on http://127.0.0.1:3000/ (Press CTRL+C to quit)

Fetching lambci/lambda:python3.7 Docker container
image.....
2019-07-12 15:28:56 Mounting /<working-development-path>/sam-app/.aws-sam/build/
HelloWorldFunction as /var/task:ro,delegated inside runtime container
START RequestId: 52fdcf07-2182-154f-163f-5f0f9a621d72 Version: $LATEST
END RequestId: 52fdcf07-2182-154f-163f-5f0f9a621d72
REPORT RequestId: 52fdcf07-2182-154f-163f-5f0f9a621d72 Duration: 4.42 ms Billed
Duration: 100 ms Memory Size: 128 MB Max Memory Used: 22 MB
2019-07-12 15:28:58 No Content-Type given. Defaulting to 'application/json'.
2019-07-12 15:28:58 127.0.0.1 - - [12/Jul/2019 15:28:58] "GET /hello HTTP/1.1" 200 -
```

It might take a while for the Docker image to load. After it's loaded, you can use `curl` to send a request to your application that's running on your local host:

```
curl http://127.0.0.1:3000/hello
```

**Example output:**

```
2019-07-12 15:29:57 Invoking app.lambda_handler (python3.7)
2019-07-12 15:29:57 Found credentials in shared credentials file: ~/.aws/credentials

Fetching lambci/lambda:python3.7 Docker container image.....
2019-07-12 15:29:58 Mounting /<working-development-path>/sam-app/.aws-sam/build/
HelloWorldFunction as /var/task:ro,delegated inside runtime container
START RequestId: 52fdcf07-2182-154f-163f-5f0f9a621d72 Version: $LATEST
END RequestId: 52fdcf07-2182-154f-163f-5f0f9a621d72
REPORT RequestId: 52fdcf07-2182-154f-163f-5f0f9a621d72 Duration: 7.92 ms Billed
Duration: 100 ms Memory Size: 128 MB Max Memory Used: 22 MB
{"statusCode":200,"body":{"\message\": \"hello world\"}}
```

**What AWS SAM is doing:**

The `start-api` command starts up a local endpoint that replicates your REST API endpoint. It downloads an execution container that you can run your function locally in. The end result is the same output that you saw when you called your function in the AWS Cloud.

## Making One-off Invocations

**Command to run:**

```
sam local invoke "HelloWorldFunction" -e events/event.json
```

**Example output:**

```
2019-07-01 14:08:42 Found credentials in shared credentials file: ~/.aws/credentials
2019-07-01 14:08:42 Invoking app.lambda_handler (python3.7)
```

```

Fetching lambci/lambda:python3.7 Docker container
image.....
2019-07-01 14:09:39 Mounting /<working-development-path>/sam-app/.aws-sam/build/
HelloWorldFunction as /var/task:ro,delegated inside runtime container
START RequestId: 52fdcf07-2182-154f-163f-5f0f9a621d72 Version: $LATEST
END RequestId: 52fdcf07-2182-154f-163f-5f0f9a621d72
REPORT RequestId: 52fdcf07-2182-154f-163f-5f0f9a621d72    Duration: 3.51 ms    Billed
Duration: 100 ms    Memory Size: 128 MB    Max Memory Used: 22 MB
{"statusCode":200,"body":{"\"message\": \"hello world\""}}

```

### What AWS SAM is doing:

The `invoke` command directly invokes your Lambda functions, and can pass input event payloads that you provide. With this command, you pass the event payload in the file `event.json` that's provided by the sample application.

Your initialized application came with a default `aws-proxy` event for API Gateway. A number of values are prepopulated for you. In this case, the `HelloWorldFunction` doesn't care about the particular values, so a stubbed request is OK. You can specify a number of values to be substituted in to the request to simulate what you would expect from an actual request. This following is an example of generating your own input event and comparing the output with the default `event.json` object:

```

sam local generate-event apigateway aws-proxy --body "" --path "hello" --method GET > api-
event.json
diff api-event.json event.json

```

### Example output:

```

<  "body": "",
---
>  "body": {"\"message\": \"hello world\""},
4,6c4,6
<  "path": "/hello",
<  "httpMethod": "GET",
<  "isBase64Encoded": true,
---
>  "path": "/path/to/resource",
>  "httpMethod": "POST",
>  "isBase64Encoded": false,
11c11
<    "proxy": "/hello"
---
>    "proxy": "/path/to/resource"
56c56
<    "path": "/prod/hello",
---
>    "path": "/prod/path/to/resource",
58c58
<    "httpMethod": "GET",
---
>    "httpMethod": "POST",

```

## Troubleshooting

### SAM CLI error: "no such option: --app-template"

When executing `sam init`, you see the following error:

```
Error: no such option: --app-template
```

This means that you are using an older version of the AWS SAM CLI that does not support the `--app-template` parameter. To fix this, you can either update your version of AWS SAM CLI to 0.33.0 or later, or omit the `--app-template` parameter from the `sam init` command.

## SAM CLI error: "no such option: --guided"

When executing `sam deploy`, you see the following error:

```
Error: no such option: --guided
```

This means that you are using an older version of the AWS SAM CLI that does not support the `--guided` parameter. To fix this, you can either update your version of AWS SAM CLI to 0.33.0 or later, or omit the `--guided` parameter from the `sam deploy` command.

## SAM CLI error: "Failed to create managed resources: Unable to locate credentials"

When executing `sam deploy`, you see the following error:

```
Error: Failed to create managed resources: Unable to locate credentials
```

This means that you have not set up AWS credentials to enable the AWS SAM CLI to make AWS service calls. To fix this, you must set up AWS credentials. For more information, see [Setting Up AWS Credentials \(p. 10\)](#).

## SAM CLI error: "Running AWS SAM projects locally requires Docker. Have you got it installed?"

When executing `sam local start-api`, you see the following error:

```
Error: Running AWS SAM projects locally requires Docker. Have you got it installed?
```

This means that you do not have Docker properly installed. Docker is required to test your application locally. To fix this, follow the instructions for installing Docker for your development host.

For instructions on installing Docker on your development host, go to [Installing the AWS SAM CLI \(p. 3\)](#), choose the appropriate platform, and follow the instructions in the section titled **Install Docker**.

## Curl Error: "Missing Authentication Token"

When trying to invoke the API Gateway endpoint, you see the following error:



```
{"message": "Missing Authentication Token"}
```

This means that you've attempted to send a request to the correct domain, but the URI isn't recognizable. To fix this, verify the full URL, and update the `curl` command with the correct URL.

## Curl Error: "curl: (6) Could not resolve: ..."

When trying to invoke the API Gateway endpoint, you see the following error:

```
curl: (6) Could not resolve: endpointdomain (Domain name not found)
```

This means that you've attempted to send a request to an invalid domain. This can happen if your serverless application failed to deploy successfully, or if you have a typo in your `curl` command. Verify that the application was deployed successfully by using the AWS CloudFormation console or AWS CLI, and that your `curl` command is correct.

## Clean Up

If you no longer need the AWS resources you created by running this tutorial, you can remove them by deleting the AWS CloudFormation stack that you deployed.

To delete the AWS CloudFormation stack created with this tutorial using the AWS Management Console, follow these steps:

1. Sign in to the AWS Management Console and open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
2. In the left navigation pane, choose **Stacks**.
3. In the list of stacks, choose **aws-sam-getting-started**.
4. Choose **Delete**.

When done, the status of the of the stack will change to **DELETE\_COMPLETE**.

Alternatively, you can delete the AWS CloudFormation stack by executing the following AWS CLI command:

```
aws cloudformation delete-stack --stack-name aws-sam-getting-started --region region
```

## Verify Deleted Stack

For both methods of deleting the AWS CloudFormation stack, you can verify it was deleted by going to the <https://console.aws.amazon.com/cloudformation>, choosing **Stacks** in the left navigation pane, and choosing **Deleted** in the dropdown to the right of the search text box. You should see your stack name **aws-sam-getting-started** in the list of deleted stacks.

## Conclusion

In this tutorial, you've done the following:

1. Created, built, and deployed a serverless application to AWS with AWS SAM.
2. Tested your application locally by using the AWS SAM CLI and Docker.

3. Deleted the AWS resources that you no longer need.

## Next Steps

You're now ready to start building your own applications using the AWS SAM CLI.

To help you get started, you can download any of the example applications from the AWS SAM GitHub repository. To access this repository, see [AWS SAM example applications](#).

# AWS Serverless Application Model (AWS SAM) Specification

You use the AWS SAM specification to define your serverless application. This section provides details for the AWS SAM template sections, resources types, resource properties, data types, resource attributes, intrinsic functions, and API Gateway extensions that you can use in AWS SAM templates.

AWS SAM templates are an extension of AWS CloudFormation templates, with some additional components that make them easier to work with. For the full reference for AWS CloudFormation templates, see [AWS CloudFormation Template Reference](#) in the *AWS CloudFormation User Guide*.

## Topics

- [AWS SAM Template Anatomy](#) (p. 22)
- [AWS SAM Resource and Property Reference](#) (p. 28)
- [Resource Attributes](#) (p. 117)
- [Intrinsic Functions](#) (p. 118)
- [API Gateway Extensions](#) (p. 118)

## AWS SAM Template Anatomy

The format of an AWS SAM template closely follows the format of an AWS CloudFormation template, which is described in [Template Anatomy](#) in the *AWS CloudFormation User Guide*.

The primary differences between AWS SAM templates and AWS CloudFormation templates are the following:

- **Transform declaration.** The declaration `Transform: AWS::Serverless-2016-10-31` is required for AWS SAM templates. This declaration identifies an AWS CloudFormation template as an AWS SAM template. For more information about transforms, see [Transform](#) in the *AWS CloudFormation User Guide*.
- **Globals section.** The Globals section is unique to AWS SAM. It defines properties that are common to all your serverless functions and APIs. All the `AWS::Serverless::Function`, `AWS::Serverless::Api`, and `AWS::Serverless::SimpleTable` resources inherit the properties that are defined in the Globals section. For more information about the Globals section, see [Globals Section of the Template](#) (p. 24) in the *AWS Serverless Application Model Developer Guide*.
- **Resources section.** In AWS SAM templates the Resources section can contain a combination of AWS CloudFormation resources and AWS SAM resources. For more information about AWS CloudFormation resources, see [AWS Resource and Property Types Reference](#) in the *AWS CloudFormation User Guide*. For more information about AWS SAM resources see [AWS SAM Resource and Property Reference](#) (p. 28) in the *AWS Serverless Application Model Developer Guide*.

All other sections of an AWS SAM template correspond to the AWS CloudFormation template section of the same name.

## YAML

The following example shows a YAML-formatted template fragment.

```
Transform: AWS::Serverless-2016-10-31

Globals:
  set of globals

Description:
  String

Metadata:
  template metadata

Parameters:
  set of parameters

Mappings:
  set of mappings

Conditions:
  set of conditions

Resources:
  set of resources

Outputs:
  set of outputs
```

## Template Sections

Templates include several major sections. `Transform` and `Resources` are the only required sections.

The sections in a template can be in any order. However, as you build your template, it can be helpful to use the logical order that's shown in the following list. This is because the values in one section might refer to values from a previous section.

### Transform (required)

For AWS SAM templates, you must include this section with a value of `AWS::Serverless-2016-10-31`.

Additional transforms are optional. For more information about transforms, see [Transform](#) in the *AWS CloudFormation User Guide*.

### Globals (optional) (p. 24)

A section in your AWS SAM template to define properties that are common to all your serverless functions, APIs, and simple tables. All the `AWS::Serverless::Function`, `AWS::Serverless::Api`, and `AWS::Serverless::SimpleTable` resources inherit the properties that are defined in the `Globals` section.

This section is unique to AWS SAM. There isn't a corresponding section in AWS CloudFormation templates.

### Description (optional)

A text string that describes the template.

This section corresponds directly with the `Description` section of AWS CloudFormation templates.

### Metadata (optional)

Objects that provide additional information about the template.

This section corresponds directly with the Metadata section of AWS CloudFormation templates.

### Parameters (optional)

Values to pass to your template at runtime (when you create or update a stack). You can refer to parameters from the `Resources` and `Outputs` sections of the template.

This section corresponds directly with the Parameters section of AWS CloudFormation templates.

### Mappings (optional)

A mapping of keys and associated values that you can use to specify conditional parameter values, similar to a lookup table. You can match a key to a corresponding value by using the `Fn::FindInMap` intrinsic function in the `Resources` and `Outputs` sections.

This section corresponds directly with the Mappings section of AWS CloudFormation templates.

### Conditions (optional)

Conditions that control whether certain resources are created or whether certain resource properties are assigned a value during stack creation or update. For example, you could conditionally create a resource that depends on whether the stack is for a production or test environment.

This section corresponds directly with the Conditions section of AWS CloudFormation templates.

### Resources (required)

Specifies the stack resources and their properties, such as an Amazon EC2 instance or an Amazon S3 bucket. You can refer to resources in the `Resources` and `Outputs` sections of the template.

This section is similar to the `Resources` section of AWS CloudFormation templates. In AWS SAM templates, this section can contain AWS SAM resources in addition to AWS CloudFormation resources.

### Outputs (optional)

Describes the values that are returned whenever you view your stack's properties. For example, you can declare an output for an S3 bucket name, and then call the `aws cloudformation describe-stacks` AWS CLI command to view the name.

This section corresponds directly with the `Outputs` section of AWS CloudFormation templates.

## Globals Section of the Template

Resources in an AWS SAM template tend to have shared configuration, such as `Runtime`, `Memory`, `VPCConfig`, `Environment`, and `Cors`. Instead of duplicating this information in every resource, you can write them once in the `Globals` section and let your resources inherit them.

The `Globals` section is supported by the `AWS::Serverless::Function`, `AWS::Serverless::Api`, and `AWS::Serverless::SimpleTable` resources.

Example:

```
Globals:
  Function:
    Runtime: nodejs6.10
    Timeout: 180
    Handler: index.handler
```

```
Environment:
  Variables:
    TABLE_NAME: data-table

Resources:
  HelloWorldFunction:
    Type: AWS::Serverless::Function
    Properties:
      Environment:
        Variables:
          MESSAGE: "Hello From SAM"

  ThumbnailFunction:
    Type: AWS::Serverless::Function
    Properties:
      Events:
        Thumbnail:
          Type: Api
          Properties:
            Path: /thumbnail
            Method: POST
```

In this example, both `HelloWorldFunction` and `ThumbnailFunction` use "nodejs6.10" for Runtime, "180" seconds for Timeout, and "index.handler" for Handler. `HelloWorldFunction` adds the `MESSAGE` environment variable, in addition to the inherited `TABLE_NAME`. `ThumbnailFunction` inherits all the `Globals` properties and adds an API event source.

## Supported Resources and Properties

Currently, AWS SAM supports the following resources and properties:

```
Globals:
  Function:
    # Properties of AWS::Serverless::Function
    Handler:
    Runtime:
    CodeUri:
    DeadLetterQueue:
    Description:
    MemorySize:
    Timeout:
    VpcConfig:
    Environment:
    Tags:
    Tracing:
    KmsKeyArn:
    Layers:
    AutoPublishAlias:
    DeploymentPreference:
    PermissionsBoundary:
    ReservedConcurrentExecutions:

  Api:
    # Properties of AWS::Serverless::Api
    # Also works with Implicit APIs
    Auth:
    Name:
    DefinitionUri:
    CacheClusterEnabled:
    CacheClusterSize:
    Variables:
    EndpointConfiguration:
    MethodSettings:
```

```
BinaryMediaTypes:
MinimumCompressionSize:
Cors:
GatewayResponses:
AccessLogSetting:
CanarySetting:
TracingEnabled:
OpenApiVersion:

SimpleTable:
  # Properties of AWS::Serverless::SimpleTable
  SSESpecification:
```

## Implicit APIs

*Implicit APIs* are APIs that are created by AWS SAM when you declare an API in the `Events` section. You can use `Globals` to override all the properties of implicit APIs.

## Unsupported Properties

The following properties are **not** supported in the `Globals` section. We made the explicit call to not support them because it either made the template hard to understand, or opened a scope for potential security issues.

```
Function:
  Role:
  Policies:
  FunctionName:
  Events:

Api:
  StageName:
  DefinitionBody:
```

## Overridable Properties

Properties that are declared in the `Globals` section can be overridden by the resource. For example, you can add new variables to an environment variable map, or you can override globally declared variables. But the resource **cannot** remove a property that's specified in the `Globals` environment variables map. More generally, the `Globals` section declares properties that are shared by all your resources. Some resources can provide new values for globally declared properties, but they can't completely remove them. If some resources use a property but others don't, then you must not declare them in the `Globals` section.

The following sections describe how overriding works for different data types.

### Primitive Data Types Are Replaced

Primitive data types include strings, numbers, Booleans, and so on.

The value specified in the `Resources` section **replaces** the value in the `Globals` section.

Example:

```
Globals:
  Function:
    Runtime: nodejs4.3

Resources:
```

```
MyFunction:
  Type: AWS::Serverless::Function
  Properties:
    Runtime: python3.6
```

The Runtime for MyFunction is set to python3.6.

## Maps Are Merged

Maps are also known as dictionaries or collections of key-value pairs.

Map entries in the resource are **merged** with global map entries. If there are duplicates, the resource entry overrides the global entry.

Example:

```
Globals:
  Function:
    Environment:
      Variables:
        STAGE: Production
        TABLE_NAME: global-table

Resources:
  MyFunction:
    Type: AWS::Serverless::Function
    Properties:
      Environment:
        Variables:
          TABLE_NAME: resource-table
          NEW_VAR: hello
```

The environment variables of MyFunction are set to the following:

```
{
  "STAGE": "Production",
  "TABLE_NAME": "resource-table",
  "NEW_VAR": "hello"
}
```

## Lists Are Additive

Lists are also known as arrays.

Entries in the Globals section are **prepended** to the list in the Resource section.

Example:

```
Globals:
  Function:
    VpcConfig:
      SecurityGroupIds:
        - sg-123
        - sg-456

Resources:
  MyFunction:
    Type: AWS::Serverless::Function
    Properties:
      VpcConfig:
        SecurityGroupIds:
```



```
- sg-first
```

The SecurityGroupIds for MyFunction's VpcConfig are set to the following:

```
[ "sg-123", "sg-456", "sg-first" ]
```

## AWS SAM Resource and Property Reference

This section contains reference information for the AWS SAM resource and property types.

### Topics

- [AWS::Serverless::Api](#) (p. 28)
- [AWS::Serverless::Application](#) (p. 53)
- [AWS::Serverless::Function](#) (p. 56)
- [AWS::Serverless::HttpApi](#) (p. 102)
- [AWS::Serverless::LayerVersion](#) (p. 111)
- [AWS::Serverless::SimpleTable](#) (p. 114)

## AWS::Serverless::Api

Creates a collection of Amazon API Gateway resources and methods that can be invoked through HTTPS endpoints.

An [AWS::Serverless::Api](#) (p. 28) resource need not be explicitly added to a AWS Serverless Application Definition template. A resource of this type is implicitly created from the union of [Api](#) events defined on [AWS::Serverless::Function](#) (p. 56) resources defined in the template that do not refer to an [AWS::Serverless::Api](#) (p. 28) resource.

An [AWS::Serverless::Api](#) (p. 28) resource should be used to define and document the API using OpenApi, which provides more ability to configure the underlying Amazon API Gateway resources.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
Type: AWS::Serverless::Api
Properties:
  AccessLogSetting: AccessLogSetting
  Auth: ApiAuth (p. 34)
  BinaryMediaTypes: List
  CacheClusterEnabled: Boolean
  CacheClusterSize: String
  CanarySetting: CanarySetting
  Cors: String | CorsConfiguration (p. 48)
  DefinitionBody: String
  DefinitionUri: String | ApiDefinition (p. 47)
  Domain: DomainConfiguration (p. 50)
  EndpointConfiguration: String
  GatewayResponses: Map
  MethodSettings: MethodSettings
  MinimumCompressionSize: Integer
```

```
Models: Map
Name: String
OpenApiVersion: String
StageName: String
Tags: Map
TracingEnabled: Boolean
Variables: Map
```

## Properties

### AccessLogSetting

Configures Access Log Setting for a stage.

Type: [AccessLogSetting](#)

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [AccessLogSetting](#) property of an `AWS::ApiGateway::Stage`.

### Auth

Configure authorization to control access to your API Gateway API.

For more information about configuring access using AWS SAM see [Controlling Access to API Gateway APIs \(p. 125\)](#) in the AWS Serverless Application Model Developer Guide.

Type: [ApiAuth \(p. 34\)](#)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### BinaryMediaTypes

List of MIME types that your API could return. Use this to enable binary support for APIs. Use ~1 instead of / in the mime types.

Type: List

Required: No

*CloudFormation Compatibility:* This property is similar to the [BinaryMediaTypes](#) property of an `AWS::ApiGateway::RestApi`. The list of `BinaryMediaTypes` is added to both the AWS CloudFormation resource and the OpenAPI document.

### CacheClusterEnabled

Indicates whether cache clustering is enabled for the stage.

Type: Boolean

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [CacheClusterEnabled](#) property of an `AWS::ApiGateway::Stage`.

### CacheClusterSize

The stage's cache cluster size.

Type: String

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [CacheClusterSize](#) property of an `AWS::ApiGateway::Stage`.

#### CanarySetting

Configure a canary setting to a stage of a regular deployment.

Type: [CanarySetting](#)

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [CanarySetting](#) property of an `AWS::ApiGateway::Stage`.

#### Cors

Manage Cross-origin resource sharing (CORS) for all your API Gateway APIs. Specify the domain to allow as a string or specify a dictionary with additional Cors configuration. NOTE: Cors requires SAM to modify your OpenAPI definition. So, it works only inline OpenApi defined with DefinitionBody.

For more information about CORS, see [Enable CORS for an API Gateway REST API Resource](#) in the Amazon API Gateway Developer Guide..

**NOTE:** API Gateway requires literal values to be a quoted string, so you must include single quotes in the `Allow___` values. For example, `"www.example.com"` is correct whereas `www.example.com` is not correct.

Type: String | [CorsConfiguration](#) (p. 48)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### DefinitionBody

OpenAPI specification that describes your API. If neither `DefinitionUri` nor `DefinitionBody` are specified, SAM will generate a `DefinitionBody` for you based on your template configuration.

Type: String

Required: No

*CloudFormation Compatibility:* This property is similar to the [Body](#) property of an `AWS::ApiGateway::RestApi`. If certain properties are provided, content may be inserted or modified into the `DefinitionBody` before being passed to CloudFormation. Properties include `Auth`, `BinaryMediaTypes`, `Cors`, `GatewayResponses`, `Models`, and an `EventSource` of type `Api` on for a corresponding `AWS::Serverless::Function`.

#### DefinitionUri

AWS S3 Uri, local file path, or location object of the the OpenAPI document defining the API. The AWS S3 object this property references must be a valid OpenAPI file. If neither `DefinitionUri` nor `DefinitionBody` are specified, SAM will generate a `DefinitionBody` for you based on your template configuration.

If a local file path is provided, the template must go through the workflow that includes the `sam deploy` or `sam package` command, in order for the definition to be transformed properly.

Intrinsic functions are not supported in external OpenApi files referenced by `DefinitionUri`. Use instead the `DefinitionBody` property with the [Include Transform](#) to import an OpenApi definition into the template.

Type: String | [ApiDefinition](#) (p. 47)

Required: No

*CloudFormation Compatibility:* This property is similar to the [BodyS3Location](#) property of an `AWS::ApiGateway::RestApi`. The nested Amazon S3 properties are named differently.

#### Domain

Configures a custom domain for this API Gateway API.

Type: [DomainConfiguration](#) (p. 50)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### EndpointConfiguration

Specify the type of endpoint for API endpoint.

Valid values are REGIONAL, EDGE, or PRIVATE.

Type: String

Required: No

*CloudFormation Compatibility:* This property is similar to the [EndpointConfiguration](#) property of an `AWS::ApiGateway::RestApi`. AWS SAM only accepts a single endpoint configuration string.

#### GatewayResponses

Configures Gateway Responses for an API. Gateway Responses are responses returned by API Gateway, either directly or through the use of Lambda Authorizers. For more information, see the documentation for the [Api Gateway OpenApi extension for Gateway Responses](#).

Type: Map

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### MethodSettings

Configures all settings for API stage including Logging, Metrics, CacheTTL, Throttling.

Type: [MethodSettings](#)

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [MethodSettings](#) property of an `AWS::ApiGateway::Stage`.

#### MinimumCompressionSize

Allow compression of response bodies based on client's Accept-Encoding header. Compression is triggered when response body size is greater than or equal to your configured threshold. The maximum body size threshold is 10 MB (10,485,760 Bytes). - The following compression types are supported: gzip, deflate, and identity.

*Type:* Integer

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [MinimumCompressionSize](#) property of an `AWS::ApiGateway::RestApi`.

#### Models

The schemas to be used by your API methods. These schemas can be described using JSON or YAML. See the Examples section at the bottom of this page for example models.

*Type:* Map

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### Name

A name for the API Gateway RestApi resource

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Name](#) property of an `AWS::ApiGateway::RestApi`.

#### OpenApiVersion

Version of OpenApi to use. This can either be 2.0 for the Swagger specification, or one of the OpenApi 3.0 versions, like 3.0.1. For more information about OpenAPI, see the [OpenAPI Specification](#).

**Note:** Setting this property to any valid value will also remove the stage `Stage` that SAM creates.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### StageName

The name of the stage, which API Gateway uses as the first path segment in the invoke Uniform Resource Identifier (URI).

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is similar to the [StageName](#) property of an `AWS::ApiGateway::Stage`. It is required in SAM, but not required in API Gateway

*Additional Notes:* The Implicit API has a stage name of "prod"

#### Tags

A map (string to string) that specifies the tags to be added to this API Gateway stage. Keys and values are limited to alphanumeric characters. Keys can be 1 to 127 Unicode characters in length and cannot be prefixed with `aws:`. Values can be 1 to 255 Unicode characters in length.

*Type:* Map

*Required:* No

*CloudFormation Compatibility:* This property is similar to the [Tags](#) property of an `AWS::ApiGateway::Stage`. The `Tags` property in SAM consists of Key:Value pairs; in CloudFormation it consists of a list of Tag objects. When the stack is created, SAM will automatically add a `lambda:createdBy:SAM` tag to this API Gateway stage.

`TracingEnabled`

Indicates whether active tracing with X-Ray is enabled for the stage.

*Type:* Boolean

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [TracingEnabled](#) property of an `AWS::ApiGateway::Stage`.

`Variables`

A map (string to string) that defines the stage variables, where the variable name is the key and the variable value is the value. Variable names are limited to alphanumeric characters. Values must match the following regular expression: `[A-Za-z0-9._~:/?#&=,- ]+`.

*Type:* Map

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Variables](#) property of an `AWS::ApiGateway::Stage`.

## Examples

### SimpleApiExample

A Hello World SAM template that contains a Lambda Function with an API endpoint.

YAML

```
AWSTemplateFormatVersion: '2010-09-09'
Transform: AWS::Serverless-2016-10-31
Description: AWS SAM template with a simple API definition
Resources:
  ApiGatewayApi:
    Type: AWS::Serverless::Api
    Properties:
      StageName: prod
  ApiFunction: # Adds a GET api endpoint at "/" to the ApiGatewayApi via an Api event
    Type: AWS::Serverless::Function
    Properties:
      Events:
        ApiEvent:
          Type: Api
          Properties:
            Path: /
            Method: get
            RestApiId:
              Ref: ApiGatewayApi
      Runtime: python3.7
      Handler: index.handler
      InlineCode: |
        def handler(event, context):
```

```
return {'body': 'Hello World!', 'statusCode': 200}
```

## ApiCorsExample

AWS SAM template with API defined in an external Swagger file along with Lambda integrations and CORS configurations. See [GitHub](#) for the full example.

### YAML

```
Type: AWS::Serverless::Api
Properties:
  StageName: Prod
  # Allows www.example.com to call these APIs
  # SAM will automatically add AllowMethods with a list of methods for this API
  Cors: "'www.example.com'"
  DefinitionBody: # Pull in an OpenApi definition from S3
    'Fn::Transform':
      Name: 'AWS::Include'
      # Replace "bucket" with your bucket name
      Parameters:
        Location: s3://bucket/swagger.yaml
```

## ApiCognitoAuthExample

AWS SAM template with an API that uses AWS Cognito to authorize requests against the API. See [GitHub](#) for the full example.

### YAML

```
Type: AWS::Serverless::Api
Properties:
  StageName: Prod
  Cors: ""
  Auth:
    DefaultAuthorizer: MyCognitoAuthorizer
    Authorizers:
      MyCognitoAuthorizer:
        UserPoolArn:
          Fn::GetAtt: [MyCognitoUserPool, Arn]
```

## ApiAuth

Configure authorization to control access to your API Gateway API.

For more information and examples for configuring access using AWS SAM see [Controlling Access to API Gateway APIs \(p. 125\)](#) in the AWS Serverless Application Model Developer Guide.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
AddDefaultAuthorizerToCorsPreflight: Boolean
ApiKeyRequired: Boolean
Authorizers: CognitoAuthorizer (p. 37) | LambdaTokenAuthorizer (p. 42)
              | LambdaRequestAuthorizer (p. 39)
DefaultAuthorizer: String
InvokeRole: String
```

`ResourcePolicy`: [ResourcePolicyStatement](#) (p. 45)

## Properties

### AddDefaultAuthorizerToCorsPreflight

If the `DefaultAuthorizer` and `Cors` properties are set, then setting `AddDefaultAuthorizerToCorsPreflight` will cause the default authorizer to be added to the `Options` property in the `OpenAPI` section.

*Type*: Boolean

*Required*: No

*Default*: True

*CloudFormation Compatibility*: This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### ApiKeyRequired

If set to true then an API key is required for all API events. For more information about API keys see [Create and Use Usage Plans with API Keys](#) in the API Gateway Developer Guide.

*Type*: Boolean

*Required*: No

*CloudFormation Compatibility*: This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Authorizers

The authorizer used to control access to your API Gateway API.

For more information, see [Controlling Access to API Gateway APIs](#) (p. 125) in the AWS Serverless Application Model Developer Guide.

*Type*: [CognitoAuthorizer](#) (p. 37) | [LambdaTokenAuthorizer](#) (p. 42) | [LambdaRequestAuthorizer](#) (p. 39)

*Required*: No

*Default*: None

*CloudFormation Compatibility*: This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

*Additional Notes*: SAM adds the `Authorizers` to the `OpenApi` definition of an `Api`.

### DefaultAuthorizer

Specify a default authorizer for an API Gateway API, which will be used for authorizing API calls by default.

**Note**: If the `Api EventSource` for the function associated with this `Api` is configured to use IAM Permissions, then this property must be set to `AWS_IAM`, otherwise an error will result.

*Type*: String

*Required*: No



*Default:* None

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### InvokeRole

Sets integration credentials for all resources and methods to this value.

Supported values: CALLER\_CREDENTIALS, NONE, IAM Role Arn.

CALLER\_CREDENTIALS maps to `arn:aws:iam::*:user/*`, which uses the caller credentials to invoke the endpoint.

*Type:* String

*Required:* No

*Default:* CALLER\_CREDENTIALS

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### ResourcePolicy

Configure Resource Policy for all methods and paths on an API.

*Type:* [ResourcePolicyStatement \(p. 45\)](#)

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

*Additional Notes:* This setting can also be defined on individual `AWS::Serverless::Function` using the [ApiFunctionAuth \(p. 75\)](#). This is required for APIs with `EndpointConfiguration: PRIVATE`.

## Examples

### CognitoAuth

#### Cognito Auth Example

#### YAML

```
Auth:
  Authorizers:
    MyCognitoAuth:
      UserPoolArn:
        Fn::GetAtt:
          - MyUserPool
          - Arn
      AuthType: "COGNITO_USER_POOLS"
  DefaultAuthorizer: MyCognitoAuth
  InvokeRole: CALLER_CREDENTIALS
  AddDefaultAuthorizerToCorsPreflight: false
  ApiKeyRequired: false
  ResourcePolicy:
    CustomStatements: [{
      "Effect": "Allow",
      "Principal": "*",
      "Action": "execute-api:Invoke",
```

```
"Resource": "execute-api:/Prod/PUT/get",
"Condition": {
  "IpAddress": {
    "aws:SourceIp": "1.2.3.4"
  }
}
}]
IpRangeBlacklist: ['10.20.30.40']
```

## CognitoAuthorizer

Define a Amazon Cognito User Pool authorizer.

For more information and examples, see [Controlling Access to API Gateway APIs \(p. 125\)](#) in the AWS Serverless Application Model Developer Guide.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
AuthorizationScopes: List
Identity: CognitoAuthorizationIdentity (p. 38)
UserPoolArn: String
```

### Properties

#### AuthorizationScopes

List of authorization scopes for this authorizer.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### Identity

This property can be used to specify an `IdentitySource` in an incoming request for an authorizer

*Type:* [CognitoAuthorizationIdentity \(p. 38\)](#)

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### UserPoolArn

Can refer to a user pool/specify a userpool arn to which you want to add this cognito authorizer

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### CognitoAuth

#### Cognito Auth Example

##### YAML

```
Auth:
  Authorizers:
    MyCognitoAuth:
      AuthorizationScopes:
        - scope1
        - scope2
      UserPoolArn:
        Fn::GetAtt:
          - MyCognitoUserPool
          - Arn
      Identity:
        Header: MyAuthorizationHeader
        ValidationExpression: myauthvalidationexpression
```

### CognitoAuthorizationIdentity

This property can be used to specify an IdentitySource in an incoming request for an authorizer. For more information about IdentitySource see the [ApiGateway Authorizer OpenApi extension](#).

#### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

##### YAML

```
Header: String
ReauthorizeEvery: Integer
ValidationExpression: String
```

## Properties

### Header

Specify the header name for Authorization in the OpenApi definition.

*Type:* String

*Required:* No

*Default:* Authorization

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### ReauthorizeEvery

The time-to-live (TTL) period, in seconds, that specifies how long API Gateway caches authorizer results. If you specify a value greater than 0, API Gateway caches the authorizer responses. By default, API Gateway sets this property to 300. The maximum value is 3600, or 1 hour.

*Type:* Integer

*Required:* No

*Default:* 300

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### ValidationExpression

Specify a validation expression for validating the incoming Identity

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### CognitoAuthIdentity

#### YAML

```
Identity:
  Header: MyCustomAuthHeader
  ValidationExpression: Bearer.*
  ReauthorizeEvery: 30
```

### LambdaRequestAuthorizer

Configure a Lambda Authorizer to control access to your Api with a Lambda function.

For more information and examples, see [Controlling Access to API Gateway APIs \(p. 125\)](#) in the AWS Serverless Application Model Developer Guide.

#### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
AuthorizationScopes: List
FunctionArn: String
FunctionInvokeRole: String
FunctionPayloadType: String
Identity: LambdaRequestAuthorizationIdentity (p. 41)
```

## Properties

#### AuthorizationScopes

List of authorization scopes for this authorizer.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### FunctionArn

Specify the function arn of the Lambda function which gives the authorization to Api

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### FunctionInvokeRole

Adds authorizer credentials to the OpenApi definition of the Lambda authorizer.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### FunctionPayloadType

This property can be used to define the type of Lambda Authorizer for an Api.

Supported values: TOKEN and REQUEST

*Type:* String

*Required:* No

*Default:* TOKEN

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### Identity

This property can be used to specify an `IdentitySource` in an incoming request for an authorizer

*Type:* [LambdaRequestAuthorizationIdentity \(p. 41\)](#)

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### LambdaRequestAuth

#### YAML

```
Authorizer:
  MyLambdaRequestAuth:
    FunctionPayloadType: REQUEST
    FunctionArn:
      Fn::GetAtt:
        - MyAuthFunction
        - Arn
    FunctionInvokeRole:
```

```
Fn::GetAtt:
  - LambdaAuthInvokeRole
  - Arn
Identity:
  Headers:
    - Authorization1
```

## LambdaRequestAuthorizationIdentity

This property can be used to specify an IdentitySource in an incoming request for an authorizer. For more information about IdentitySource see the [ApiGateway Authorizer OpenApi extension](#).

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
Context: List
Headers: List
QueryString: List
ReauthorizeEvery: Integer
StageVariables: List
```

### Properties

#### Context

Converts the given context strings to the mapping expressions of format `context.contextString`.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### Headers

Converts the headers to comma-separated string of mapping expressions of format `method.request.header.name`.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### QueryString

Converts the given query strings to comma-separated string of mapping expressions of format `method.request.querystring.queryString`.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### ReauthorizeEvery

The time-to-live (TTL) period, in seconds, that specifies how long API Gateway caches authorizer results. If you specify a value greater than 0, API Gateway caches the authorizer responses. By default, API Gateway sets this property to 300. The maximum value is 3600, or 1 hour.

*Type:* Integer

*Required:* No

*Default:* 300

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### StageVariables

Converts the given stage variables to comma-separated string of mapping expressions of format `stageVariables.stageVariable`.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Examples

#### LambdaRequestIdentity

##### YAML

```
Identity:
  QueryStrings:
    - auth
  Headers:
    - Authorization
  StageVariables:
    - VARIABLE
  Context:
    - authcontext
  ReauthorizeEvery: 100
```

#### LambdaTokenAuthorizer

Configure a Lambda Authorizer to control access to your Api with a Lambda function.

For more information and examples, see [Controlling Access to API Gateway APIs \(p. 125\)](#) in the AWS Serverless Application Model Developer Guide.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

##### YAML

```
AuthorizationScopes: List
FunctionArn: String
FunctionInvokeRole: String
```

```
FunctionPayloadType: String  
Identity: LambdaTokenAuthorizationIdentity (p. 44)
```

## Properties

### AuthorizationScopes

List of authorization scopes for this authorizer.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### FunctionArn

Specify the function arn of the Lambda function which gives the authorization to Api

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### FunctionInvokeRole

Adds authorizer credentials to the OpenApi definition of the Lambda authorizer.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### FunctionPayloadType

This property can be used to define the type of Lambda Authorizer for an Api.

Supported values: TOKEN and REQUEST

*Type:* String

*Required:* No

*Default:* TOKEN

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Identity

This property can be used to specify an IdentitySource in an incoming request for an authorizer.

*Type:* [LambdaTokenAuthorizationIdentity \(p. 44\)](#)

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.



## Examples

### LambdaTokenAuth

#### YAML

```
Authorizers:
  MyLambdaTokenAuth:
    FunctionArn:
      Fn::GetAtt:
        - MyAuthFunction
        - Arn
    Identity:
      Header: MyCustomAuthHeader # OPTIONAL; Default: 'Authorization'
      ValidationExpression: mycustomauthexpression # OPTIONAL
      ReauthorizeEvery: 20 # OPTIONAL; Service Default: 300
```

### BasicLambdaTokenAuth

#### YAML

```
Authorizers:
  MyLambdaTokenAuth:
    FunctionArn:
      Fn::GetAtt:
        - MyAuthFunction
        - Arn
```

### LambdaTokenAuthorizationIdentity

This property can be used to specify an IdentitySource in an incoming request for an authorizer. For more information about IdentitySource see the [ApiGateway Authorizer OpenApi extension](#).

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
ReauthorizeEvery: Integer
ValidationExpression: String
```

## Properties

### ReauthorizeEvery

The time-to-live (TTL) period, in seconds, that specifies how long API Gateway caches authorizer results. If you specify a value greater than 0, API Gateway caches the authorizer responses. By default, API Gateway sets this property to 300. The maximum value is 3600, or 1 hour.

*Type:* Integer

*Required:* No

*Default:* 300

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### ValidationExpression

Specify a validation expression for validating the incoming Identity.

Type: String

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Examples

#### LambdaTokenIdentity

##### YAML

```
Identity:
  Header: Auth
  ValidationExpression: Bearer.*
  ReauthorizeEvery: 30
```

### ResourcePolicyStatement

Configure Resource Policy for all methods and paths on an API.

#### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

##### YAML

```
AwsAccountBlacklist: List
AwsAccountWhitelist: List
CustomStatements: List
IpRangeBlacklist: List
IpRangeWhitelist: List
SourceVpcBlacklist: List
SourceVpcWhitelist: List
```

### Properties

#### AwsAccountBlacklist

Resource Policy statements will be generated and attached to the Api for blacklisting the given list of Aws accounts

Type: List

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### AwsAccountWhitelist

Resource Policy statements will be generated and attached to the Api for whitelisting the given list of Aws accounts

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### CustomStatements

A list of resource policy statements can be given for an Api

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### IpRangeBlacklist

Resource Policy statements will be generated and attached to the Api for blacklisting the given list of Ip addresses or ranges

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### IpRangeWhitelist

Resource Policy statements will be generated and attached to the Api for whitelisting the given list of Ip addresses or ranges

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### SourceVpcBlacklist

Resource Policy statements will be generated and attached to the Api for blacklisting the given list of Source Vpcs or Vpc endpoints

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### SourceVpcWhitelist

Resource Policy statements will be generated and attached to the Api for whitelisting the given list of Source Vpcs or Vpc endpoints

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### SourceVpcBlacklist

Blacklisting source VPC or VPC endpoint

#### YAML

```
Auth:
  ResourcePolicy:
    CustomStatements: [{
      "Effect": "Allow",
      "Principal": "*",
      "Action": "execute-api:Invoke",
      "Resource": "execute-api:/Prod/PUT/get",
      "Condition": {
        "IpAddress": {
          "aws:SourceIp": "1.2.3.4"
        }
      }
    }]
  IpRangeBlacklist: ['10.20.30.40', '1.2.3.4']
  SourceVpcBlacklist: ["vpce-1a2b3c4d"]
  AwsAccountWhitelist: ['123456789101']
```

## ApiDefinition

An OpenAPI document defining the API.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Bucket: String
Key: String
Version: String
```

## Properties

### Bucket

The name of the Amazon S3 bucket where the OpenAPI file is stored.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [Bucket](#) property of the `AWS::ApiGateway::RestApi` S3Location data type.

### Key

The Amazon S3 key of the OpenAPI file.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [key](#) property of the `AWS::ApiGateway::RestApi` S3Location data type.

#### Version

For versioned objects, the version of the OpenAPI file.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [version](#) property of the `AWS::ApiGateway::RestApi` S3Location data type.

## Examples

### Definition Uri example

#### API Definition example

#### YAML

```
DefinitionUri:
  Bucket: mybucket-name
  Key: mykey-name
  Version: 121212
```

## CorsConfiguration

Manage cross-origin resource sharing (CORS) for your API Gateway APIs. Specify the domain to allow as a string or specify a dictionary with additional Cors configuration. NOTE: Cors requires SAM to modify your OpenAPI definition, so it only works with inline OpenApi defined in the `DefinitionBody` property.

For more information about CORS, see [Enable CORS for an API Gateway REST API Resource](#) in the Amazon API Gateway Developer Guide.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
AllowCredentials: String
AllowHeaders: String
AllowMethods: String
AllowOrigin: String
MaxAge: String
```

## Properties

### AllowCredentials

Boolean indicating whether request is allowed to contain credentials.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### AllowHeaders

String of headers to allow.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### AllowMethods

String containing the HTTP methods to allow.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### AllowOrigin

String of origin to allow.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### MaxAge

String containing the number of seconds to cache CORS Preflight request.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### CorsConfiguration

Cors Configuration example.

#### YAML

```
Cors:
  AllowMethods: "'POST, GET'"
  AllowHeaders: "'X-Forwarded-For'"
  AllowOrigin: "'www.example.com'"
  MaxAge: "'600'"
```

```
AllowCredentials: True
```

## DomainConfiguration

Configuration for a custom domain for an API.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
BasePath: List
CertificateArn: String
DomainName: String
EndpointConfiguration: String
Route53: Route53Configuration (p. 51)
```

### Properties

#### BasePath

List of basepaths to be configured with the API Gateway Domain Name.

Type: List

Required: No

Default: /

*CloudFormation Compatibility:* This property is similar to the [BasePath](#) property of an `AWS::ApiGateway::BasePathMapping`. SAM will create multiple `AWS::ApiGateway::BasePathMapping` resources, one per `BasePath` specified in this property.

#### CertificateArn

The reference to an AWS-managed certificate for use by the endpoint for this domain name. AWS Certificate Manager is the only supported source.

Type: String

Required: Yes

*CloudFormation Compatibility:* This property is similar to the [CertificateArn](#) property of an `AWS::ApiGateway::DomainName`. If `EndpointConfiguration` is set to `REGIONAL` (the default value), `CertificateArn` maps to [RegionalCertificateArn](#) in `AWS::ApiGateway::DomainName`. If the `EndpointConfiguration` is set to `EDGE`, `CertificateArn` maps to [CertificateArn](#) in `AWS::ApiGateway::DomainName`.

*Additional Notes:* For an `EDGE` endpoint, the certificate must be created in the `us-east-1` region.

#### DomainName

The custom domain name for your API Gateway API. Uppercase letters are not supported.

Type: String

Required: Yes

*CloudFormation Compatibility:* This property is passed directly to the [DomainName](#) property of an `AWS::ApiGateway::DomainName`.

#### EndpointConfiguration

Property to define the type of API Gateway endpoint to be mapped to the custom domain. The value of this property controls how the `CertificateArn` property gets mapped in AWS CloudFormation. See `CertificateArn` above.

Valid values are `REGIONAL` or `EDGE`.

*Type:* String

*Required:* No

*Default:* `REGIONAL`

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### Route53

Property that adds Route53 configuration based on the values defined.

*Type:* [Route53Configuration](#) (p. 51)

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### DomainName

DomainName example

#### YAML

```
Domain:
  DomainName: www.example.com
  CertificateArn: arn-example
  EndpointConfiguration: EDGE
  Route53:
    HostedZoneId: xyz
  BasePath:
    - /foo
    - /bar
```

### Route53Configuration

#### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
DistributionDomainName: String
EvaluateTargetHealth: Boolean
HostedZoneId: String
```



## Properties

### DistributionDomainName

Use this to configure a custom distribution of the API Custom Domain Name.

*Type:* String

*Required:* No

*Default:* Use the API Gateway Distribution

*CloudFormation Compatibility:* This property is passed directly to the [DNSName](#) property of an `AWS::Route53::RecordSetGroup` `AliasTarget`.

*Additional Notes:* Domain name of a [cloudfront distribution](#).

### EvaluateTargetHealth

When `EvaluateTargetHealth` is true, an alias record inherits the health of the referenced AWS resource, such as an ELB load balancer or another record in the hosted zone.

*Type:* Boolean

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [EvaluateTargetHealth](#) property of an `AWS::Route53::RecordSetGroup` `AliasTarget`.

*Additional Notes:* You can't set `EvaluateTargetHealth` to true when the alias target is a CloudFront distribution.

### HostedZoneId

HostedZoneId for the domain name.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [HostedZoneId](#) property of an `AWS::Route53::RecordSetGroup` `RecordSet`.

## Examples

### Route 53 Configuration Example

Shows how to configure route 53

#### YAML

```
Domain:
  DomainName: www.example.com
  CertificateArn: arn-example
  EndpointConfiguration: EDGE
Route53:
  HostedZoneId: xyz
  EvaluateTargetHealth: true
```

```
DistributionDomainName: xyz
```

## AWS::Serverless::Application

Embeds a serverless application from the [AWS Serverless Application Repository](#) or from an Amazon S3 bucket as a nested application. Nested applications are deployed as nested [AWS::CloudFormation::Stack](#) resources, which can contain multiple other resources including other [AWS::Serverless::Application](#) (p. 53) resources.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Type: AWS::Serverless::Application
Properties:
  Location: String | ApplicationLocationObject (p. 55)
  NotificationARNs: List
  Parameters: Map
  Tags: Map
  TimeoutInMinutes: Integer
```

### Properties

#### Location

Template URL, file path, or location object of a nested application.

If a template URL is provided, it must follow the format specified in the [CloudFormation TemplateUrl](#) documentation and contain a valid CloudFormation or SAM template. An [ApplicationLocationObject](#) (p. 55) can be used to specify an application that has been published to the [AWS Serverless Application Repository](#).

If a local file path is provided, the template must go through the workflow that includes the `sam deploy` or `sam package` command, in order for the application to be transformed properly.

Type: *String* | [ApplicationLocationObject](#) (p. 55)

Required: Yes

**CloudFormation Compatibility:** This property is similar to the [TemplateURL](#) property of an `AWS::CloudFormation::Stack`. The CloudFormation version does not take an [ApplicationLocationObject](#) (p. 55) to retrieve an application from the AWS Serverless Application Repository.

#### NotificationARNs

A list of existing Amazon SNS topics where notifications about stack events are sent.

Type: *List*

Required: No

**CloudFormation Compatibility:** This property is passed directly to the [NotificationARNs](#) property of an `AWS::CloudFormation::Stack`.

#### Parameters

Application parameter values.

Type: Map

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [Parameters](#) property of an `AWS::CloudFormation::Stack`.

#### Tags

A map (string to string) that specifies the tags to be added to this application. Keys and values are limited to alphanumeric characters. Keys can be 1 to 127 Unicode characters in length and cannot be prefixed with `aws:`. Values can be 1 to 255 Unicode characters in length.

Type: Map

Required: No

*CloudFormation Compatibility:* This property is similar to the [Tags](#) property of an `AWS::CloudFormation::Stack`. The Tags property in SAM consists of Key:Value pairs; in CloudFormation it consists of a list of Tag objects. When the stack is created, SAM will automatically add a `lambda:createdBy:SAM` tag to this application. In addition, if this application is from the AWS Serverless Application Repository, then SAM will also automatically the two additional tags `serverlessrepo:applicationId:ApplicationId` and `serverlessrepo:semanticVersion:SemanticVersion`.

#### TimeoutInMinutes

The length of time, in minutes, that AWS CloudFormation waits for the nested stack to reach the `CREATE_COMPLETE` state. The default is no timeout. When AWS CloudFormation detects that the nested stack has reached the `CREATE_COMPLETE` state, it marks the nested stack resource as `CREATE_COMPLETE` in the parent stack and resumes creating the parent stack. If the timeout period expires before the nested stack reaches `CREATE_COMPLETE`, AWS CloudFormation marks the nested stack as failed and rolls back both the nested stack and parent stack.

Type: Integer

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [TimeoutInMinutes](#) property of an `AWS::CloudFormation::Stack`.

## Return Values

### Ref

When the logical ID of this resource is provided to the Ref intrinsic function, it returns the resource name of the underlying `AWS::CloudFormation::Stack` resource.

For more information about using the Ref function, see [Ref](#).

### Fn::GetAtt

Fn::GetAtt returns a value for a specified attribute of this type. The following are the available attributes and sample return values.

For more information about using Fn::GetAtt, see [Fn::GetAtt](#).

## GetAtt

`Outputs.ApplicationOutputName`

The value of the stack output with name `ApplicationOutputName`.

## Examples

### SAR Application

Application that uses a template from the Serverless Application Repository

#### YAML

```
Type: AWS::Serverless::Application
Properties:
  Location:
    ApplicationId: 'arn:aws:serverlessrepo:us-east-1:012345678901:applications/my-
application'
    SemanticVersion: 1.0.0
  Parameters:
    StringParameter: parameter-value
    IntegerParameter: 2
```

### Normal-Application

Application from an S3 url

#### YAML

```
Type: AWS::Serverless::Application
Properties:
  Location: https://s3.amazonaws.com/demo-bucket/template.yaml
```

## ApplicationLocationObject

An application that has been published to the [AWS Serverless Application Repository](#).

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
ApplicationId: String
SemanticVersion: String
```

### Properties

`ApplicationId`

The Amazon Resource Name (ARN) of the application.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

SemanticVersion

The semantic version of the application.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### my-application

Example application location object

#### YAML

```
Location:
  ApplicationId: 'arn:aws:serverlessrepo:us-east-1:012345678901:applications/my-
application'
  SemanticVersion: 1.0.0
```

## AWS::Serverless::Function

Creates a Lambda function, IAM execution role, and event source mappings which trigger the function.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Type: AWS::Serverless::Function
Properties:
  AssumeRolePolicyDocument: JSON
  AutoPublishAlias: String
  CodeUri: String | FunctionCode (p. 101)
  DeadLetterQueue: Map | DeadLetterQueue (p. 63)
  DeploymentPreference: DeploymentPreference (p. 64)
  Description: String
  Environment: Environment
  EventInvokeConfig: EventInvokeConfiguration (p. 67)
  Events: EventSource (p. 72)
  FunctionName: String
  Handler: String
  InlineCode: String
  KmsKeyArn: String
  Layers: List
  MemorySize: Integer
  PermissionsBoundary: String
  Policies: String | List | Map
  ProvisionedConcurrencyConfig: ProvisionedConcurrencyConfig
  ReservedConcurrentExecutions: Integer
```

```
Role: String
Runtime: String
Tags: Map
Timeout: Integer
Tracing: String
VersionDescription: String
VpcConfig: VpcConfig
```

## Properties

### AssumeRolePolicyDocument

Adds an AssumeRolePolicyDocument for the default created Role for this function. If this property is not specified SAM adds a default assume role for this function.

Type: JSON

Required: No

*CloudFormation Compatibility:* This property is similar to the `AssumeRolePolicyDocument` property of an `AWS::IAM::Role`. SAM adds this property to the generated IAM Role for this Function. If a Role ARN is provided for this Function, this property does nothing.

### AutoPublishAlias

Name of the Lambda alias. For more information about Lambda aliases, see [AWS Lambda Function Aliases](#).

This AWS SAM property generates two additional resources: an `AWS::Lambda::Version` resource and an `AWS::Lambda::Alias` resource.

The `AWS::Lambda::Version` resource has a logical id of `<function-logical-id>Version<sha>`, where the `<sha>` is 10 digits of the SHA256 of CodeUri. This resource can be referenced in intrinsic functions by using the logical ID or `<function-logical-id>.Version`.

The `AWS::Lambda::Alias` resource has a logical id of `<function-logical-id>Alias<alias-name>`, where `<alias-name>` is the alias name specified in this property. This resource can be referenced in intrinsic functions by using the logical ID or `<function-logical-id>.Alias`.

For examples that use this property, see [Deploying Serverless Applications Gradually \(p. 146\)](#).

Type: String

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### CodeUri

AWS S3 Uri, local file path, or [FunctionCode \(p. 101\)](#) object of the function code.

If an AWS S3 Uri or [FunctionCode \(p. 101\)](#) object is provided, the AWS S3 object referenced must be a valid [Lambda deployment package](#).

If a local file path is provided, the template must go through the workflow that includes the `sam deploy` or `sam package` command, in order for the code to be transformed properly.

**Note:** Either CodeUri or InlineCode is required.

Type: String | [FunctionCode \(p. 101\)](#)

*Required:* Conditional

*CloudFormation Compatibility:* This property is similar to the [Code](#) property of an `AWS::Lambda::Function`. The nested Amazon S3 properties are named differently.

#### DeadLetterQueue

Configures SNS topic or SQS queue where Lambda sends events that it can't process. For more information about dead letter queue functionality, see [AWS Lambda Function Dead Letter Queues](#).

*Type:* Map | [DeadLetterQueue \(p. 63\)](#)

*Required:* No

*CloudFormation Compatibility:* This property is similar to the [DeadLetterConfig](#) property of an `AWS::Lambda::Function`. In CloudFormation the type is derived from the `TargetArn`, whereas in SAM you must pass the type along with the `TargetArn`.

#### DeploymentPreference

Settings to enable gradual Lambda deployments.

If a `DeploymentPreference` object is specified, SAM will create an [AWS::CodeDeploy::Application](#) called `ServerlessDeploymentApplication` (one per stack), an [AWS::CodeDeploy::DeploymentGroup](#) called `<function-logical-id>DeploymentGroup`, and an [AWS::IAM::Role](#) called `CodeDeployServiceRole`.

*Type:* [DeploymentPreference \(p. 64\)](#)

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

*See Also:* See the [Deploying Serverless Applications Gradually \(p. 146\)](#) documentation for more information about this property.

#### Description

Description of the function.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Description](#) property of an `AWS::Lambda::Function`.

#### Environment

Configuration for the runtime environment.

*Type:* [Environment](#)

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Environment](#) property of an `AWS::Lambda::Function`.

#### EventInvokeConfig

The object describing event invoke config on a Lambda function.

*Type:* [EventInvokeConfiguration \(p. 67\)](#)

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### Events

Specifies the events that trigger this function. Events consist of a type and a set of properties that depend on the type.

Type: [EventSource](#) (p. 72)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### FunctionName

A name for the function. If you don't specify a name, a unique name will be generated for you.

Type: String

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [FunctionName](#) property of an `AWS::Lambda::Function`.

#### Handler

Function within your code that is called to begin execution.

Type: String

Required: Yes

*CloudFormation Compatibility:* This property is passed directly to the [Handler](#) property of an `AWS::Lambda::Function`.

#### InlineCode

Lambda function code written directly in the template.

**Note:** Either `CodeUri` or `InlineCode` is required.

Type: String

Required: Conditional

*CloudFormation Compatibility:* This property is passed directly to the [ZipFile](#) property of the `AWS::Lambda::Function` Code data type.

#### KmsKeyArn

The Amazon Resource Name (ARN) of an AWS Key Management Service (AWS KMS) key that Lambda uses to encrypt and decrypt your function's environment variables.

Type: String

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [KmsKeyArn](#) property of an `AWS::Lambda::Function`.

#### Layers

List of `LayerVersion` ARNs that should be used by this function. The order specified here is the order that they will be imported when running the Lambda function.



Type: List

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [Layers](#) property of an `AWS::Lambda::Function`.

#### MemorySize

Size of the memory allocated per invocation of the function in MB.

Type: Integer

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [MemorySize](#) property of an `AWS::Lambda::Function`.

#### PermissionsBoundary

ARN of a permissions boundary to use for this function's execution role. This property only works if the Role is generated for you.

Type: String

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [PermissionsBoundary](#) property of an `AWS::IAM::Role`.

#### Policies

One or more policies that this function needs, which will be appended to the default role for this function.

This property accepts a single string or a list of strings, and can be the name of AWS managed IAM policies or SAM Policy Templates, or inline IAM policy document(s) formatted in YAML. If the Role property is set, this property has no meaning.

For more information about AWS managed IAM policies, see [AWS Managed Policies](#). For more information about AWS SAM Policy Templates, see [AWS SAM Policy Templates \(p. 178\)](#). For more information about inline policies, see [Inline Policies](#).

Type: String | List | Map

Required: No

*CloudFormation Compatibility:* This property is similar to the [Policies](#) property of an `AWS::IAM::Role`. AWS SAM supports AWS Managed policy names and SAM Policy Templates in addition to JSON policy documents; CloudFormation only accepts JSON policy documents.

#### ProvisionedConcurrencyConfig

Provisioned concurrency configuration of a function's alias.

**Note:** `ProvisionedConcurrencyConfig` can only be specified if the `AutoPublishAlias` is set, otherwise an error will result.

Type: [ProvisionedConcurrencyConfig](#)

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [ProvisionedConcurrencyConfig](#) property of an `AWS::Lambda::Alias`.

#### ReservedConcurrentExecutions

The maximum of concurrent executions you want to reserve for the function.

For more information about this property see [AWS Lambda Function Scaling](#) in the AWS Lambda Developer Guide.

*Type:* Integer

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [ReservedConcurrentExecutions](#) property of an `AWS::Lambda::Function`.

#### Role

ARN of an IAM role to use as this function's execution role.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is similar to the [Role](#) property of an `AWS::Lambda::Function`. This is required in AWS CloudFormation but not in AWS SAM. If a role is not specified, one is created for you with a logical id of `<function-logical-id>Role`.

#### Runtime

The runtime environment.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [Runtime](#) property of an `AWS::Lambda::Function`.

#### Tags

A map (string to string) that specifies the tags added to the Lambda function and the corresponding IAM execution role. Keys and values are limited to alphanumeric characters. Keys can be 1 to 127 Unicode characters in length and cannot be prefixed with `aws:`. Values can be 1 to 255 Unicode characters in length.

*Type:* Map

*Required:* No

*CloudFormation Compatibility:* This property is similar to the [Tags](#) property of an `AWS::Lambda::Function`. The Tags property in SAM consists of Key:Value pairs; in CloudFormation it consists of a list of Tag objects. When the stack is created, SAM automatically adds a `lambda:createdBy:SAM` tag to this Lambda function and the corresponding IAM execution role.

#### Timeout

Maximum time that the function can run before it is killed in seconds.

*Type:* Integer

*Required:* No

*Default:* 3

*CloudFormation Compatibility:* This property is passed directly to the [Timeout](#) property of an `AWS::Lambda::Function`.

#### Tracing

String that specifies the function's X-Ray tracing mode. For more information about X-Ray, see [Using AWS X-Ray](#) in the AWS Lambda Developer Guide.

Supported values: Active and PassThrough

Type: String

Required: No

*CloudFormation Compatibility:* This property is similar to the [TracingConfig](#) property of an `AWS::Lambda::Function`. If Tracing is set to Active then AWS SAM adds the `arn:aws:iam::aws:policy/AWSXrayWriteOnlyAccess` policy to the Lambda role.

#### VersionDescription

A string that specifies the Description field which will be added on the new lambda version resource.

Type: String

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [Description](#) property of an `AWS::Lambda::Version`.

#### VpcConfig

Configuration to enable this function to access private resources within your VPC.

Type: [VpcConfig](#)

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [VpcConfig](#) property of an `AWS::Lambda::Function`.

## Return Values

### Ref

When the logical ID of this resource is provided to the Ref intrinsic function, it returns the resource name of the underlying Lambda function.

For more information about using the Ref function, see [Ref](#).

### Fn::GetAtt

Fn::GetAtt returns a value for a specified attribute of this type. The following are the available attributes and sample return values.

For more information about using Fn::GetAtt, see [Fn::GetAtt](#).

#### GetAtt

##### Arn

The Amazon Resource Name (ARN) of the underlying Lambda function.

## Examples

### Simple Function

Base case example of an AWS::Serverless::Function resource.

#### YAML

```
Type: AWS::Serverless::Function
Properties:
  Handler: index.handler
  Runtime: python3.6
  CodeUri: s3://bucket/key
```

### Function Properties Example

Example of an AWS::Serverless::Function that uses InlineCode, Tracing, Policies, and Layers.

#### YAML

```
Type: AWS::Serverless::Function
Properties:
  Handler: index.handler
  Runtime: python3.6
  InlineCode: |
    def handler(event, context):
      print("Hello, world!")
  ReservedConcurrentExecutions: 30
  Layers:
    - Ref: MyLayer
  Tracing: Active
  Timeout: 120
  Policies:
    - AWSLambdaExecute
    - Version: '2012-10-17'
      Statement:
        - Effect: Allow
          Action:
            - s3:GetObject
            - s3:GetObjectACL
          Resource: 'arn:aws:s3:::my-bucket/*'
```

## DeadLetterQueue

Specifies an SQS queue or SNS topic that AWS Lambda (Lambda) sends events to when it can't process them. For more information about dead letter queue functionality, see [AWS Lambda Function Dead Letter Queues](#).

SAM will automatically add appropriate permission to your Lambda function execution role to give Lambda service access to the resource. sqs:SendMessage will be added for SQS queues and sns:Publish for SNS topics.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
TargetArn: String
```

```
Type: String
```

## Properties

### TargetArn

The Amazon Resource Name (ARN) of an Amazon SQS queue or Amazon SNS topic.

Type: String

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [TargetArn](#) property of the `AWS::Lambda::Function DeadLetterConfig` data type.

### Type

The type of dead letter queue.

Supported values: SNS, SQS.

Type: String

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### DeadLetterQueue

Dead Letter Queue example for an SNS topic.

#### YAML

```
DeadLetterQueue:
  Type: SNS
  TargetArn: arn:aws:sns:us-east-2:123456789012:my-topic
```

## DeploymentPreference

Specifies the configurations to enable gradual Lambda deployments. For more information about configuring gradual Lambda deployments, see [Deploying Serverless Applications Gradually \(p. 146\)](#).

**Note:** You must specify an `AutoPublishAlias` in your [AWS::Serverless::Function \(p. 56\)](#) to use a `DeploymentPreference` object, otherwise an error will result.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Alarms: List
Enabled: Boolean
```

```
Hooks: Hooks \(p. 66\)  
Role: String  
TriggerConfigurations: List  
Type: String
```

## Properties

### Alarms

A list of CloudWatch alarms that you want to be triggered by any errors raised by the deployment.

Type: List

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Enabled

Whether this deployment preference is enabled.

Type: Boolean

Required: No

Default: True

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Hooks

Validation Lambda functions that are run before and after traffic shifting.

Type: [Hooks \(p. 66\)](#)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Role

An IAM role ARN that CodeDeploy will use for traffic shifting. An IAM role will not be created if this is provided.

Type: String

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### TriggerConfigurations

A list of trigger configurations you want to associate with the deployment group. Used to notify an SNS topic on lifecycle events.

Type: List

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [TriggerConfigurations](#) property of an `AWS::CodeDeploy::DeploymentGroup`.

#### Type

There are two categories of deployment types at the moment: Linear and Canary. For more information about available deployment types see [Deploying Serverless Applications Gradually](#) (p. 146).

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### DeploymentPreference

Example deployment preference

#### YAML

```
DeploymentPreference:
  Enabled: True
  Type: Canary10Percent10Minutes
  Alarms:
    - Ref: AliasErrorMetricGreaterThanZeroAlarm
    - Ref: LatestVersionErrorMetricGreaterThanZeroAlarm
  Hooks:
    PreTraffic:
      Ref: PreTrafficLambdaFunction
    PostTraffic:
      Ref: PostTrafficLambdaFunction
```

## Hooks

Validation Lambda functions that are run before and after traffic shifting.

**Note:** The Lambda functions referenced in this property configure the `CodeDeployLambdaAliasUpdate` object of the resulting [AWS::Lambda::Alias](#) resource. For more information, see [CodeDeployLambdaAliasUpdate Policy](#) in the AWS CloudFormation User Guide.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
PostTraffic: String
PreTraffic: String
```

## Properties

### PostTraffic

Lambda function that is run after traffic shifting.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### PreTraffic

Lambda function that is run before traffic shifting.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### Hooks

Example hook functions

#### YAML

```
Hooks:
  PreTraffic:
    Ref: PreTrafficLambdaFunction
  PostTraffic:
    Ref: PostTrafficLambdaFunction
```

## EventInvokeConfiguration

Configuration options for [asynchronous](#) Lambda Alias or Version invocations.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
DestinationConfig: EventInvokeDestinationConfiguration (p. 68)
MaximumEventAgeInSeconds: Integer
MaximumRetryAttempts: Integer
```

## Properties

#### DestinationConfig

A configuration object that specifies the destination of an event after Lambda processes it.

*Type:* [EventInvokeDestinationConfiguration \(p. 68\)](#)

*Required:* No



*CloudFormation Compatibility:* This property is similar to the [DestinationConfig](#) property of an `AWS::Lambda::EventInvokeConfig`. SAM requires an extra parameter, "Type", that does not exist in CloudFormation.

#### MaximumEventAgeInSeconds

The maximum age of a request that Lambda sends to a function for processing.

*Type:* Integer

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [MaximumEventAgeInSeconds](#) property of an `AWS::Lambda::EventInvokeConfig`.

#### MaximumRetryAttempts

The maximum number of times to retry before the function returns an error.

*Type:* Integer

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [MaximumRetryAttempts](#) property of an `AWS::Lambda::EventInvokeConfig`.

## Examples

### MaximumEventAgeInSeconds

#### MaximumEventAgeInSeconds example

##### YAML

```
EventInvokeConfig:
  MaximumEventAgeInSeconds: 60
  MaximumRetryAttempts: 2
  DestinationConfig:
    OnSuccess:
      Type: SQS
      Destination: arn:aws:sqs:us-west-2:012345678901:my-queue
    OnFailure:
      Type: Lambda
      Destination:
        Ref: DestinationLambda
```

## EventInvokeDestinationConfiguration

A configuration object that specifies the destination of an event after Lambda processes it.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

##### YAML

```
OnFailure: OnFailure (p. 69)
OnSuccess: OnSuccess (p. 70)
```

## Properties

### OnFailure

A destination for events that failed processing.

Type: [OnFailure \(p. 69\)](#)

Required: No

*CloudFormation Compatibility:* This property is similar to the [OnFailure](#) property of an `AWS::Lambda::EventInvokeConfig`. Requires `Type`, an additional SAM-only property.

### OnSuccess

A destination for events that were processed successfully.

Type: [OnSuccess \(p. 70\)](#)

Required: No

*CloudFormation Compatibility:* This property is similar to the [OnSuccess](#) property of an `AWS::Lambda::EventInvokeConfig`. Requires `Type`, an additional SAM-only property.

## Examples

### OnSuccess

#### OnSuccess example

#### YAML

```
EventInvokeConfig:
  DestinationConfig:
    OnSuccess:
      Type: SQS
      Destination: arn:aws:sqs:us-west-2:012345678901:my-queue
    OnFailure:
      Type: Lambda
      Destination:
        Ref: DestinationLambda
```

### OnFailure

A destination for events that failed processing.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Destination: String
Type: String
```

## Properties

### Destination

The Amazon Resource Name (ARN) of the destination resource.

Type: String

Required: Conditional

*CloudFormation Compatibility:* This property is similar to the [OnFailure](#) property of an `AWS::Lambda::EventInvokeConfig`. SAM will add any necessary permissions to the auto-generated IAM Role associated with this function to access the resource referenced in this property.

*Additional Notes:* If the type is Lambda/EventBridge, Destination is required.

### Type

Type of the resource referenced in the destination. Supported types are SQS, SNS, Lambda, and EventBridge.

Type: String

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

*Additional Notes:* If the type is SQS/SNS and the Destination property is left blank, then the SQS/SNS resource is auto generated by SAM. To reference the resource, use `<function-logical-id>.DestinationQueue` for SQS or `<function-logical-id>.DestinationTopic` for SNS. If the type is Lambda/EventBridge, Destination is required.

## Examples

### EventInvoke Configuration Example

Example of how to use EventInvokeConfig.

Since no Destination is given for the SQS OnSuccess configuration, SAM will create a SQS queue and add any necessary permissions. Since a Destination is given for the Lambda OnFailure configuration, SAM will only add the necessary permissions to the source Lambda function to call the destination Lambda function.

### YAML

```
EventInvokeConfig:
  DestinationConfig:
    OnSuccess:
      Type: SQS
    OnFailure:
      Type: Lambda
      Destination:
        Ref: DestinationLambda
```

### OnSuccess

A destination for events that were processed successfully.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

## YAML

```
Destination: String
Type: String
```

## Properties

### Destination

The Amazon Resource Name (ARN) of the destination resource.

Type: String

Required: Conditional

*CloudFormation Compatibility:* This property is similar to the `OnSuccess` property of an `AWS::Lambda::EventInvokeConfig`. SAM will add any necessary permissions to the auto-generated IAM Role associated with this function to access the resource referenced in this property.

*Additional Notes:* If the type is Lambda/EventBridge, Destination is required.

### Type

Type of the resource referenced in the destination. Supported types are SQS, SNS, Lambda, and EventBridge.

Type: String

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

*Additional Notes:* If the type is SQS/SNS and the Destination property is left blank, then the SQS/SNS resource is auto generated by SAM. To reference the resource, use `<function-logical-id>.DestinationQueue` for SQS or `<function-logical-id>.DestinationTopic` for SNS. If the type is Lambda/EventBridge, Destination is required.

## Examples

### EventInvoke Configuration Example

Example of how to use EventInvokeConfig.

Since no Destination is given for the SQS OnSuccess configuration, SAM will create a SQS queue and add any necessary permissions. Since a Destination is given for the Lambda OnFailure configuration, SAM will only add the necessary permissions to this Lambda function to call the destination Lambda function.

## YAML

```
EventInvokeConfig:
  DestinationConfig:
    OnSuccess:
      Type: SQS
    OnFailure:
      Type: Lambda
      Destination:
```

Ref: DestinationLambda

## EventSource

The object describing the source of events which trigger the function. Each event consists of a type and a set of properties that depend on that type. For more information about the properties of each event source, see the topic corresponding to that type.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Properties: S3 (p. 94) | SNS (p. 97) | Kinesis (p. 92) | DynamoDB (p. 84)
| SQS (p. 100) | Api (p. 73) | Schedule (p. 95) | CloudWatchEvent (p. 81)
| EventBridgeRule (p. 86) | CloudWatchLogs (p. 82) | IoTRule (p. 91)
| AlexaSkill (p. 73) | Cognito (p. 83) | HttpApi (p. 88)
Type: String
```

## Properties

### Properties

Object describing properties of this event mapping. The set of properties must conform to the defined Type.

Type: S3 (p. 94) | SNS (p. 97) | Kinesis (p. 92) | DynamoDB (p. 84) | SQS (p. 100) |  
Api (p. 73) | Schedule (p. 95) | CloudWatchEvent (p. 81) | EventBridgeRule (p. 86)  
| CloudWatchLogs (p. 82) | IoTRule (p. 91) | AlexaSkill (p. 73) | Cognito (p. 83) |  
HttpApi (p. 88)

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Type

The event type.

Supported values: S3, SNS, Kinesis, DynamoDB, SQS, Api, Schedule, CloudWatchEvent, CloudWatchLogs, IoTRule, AlexaSkill, Cognito.

Type: String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### API-Event

Example of using an API Event

## YAML

```
ApiEvent:
  Type: Api
  Properties:
    Method: get
    Path: /group/{user}
    RestApiId:
      Ref: MyApi
```

## AlexaSkill

Adds an Alexa Skills Kit Trigger.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

## YAML

```
SkillId: String
```

### Properties

#### SkillId

The Alexa Skill ID for your Alexa Skill. For more information about Skill ID see [Configure the trigger for a Lambda function](#) in the Alexa Skills Kit documentation.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Examples

#### AlexaSkillTrigger

Alexa Skill Event Example

## YAML

```
AlexaSkillEvent:
  Type: AlexaSkill
```

## Api

The object describing an event source with type Api. If an [AWS::Serverless::Api \(p. 28\)](#) resource is defined, the path and method values must correspond to an operation in the OpenApi definition of the API.

If no [AWS::Serverless::Api \(p. 28\)](#) is defined, the function input and output are a representation of the HTTP request and HTTP response.

For example, using the JavaScript API, the status code and body of the response can be controlled by returning an object with the keys `statusCode` and `body`.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

## YAML

```
Auth: ApiFunctionAuth (p. 75)
Method: String
Path: String
RequestModel: RequestModel (p. 79)
RequestParameters: String | RequestParameter (p. 80)
RestApiId: String
```

## Properties

### Auth

Auth configuration for this specific Api+Path+Method.

Useful for overriding the API's `DefaultAuthorizer` setting auth config on an individual path when no `DefaultAuthorizer` is specified or overriding the default `ApiKeyRequired` setting.

Type: [ApiFunctionAuth \(p. 75\)](#)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Method

HTTP method for which this function is invoked.

Type: String

Required: Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Path

Uri path for which this function is invoked. Must start with `/`.

Type: String

Required: Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### RequestModel

Request model to use for this specific Api+Path+Method. This should reference the name of a model specified in the `Models` section of an [AWS::Serverless::Api \(p. 28\)](#) resource.

Type: [RequestModel \(p. 79\)](#)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### RequestParameters

Request parameters configuration for this specific Api+Path+Method. All parameter names must start with `method.request` and must be limited to `method.request.header`, `method.request.querystring`, or `method.request.path`.

If a parameter is a string and not a Function Request Parameter Object, then `Required` and `Caching` will default to `False`.

Type: String | [RequestParameter \(p. 80\)](#)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### RestApiId

Identifier of a RestApi resource, which must contain an operation with the given path and method. Typically, this is set to reference an [AWS::Serverless::Api \(p. 28\)](#) resource defined in this template.

If not defined, a default [AWS::Serverless::Api \(p. 28\)](#) resource is created using a generated OpenApi document containing a union of all paths and methods defined by Api events defined in this template that do not specify a `RestApiId`.

This cannot reference an [AWS::Serverless::Api \(p. 28\)](#) resource defined in another template.

Type: String

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### ApiEvent

An example of Api Event

#### YAML

```
ApiEvent:
  Type: Api
  Properties:
    Path: /path
    Method: get
    RequestParameters:
      - method.request.header.Authorization
```

### ApiFunctionAuth

Configures authorization at the event level.

Configure Auth for a specific Api + Path + Method

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:



## YAML

```
ApiKeyRequired: Boolean
AuthorizationScopes: List
Authorizer: String
InvokeRole: String
ResourcePolicy: ResourcePolicyStatement (p. 77)
```

## Properties

### ApiKeyRequired

Requires an API key for this API path and method.

*Type:* Boolean

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### AuthorizationScopes

Authorization scopes to apply to this Api + Path + Method.

Scopes listed here will override any scopes applied by the `DefaultAuthorizer` if one exists.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Authorizer

The `Authorizer` for a specific Function

If you have specified a Global Authorizer on the API and want to make a specific Function public, override by setting `Authorizer` to `NONE`.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### InvokeRole

Specifies the `InvokeRole` to use for `AWS_IAM` authorization.

*Type:* String

*Required:* No

*Default:* `CALLER_CREDENTIALS`

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

*Additional Notes:* `CALLER_CREDENTIALS` maps to `arn:aws:iam::*:user/*`, which uses the caller credentials to invoke the endpoint.

#### ResourcePolicy

Configure Resource Policy for this path on an API.

Type: [ResourcePolicyStatement](#) (p. 77)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### Function-Auth

Specifying Authorization at Function level

#### YAML

```
Auth:
  ApiKeyRequired: true
  Authorizer: NONE
```

### ResourcePolicyStatement

Configure Resource Policy for all methods and paths on an API.

#### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
AwsAccountBlacklist: List
AwsAccountWhitelist: List
CustomStatements: List
IpRangeBlacklist: List
IpRangeWhitelist: List
SourceVpcBlacklist: List
SourceVpcWhitelist: List
```

## Properties

### AwsAccountBlacklist

Resource Policy statements will be generated and attached to the Api for blacklisting the given list of Aws accounts

Type: List

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### AwsAccountWhitelist

Resource Policy statements will be generated and attached to the Api for whitelisting the given list of Aws accounts

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### CustomStatements

A list of resource policy statements can be given for an Api

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### IpRangeBlacklist

Resource Policy statements will be generated and attached to the Api for blacklisting the given list of Ip addresses or ranges

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### IpRangeWhitelist

Resource Policy statements will be generated and attached to the Api for whitelisting the given list of Ip addresses or ranges

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### SourceVpcBlacklist

Resource Policy statements will be generated and attached to the Api for blacklisting the given list of Source Vpcs or Vpc endpoints

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### SourceVpcWhitelist

Resource Policy statements will be generated and attached to the Api for whitelisting the given list of Source Vpcs or Vpc endpoints

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### SourceVpcBlacklist

Blacklisting source VPC or VPC endpoint

#### YAML

```
Auth:
  ResourcePolicy:
    CustomStatements: [{
      "Effect": "Allow",
      "Principal": "*",
      "Action": "execute-api:Invoke",
      "Resource": "execute-api:/Prod/PUT/get",
      "Condition": {
        "IpAddress": {
          "aws:SourceIp": "1.2.3.4"
        }
      }
    }]
  IpRangeBlacklist: ['10.20.30.40', '1.2.3.4']
  SourceVpcBlacklist: ["vpce-1a2b3c4d"]
  AwsAccountWhitelist: ['123456789101']
```

## RequestModel

Configure Request Model for a specific Api+Path+Method.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Model: String
Required: Boolean
```

## Properties

### Model

Name of a model defined in the Models property of the [AWS::Serverless::Api](#) (p. 28).

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Required

adds a required property in the parameters section of Open Api definition for given Api endpoint

*Type:* Boolean

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### Request Model

Request Model Example

#### YAML

```
RequestModel:
  Model: User
  Required: true
```

### RequestParameter

Configure Request Parameter for a specific Api+Path+Method.

Either *Required* or *Caching* property needs to be specified for request parameter

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Caching: Boolean
Required: Boolean
```

## Properties

### Caching

Adds `cacheKeyParameters` section to the API Gateway OpenApi definition

*Type:* Boolean

*Required:* Conditional

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Required

This field specifies whether a parameter is required

*Type:* Boolean

*Required:* Conditional

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### Request Parameter

Example of setting Request Parameters

#### YAML

```
RequestParameters:
  - method.request.header.Authorization:
      Required: true
      Caching: true
```

## CloudWatchEvent

The object describing an event source with type CloudWatchEvent

SAM generates [AWS::Events::Rule](#) resource when this event type is set

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
EventBusName: String
Input: String
InputPath: String
Pattern: EventPattern
```

## Properties

### EventBusName

The event bus to associate with this rule. If you omit this, the default event bus is used.

*Type:* String

*Required:* No

*Default:* default

*CloudFormation Compatibility:* This property is passed directly to the [EventBusName](#) property of an `AWS::Events::Rule`.

### Input

Valid JSON text passed to the target. If you use this property, nothing from the event text itself is passed to the target.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Input](#) property of an `AWS::Events::Rule` Target.

### InputPath

When you don't want to pass the entire matched event, `InputPath` describes which part of the event to pass to the target.

Type: String

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [InputPath](#) property of an `AWS::Events::Rule` Target.

#### Pattern

Describes which events are routed to the specified target. For more information, see [Events and Event Patterns in EventBridge](#) in the Amazon EventBridge User Guide

Type: [EventPattern](#)

Required: Yes

*CloudFormation Compatibility:* This property is passed directly to the [EventPattern](#) property of an `AWS::Events::Rule`.

## Examples

### CloudWatchEvent

CloudWatch Event Example

#### YAML

```
CWEvent:
  Type: CloudWatchEvent
  Properties:
    Input: '{"Key": "Value"}'
    Pattern:
      detail:
        state:
          - terminated
```

### CloudWatchLogs

This event generates a [AWS::Logs::SubscriptionFilter](#) resource and specifies a subscription filter and associates it with the specified log group.

#### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
FilterPattern: String
LogGroupName: String
```

## Properties

#### FilterPattern

The filtering expressions that restrict what gets delivered to the destination AWS resource. For more information about the filter pattern syntax, see [Filter and Pattern Syntax](#).

Type: String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [FilterPattern](#) property of an `AWS::Logs::SubscriptionFilter`.

`LogGroupName`

The log group to associate with the subscription filter. All log events that are uploaded to this log group are filtered and delivered to the specified AWS resource if the filter pattern matches the log events.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [LogGroupName](#) property of an `AWS::Logs::SubscriptionFilter`.

## Examples

### Cloudwatchlogs Subscription filter

Cloudwatchlogs Subscription filter Example

#### YAML

```
CWLog:
  Type: CloudWatchLogs
  Properties:
    LogGroupName:
      Ref: CloudWatchLambdaLogsGroup
    FilterPattern: My pattern
```

## Cognito

The object describing an event source with type Cognito.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Trigger: List
UserPool: String
```

## Properties

### Trigger

The Lambda trigger configuration information for the new user pool.

*Type:* List

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [LambdaConfig](#) property of an `AWS::Cognito::UserPool`.



## UserPool

Reference to UserPool defined in the same template

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### Cognito Event

#### Cognito Event Example

#### YAML

```
CognitoUserPoolPreSignup:
  Type: Cognito
  Properties:
    UserPool:
      Ref: MyCognitoUserPool
    Trigger: PreSignUp
```

## DynamoDB

DynamoDB streaming event source type.

SAM generates [AWS::Lambda::EventSourceMapping](#) resource when this event type is set

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
BatchSize: Integer
BisectBatchOnFunctionError: Boolean
DestinationConfig: DestinationConfig
Enabled: Boolean
MaximumBatchingWindowInSeconds: Integer
MaximumRecordAgeInSeconds: Integer
MaximumRetryAttempts: Integer
ParallelizationFactor: Integer
StartingPosition: String
Stream: String
```

## Properties

### BatchSize

The maximum number of items to retrieve in a single batch.

*Type:* Integer

*Required:* No

*Default:* 100

*CloudFormation Compatibility:* This property is passed directly to the [BatchSize](#) property of an `AWS::Lambda::EventSourceMapping`.

*Minimum:* 1

*Maximum:* 1000

#### `BisectBatchOnFunctionError`

If the function returns an error, split the batch in two and retry.

*Type:* Boolean

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [BisectBatchOnFunctionError](#) property of an `AWS::Lambda::EventSourceMapping`.

#### `DestinationConfig`

An Amazon SQS queue or Amazon SNS topic destination for discarded records.

*Type:* [DestinationConfig](#)

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [DestinationConfig](#) property of an `AWS::Lambda::EventSourceMapping`.

#### `Enabled`

Disables the event source mapping to pause polling and invocation.

*Type:* Boolean

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Enabled](#) property of an `AWS::Lambda::EventSourceMapping`.

#### `MaximumBatchingWindowInSeconds`

The maximum amount of time to gather records before invoking the function, in seconds.

*Type:* Integer

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [MaximumBatchingWindowInSeconds](#) property of an `AWS::Lambda::EventSourceMapping`.

#### `MaximumRecordAgeInSeconds`

The maximum age of a record that Lambda sends to a function for processing.

*Type:* Integer

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [MaximumRecordAgeInSeconds](#) property of an `AWS::Lambda::EventSourceMapping`.

#### `MaximumRetryAttempts`

The maximum number of times to retry when the function returns an error.

*Type:* Integer

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [MaximumRetryAttempts](#) property of an `AWS::Lambda::EventSourceMapping`.

`ParallelizationFactor`

The number of batches to process from each shard concurrently.

*Type:* Integer

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [ParallelizationFactor](#) property of an `AWS::Lambda::EventSourceMapping`.

`StartingPosition`

The position in a stream from which to start reading.

Supported values: `TRIM_HORIZON`, `LATEST`.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [StartingPosition](#) property of an `AWS::Lambda::EventSourceMapping`.

`Stream`

ARN of the DynamoDB stream.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [EventSourceArn](#) property of an `AWS::Lambda::EventSourceMapping`.

## Examples

### DynamoDB Event

DynamoDB Event

### YAML

```
Properties:
  Stream: arn:aws:dynamodb:us-east-1:123456789012:table/TestTable/
stream/2016-08-11T21:21:33.291
  StartingPosition: TRIM_HORIZON
  BatchSize: 10
  Enabled: false
```

## EventBridgeRule

The object describing an event source with type `EventBridgeRule`

SAM generates `AWS::Events::Rule` resource when this event type is set

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
Input: String
InputPath: String
Pattern: EventPattern
```

## Properties

### Input

Valid JSON text passed to the target. If you use this property, nothing from the event text itself is passed to the target.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Input](#) property of an `AWS::Events::Rule` Target.

### InputPath

When you don't want to pass the entire matched event, `InputPath` describes which part of the event to pass to the target.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [InputPath](#) property of an `AWS::Events::Rule` Target.

### Pattern

Pattern describing which EventBridge events trigger the function. Only matching events trigger the function. For more information, see [Events and Event Patterns in EventBridge](#) in the Amazon EventBridge User Guide

*Type:* [EventPattern](#)

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [EventPattern](#) property of an `AWS::Events::Rule`.

## Examples

### EventBridgeRule

EventBridgeRule Event Example

### YAML

```
EBRule:
  Type: EventBridgeRule
```

```
Properties:
  Input: '{"Key": "Value"}'
  Pattern:
    detail:
      state:
        - terminated
```

## HttpApi

HTTP APIs are in beta for Amazon API Gateway and are subject to change.

The object describing an event source with type `HttpApi`.

If an OpenApi definition for the specified path and method exists on the API, SAM will add the Lambda integration and security section (if applicable) for you.

If no OpenApi definition for the specified path and method exists on the API, SAM will create this definition for you.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
ApiId: String
Auth: HttpApiFunctionAuth (p. 90)
Method: String
Path: String
```

## Properties

### ApiId

Identifier of an [AWS::Serverless::HttpApi \(p. 102\)](#) resource defined in this template.

If not defined, a default [AWS::Serverless::HttpApi \(p. 102\)](#) resource is created called `ServerlessHttpApi` using a generated OpenApi document containing a union of all paths and methods defined by Api events defined in this template that do not specify an `ApiId`.

This cannot reference an [AWS::Serverless::HttpApi \(p. 102\)](#) resource defined in another template.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Auth

Auth configuration for this specific Api+Path+Method.

Useful for overriding the API's `DefaultAuthorizer` or setting auth config on an individual path when no `DefaultAuthorizer` is specified.

*Type:* [HttpApiFunctionAuth \(p. 90\)](#)

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### Method

HTTP method for which this function is invoked.

If no `Path` and `Method` are specified, SAM will create a default API path that routes any request that doesn't map to a different endpoint to this Lambda function. Only one of these default paths can exist per API.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### Path

Uri path for which this function is invoked. Must start with `/`.

If no `Path` and `Method` are specified, SAM will create a default API path that routes any request that doesn't map to a different endpoint to this Lambda function. Only one of these default paths can exist per API.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### Default HttpApi Event

HttpApi Event that uses the default path. All unmapped paths and methods on this API will route to this endpoint.

#### YAML

```
Events:
  HttpApiEvent:
    Type: HttpApi
```

### HttpApi

HttpApi Event that uses a specific path and method.

#### YAML

```
Events:
  HttpApiEvent:
    Type: HttpApi
    Properties:
      Path: /
      Method: GET
```

## HttpApi Authorization

HttpApi Event that uses an Authorizer.

### YAML

```
Events:
  HttpApiEvent:
    Type: HttpApi
    Properties:
      Path: /authenticated
      Method: GET
      Auth:
        Authorizer: OpenIdAuth
        AuthorizationScopes:
          - scope1
          - scope2
```

## HttpApiFunctionAuth

HTTP APIs are in beta for Amazon API Gateway and are subject to change.

Configures authorization at the event level.

Configure Auth for a specific Api + Path + Method

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
AuthorizationScopes: List
Authorizer: String
```

## Properties

### AuthorizationScopes

Authorization scopes to apply to this Api + Path + Method.

Scopes listed here will override any scopes applied by the `DefaultAuthorizer` if one exists.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Authorizer

The `Authorizer` for a specific Function

If you have specified a Global Authorizer on the API and want to make a specific Function public, override by setting `Authorizer` to `NONE`.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### Function-Auth

Specifying Authorization at Function level

#### YAML

```
Auth:
  Authorizer: OpenIdAuth
  AuthorizationScopes:
    - scope1
    - scope2
```

## IoTRule

Creates an [AWS::IoT::TopicRule](#) resource to declare an AWS IoT rule. For more information see [AWS CloudFormation documentation](#)

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
AwsIotSqlVersion: String
Sql: String
```

## Properties

### AwsIotSqlVersion

The version of the SQL rules engine to use when evaluating the rule.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [AwsIotSqlVersion](#) property of an `AWS::IoT::TopicRule` TopicRulePayload.

### Sql

The SQL statement used to query the topic. For more information, see [AWS IoT SQL Reference](#) in the AWS IoT Developer Guide.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [Sql](#) property of an `AWS::IoT::TopicRule` TopicRulePayload.



## Examples

### IOT Rule

#### IOT Rule Example

#### YAML

```
IoTRule:
  Type: IoTRule
  Properties:
    Sql: SELECT * FROM 'topic/test'
```

## Kinesis

Kinesis event source configuration.

SAM generates [AWS::Lambda::EventSourceMapping](#) resource when this event type is set

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
BatchSize: Integer
BisectBatchOnFunctionError: Boolean
DestinationConfig: DestinationConfig
Enabled: Boolean
MaximumBatchingWindowInSeconds: Integer
MaximumRecordAgeInSeconds: Integer
MaximumRetryAttempts: Integer
ParallelizationFactor: Integer
StartingPosition: String
Stream: String
```

## Properties

### BatchSize

The maximum number of items to retrieve in a single batch.

*Type:* Integer

*Required:* No

*Default:* 100

*CloudFormation Compatibility:* This property is passed directly to the [BatchSize](#) property of an `AWS::Lambda::EventSourceMapping`.

*Minimum:* 1

*Maximum:* 10000

### BisectBatchOnFunctionError

If the function returns an error, split the batch in two and retry.

Type: Boolean

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [BisectBatchOnFunctionError](#) property of an `AWS::Lambda::EventSourceMapping`.

`DestinationConfig`

An Amazon SQS queue or Amazon SNS topic destination for discarded records.

Type: [DestinationConfig](#)

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [DestinationConfig](#) property of an `AWS::Lambda::EventSourceMapping`.

`Enabled`

Disables the event source mapping to pause polling and invocation.

Type: Boolean

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [Enabled](#) property of an `AWS::Lambda::EventSourceMapping`.

`MaximumBatchingWindowInSeconds`

The maximum amount of time to gather records before invoking the function, in seconds.

Type: Integer

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [MaximumBatchingWindowInSeconds](#) property of an `AWS::Lambda::EventSourceMapping`.

`MaximumRecordAgeInSeconds`

The maximum age of a record that Lambda sends to a function for processing.

Type: Integer

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [MaximumRecordAgeInSeconds](#) property of an `AWS::Lambda::EventSourceMapping`.

`MaximumRetryAttempts`

The maximum number of times to retry when the function returns an error.

Type: Integer

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [MaximumRetryAttempts](#) property of an `AWS::Lambda::EventSourceMapping`.

`ParallelizationFactor`

The number of batches to process from each shard concurrently.

*Type:* Integer

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [ParallelizationFactor](#) property of an `AWS::Lambda::EventSourceMapping`.

**StartingPosition**

The position in a stream from which to start reading.

Supported values: TRIM\_HORIZON, LATEST, AT\_TIMESTAMP.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [StartingPosition](#) property of an `AWS::Lambda::EventSourceMapping`.

**Stream**

The ARN of the data stream or a stream consumer.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [EventSourceArn](#) property of an `AWS::Lambda::EventSourceMapping`.

## Examples

### Kinesis Event Source

Kinesis Event Source

#### YAML

```
Properties:
  Stream: arn:aws:kinesis:us-east-1:123456789012:stream/my-stream
  StartingPosition: TRIM_HORIZON
  BatchSize: 10
  Enabled: false
```

## S3

S3 event source type.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Bucket: String
Events: String | List
Filter: S3NotificationFilter
```

## Properties

### Bucket

S3 bucket name. This bucket must exist in the same template.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is similar to the [BucketName](#) property of an `AWS::S3::Bucket`. This is a required field in SAM. This field only accepts a reference to the S3 bucket created in this template

### Events

The Amazon S3 bucket event for which to invoke the AWS Lambda function. See [Amazon S3 supported event types](#) for valid values.

*Type:* String | List

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [Event](#) property of the `AWS::S3::Bucket LambdaConfiguration` data type.

### Filter

The filtering rules that determine which objects invoke the AWS Lambda function.

*Type:* [S3NotificationFilter](#)

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [NotificationFilter](#) property of the `AWS::S3::Bucket LambdaConfiguration` data type.

## Examples

### S3-Event

Example of an S3 Event

#### YAML

```
S3Event:
  Type: S3
  Properties:
    Bucket:
      Ref: ImagesBucket
    Events: s3:ObjectCreated:*
    Filter:
      S3Key:
        Rules:
          - Name: name
            Value: value
ImagesBucket:
  Type: AWS::S3::Bucket
```

## Schedule

This object describes an event source with type Schedule

SAM generates [AWS::Events::Rule](#) resource when this event type is set

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

## YAML

```
Description: String
Enabled: Boolean
Input: String
Name: String
Schedule: String
```

## Properties

### Description

The description of the rule.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Description](#) property of an `AWS::Events::Rule`.

### Enabled

Indicates whether the rule is enabled.

If the property is set to False, the rule is disabled

*Type:* Boolean

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [State](#) property of an `AWS::Events::Rule`.

### Input

Valid JSON text passed to the target. If you use this property, nothing from the event text itself is passed to the target.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Target](#) property of an `AWS::Events::Rule` `Target`.

### Name

The name of the rule. If you don't specify a name, AWS CloudFormation generates a unique physical ID and uses that ID for the rule name.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Name](#) property of an `AWS::Events::Rule`.

#### Schedule

The scheduling expression that determines when and how often the rule runs. For more information, see [Schedule Expressions for Rules](#).

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [ScheduleExpression](#) property of an `AWS::Events::Rule`.

## Examples

### CloudWatch Schedule Event

#### CloudWatch Schedule Event Example

#### YAML

```
CWSchedule:
  Type: Schedule
  Properties:
    Schedule: 'rate(1 minute)'
    Name: TestSchedule
    Description: test schedule
    Enabled: False
```

## SNS

SNS event source configuration.

SAM generates `AWS::SNS::Subscription` resource when this event type is set

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
FilterPolicy: SnsFilterPolicy
Region: String
SqsSubscription: Boolean | SqsSubscriptionObject (p. 98)
Topic: String
```

## Properties

### FilterPolicy

The filter policy JSON assigned to the subscription. For more information, see [GetSubscriptionAttributes](#) in the Amazon Simple Notification Service API Reference.

*Type:* [SnsFilterPolicy](#)

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [FilterPolicy](#) property of an `AWS::SNS::Subscription`.

#### Region

For cross-region subscriptions, the region in which the topic resides.

If no region is specified, CloudFormation uses the region of the caller as the default.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Region](#) property of an `AWS::SNS::Subscription`.

#### SqsSubscription

Set this property to true, or specify `SqsSubscriptionObject` to enable batching SNS topic notifications in an SQS queue. Setting this property to true creates a new SQS queue, whereas specifying a `SqsSubscriptionObject` uses an existing SQS queue.

*Type:* Boolean | [SqsSubscriptionObject](#) (p. 98)

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### Topic

The ARN of the topic to subscribe to.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [TopicArn](#) property of an `AWS::SNS::Topic`.

## Examples

### SNS Event Source Example

SNS Event Source Example

#### YAML

```
Properties:
  Topic: arn:aws:sns:us-east-1:123456789012:my_topic
  SqsSubscription: True
  FilterPolicy:
    store:
      - example_corp
    price_usd:
      - numeric:
          - ">="
          - 100
```

### SqsSubscriptionObject

Specify an existing SQS queue option to SNS event

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

## YAML

```
BatchSize: String
Enabled: Boolean
QueueArn: String
QueuePolicyLogicalId: String
QueueUrl: String
```

## Properties

### BatchSize

The maximum number of items to retrieve in a single batch for the SQS queue.

*Type:* String

*Required:* No

*Default:* 10

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### Enabled

Disables the SQS event source mapping to pause polling and invocation.

*Type:* Boolean

*Required:* No

*Default:* True

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### QueueArn

Specify an existing SQS queue arn.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### QueuePolicyLogicalId

Give a custom logicalId name for the [AWS::SQS::QueuePolicy](#) resource.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.



### QueueUrl

Specify the queue URL associated with the `QueueArn` property.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### Existing SQS for SNS event

Example to add existing SQS queue for subscribing to an SNS topic.

#### YAML

```
QueuePolicyLogicalId: CustomQueuePolicyLogicalId
QueueArn:
  Fn::GetAtt: MyCustomQueue.Arn
QueueUrl:
  Ref: MyCustomQueue
BatchSize: 5
```

## SQS

SQS event source type.

SAM generates [AWS::Lambda::EventSourceMapping](#) resource when this event type is set

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
BatchSize: Integer
Enabled: Boolean
Queue: String
```

## Properties

### BatchSize

The maximum number of items to retrieve in a single batch.

*Type:* Integer

*Required:* No

*Default:* 10

*CloudFormation Compatibility:* This property is passed directly to the `BatchSize` property of an `AWS::Lambda::EventSourceMapping`.

*Minimum:* 1

*Maximum:* 10

**Enabled**

Disables the event source mapping to pause polling and invocation.

*Type:* Boolean

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Enabled](#) property of an `AWS::Lambda::EventSourceMapping`.

**Queue**

The ARN of the queue.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [EventSourceArn](#) property of an `AWS::Lambda::EventSourceMapping`.

## Examples

### SQS Event

SQS Event

### YAML

```
Type: SQS
Properties:
  Queue: arn:aws:sqs:us-west-2:012345678901:my-queue
  BatchSize: 10
  Enabled: false
```

## FunctionCode

The [deployment package](#) for a Lambda function.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
Bucket: String
Key: String
Version: String
```

## Properties

**Bucket**

An Amazon S3 bucket in the same AWS Region as your function.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [S3Bucket](#) property of the `AWS::Lambda::Function` Code data type.

#### Key

The Amazon S3 key of the deployment package.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [S3Key](#) property of the `AWS::Lambda::Function` Code data type.

#### Version

For versioned objects, the version of the deployment package object to use.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [S3ObjectVersion](#) property of the `AWS::Lambda::Function` Code data type.

## Examples

### FunctionCode

Function Code example

#### YAML

```
FunctionCode:
  Bucket: mybucket-name
  Key: mykey-name
  Version: 121212
```

## AWS::Serverless::HttpApi

HTTP APIs are in beta for Amazon API Gateway and are subject to change.

Creates an API Gateway HTTP API, which enables you to create RESTful APIs with lower latency and lower costs than REST APIs. For more information about HTTP APIs see [HTTP API](#) in the API Gateway Developer Guide.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Type: AWS::Serverless::HttpApi
```

```
Properties:
  Auth: HttpApiAuth (p. 105)
  DefinitionBody: String
  DefinitionUri: String | HttpApiDefinition (p. 110)
  StageName: String
```

## Properties

### Auth

Configure authorization to control access to your API Gateway API.

For more information about configuring access see [JWT Authorizers](#) in the API Gateway Developer Guide.

Type: [HttpApiAuth](#) (p. 105)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### DefinitionBody

OpenAPI specification that describes your API. If neither `DefinitionUri` nor `DefinitionBody` are specified, SAM will generate a `DefinitionBody` for you based on your template configuration.

Type: String

Required: No

*CloudFormation Compatibility:* This property is similar to the `Body` property of an `AWS::ApiGatewayV2::Api`. If certain properties are provided, content may be inserted or modified into the `DefinitionBody` before being passed to CloudFormation. Properties include `Auth` and an `EventSource` of type `HttpApi` on for a corresponding `AWS::Serverless::Function`.

### DefinitionUri

AWS S3 Uri, local file path, or location object of the the OpenAPI document defining the API. The AWS S3 object this property references must be a valid OpenAPI file. If neither `DefinitionUri` nor `DefinitionBody` are specified, SAM will generate a `DefinitionBody` for you based on your template configuration.

If a local file path is provided, the template must go through the workflow that includes the `sam deploy` or `sam package` command, in order for the definition to be transformed properly.

Intrinsic functions are not supported in external OpenApi files referenced by `DefinitionUri`. Use instead the `DefinitionBody` property with the [Include Transform](#) to import an OpenApi definition into the template.

Type: String | [HttpApiDefinition](#) (p. 110)

Required: No

*CloudFormation Compatibility:* This property is similar to the `BodyS3Location` property of an `AWS::ApiGatewayV2::Api`. The nested Amazon S3 properties are named differently.

### StageName

The name of the API stage. If a name is not given, SAM will use the `$default` stage from Api Gateway.

Type: String

Required: No

Default: \$default

*CloudFormation Compatibility:* This property is passed directly to the [StageName](#) property of an `AWS::ApiGatewayV2::Stage`.

## Return Values

### Ref

When you pass the logical ID of this resource to the intrinsic Ref function, Ref returns the API ID of the underlying `AWS::ApiGatewayV2::Api` resource, such as `a1bcdef2gh`.

For more information about using the Ref function, see [Ref](#).

## Examples

### Simple Http Api

Bare minimum needed to set up an HttpApi endpoint backed by a Lambda function. This uses the default HTTP API that SAM creates.

#### YAML

```
AWSTemplateFormatVersion: '2010-09-09'
Description: AWS SAM template with a simple API definition
Resources:
  ApiFunction:
    Type: AWS::Serverless::Function
    Properties:
      Events:
        ApiEvent:
          Type: HttpApi
          Handler: index.handler
          InlineCode: |
            def handler(event, context):
                return {'body': 'Hello World!', 'statusCode': 200}
          Runtime: python3.7
      Transform: AWS::Serverless-2016-10-31
```

### Http Api with Auth

Example of how to set up authorization on API endpoints.

#### YAML

```
Properties:
  Auth:
    DefaultAuthorizer: OAuth2
    Authorizers:
      OAuth2:
        AuthorizationScopes:
          - scope4
        JwtConfiguration:
          issuer: "https://www.example.com/v1/connect/oauth2"
```

```
    audience:
      - MyApi
    IdentitySource: "$request.querystring.param"
  OpenIdAuth:
    AuthorizationScopes:
      - scope1
      - scope2
    OpenIdConnectUrl: "https://www.example.com/v1/connect/oidc/.well-known/openid-configuration"
  JwtConfiguration:
    issuer: "https://www.example.com/v1/connect/oidc"
    audience:
      - MyApi
    IdentitySource: "$request.querystring.param"
```

## Http Api with OpenApi Document

Shows how to add OpenApi to the document.

Note that SAM will fill in any missing lambda integrations for HttpApi events that reference this API. SAM will also add any missing paths that HttpApi events reference.

### YAML

```
Properties:
  DefinitionBody:
    info:
      version: '1.0'
      title:
        Ref: AWS::StackName
    paths:
      "/":
        get:
          security:
            - OpenIdAuth:
                - scope1
                - scope2
          responses: {}
    openapi: 3.0.1
    securitySchemes:
      OpenIdAuth:
        type: openIdConnect
        x-amazon-apigateway-authorizer:
          identitySource: "$request.querystring.param"
          type: jwt
          jwtConfiguration:
            audience:
              - MyApi
            issuer: https://www.example.com/v1/connect/oidc
            openIdConnectUrl: https://www.example.com/v1/connect/oidc/.well-known/openid-configuration
```

## HttpApiAuth

HTTP APIs are in beta for Amazon API Gateway and are subject to change.

Configure authorization to control access to your API Gateway API.

For more information about configuring access see [JWT Authorizers](#) in the API Gateway Developer Guide.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
Authorizers: OpenIdAuthorizer \(p. 108\) | OAuth2Authorizer \(p. 107\)
DefaultAuthorizer: String
```

## Properties

### Authorizers

The authorizer used to control access to your API Gateway API. OAuth 2.0 and OpenIdConnect are currently supported.

Type: [OpenIdAuthorizer \(p. 108\)](#) | [OAuth2Authorizer \(p. 107\)](#)

Required: No

Default: None

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

*Additional Notes:* SAM adds the Authorizers to the OpenApi definition of an Api.

### DefaultAuthorizer

Specify a default authorizer for an API Gateway API, which will be used for authorizing API calls by default.

Type: String

Required: No

Default: None

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### OpenId Auth

#### OpenId Auth Example

### YAML

```
Auth:
  Authorizers:
    OAuth2Authorizer:
      AuthorizationScopes:
        - scope1
        - scope2
      JwtConfiguration:
        issuer: "https://www.example.com/v1/connect/oauth2"
        audience:
          - MyApi
```

```
    IdentitySource: "$request.querystring.param"
  OpenIdAuthorizer:
    AuthorizationScopes:
      - scope1
      - scope2
    OpenIdConnectUrl: "https://www.example.com/v1/connect/oidc/.well-known/openid-configuration"
    JwtConfiguration:
      issuer: "https://www.example.com/v1/connect/oidc"
      audience:
        - MyApi
    IdentitySource: "$request.querystring.param"
  DefaultAuthorizer: OpenIdAuthorizer
```

## OAuth2Authorizer

HTTP APIs are in beta for Amazon API Gateway and are subject to change.

Definition for an OAuth 2.0 authorizer.

For more information, see the [API Gateway documentation](#).

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
AuthorizationScopes: List
IdentitySource: String
JwtConfiguration: Map
```

## Properties

### AuthorizationScopes

List of authorization scopes for this authorizer.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### IdentitySource

Identity source expression for this authorizer.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### JwtConfiguration

JSON Web Token (JWT) configuration for this authorizer.



This is passed through to the `jwtConfiguration` section of a `x-amazon-apigateway-authorizer` in the `securitySchemes` section of an OpenApi document.

*Type:* Map

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### OAuth2 Authorizer

#### OAuth2 Authorizer Example

#### YAML

```
Auth:
  Authorizers:
    OAuth2Authorizer:
      AuthorizationScopes:
        - scope1
      JwtConfiguration:
        issuer: "https://www.example.com/v1/connect/oauth2"
        audience:
          - MyApi
      IdentitySource: "$request.querystring.param"
  DefaultAuthorizer: OAuth2Authorizer
```

### OpenIdAuthorizer

HTTP APIs are in beta for Amazon API Gateway and are subject to change.

Definition for an OpenIdConnect (OIDC) authorizer, which is built on top of the OAuth 2.0 protocol.

For more information about configuring access see [JWT Authorizers](#) in the API Gateway Developer Guide.

#### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
AuthorizationScopes: List
IdentitySource: String
JwtConfiguration: Map
OpenIdConnectUrl: String
```

## Properties

### AuthorizationScopes

List of authorization scopes for this authorizer.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### IdentitySource

Identity source expression for this authorizer.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### JwtConfiguration

JSON Web Token (JWT) configuration for this authorizer.

This is passed through to the `jwtConfiguration` section of a `x-amazon-apigateway-authorizer` in the `securitySchemes` section of an OpenApi document.

*Type:* Map

*Required:* No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

#### OpenIdConnectUrl

URL to use to find the token issuer. If the lookup fails using this URL, the authorizer will fall back to using the issuer in `JwtConfiguration`.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

## Examples

### OpenId Authorizer

#### OpenId Authorizer Example

#### YAML

```
Auth:
  Authorizers:
    OpenIdAuthorizer:
      AuthorizationScopes:
        - scope1
        - scope2
      OpenIdConnectUrl: "https://www.example.com/v1/connect/oidc/.well-known/openid-configuration"
      JwtConfiguration:
```

```
issuer: "https://www.example.com/v1/connect/oidc"
audience:
  - MyApi
IdentitySource: "$request.querystring.param"
```

## HttpApiDefinition

HTTP APIs are in beta for Amazon API Gateway and are subject to change.

An OpenAPI document defining the API.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Bucket: String
Key: String
Version: String
```

### Properties

#### Bucket

The name of the Amazon S3 bucket where the OpenAPI file is stored.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [Bucket](#) property of the `AWS::ApiGatewayV2::Api BodyS3Location` data type.

#### Key

The Amazon S3 key of the OpenAPI file.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [Key](#) property of the `AWS::ApiGatewayV2::Api BodyS3Location` data type.

#### Version

For versioned objects, the version of the OpenAPI file.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [Version](#) property of the `AWS::ApiGatewayV2::Api BodyS3Location` data type.

## Examples

### Definition Uri example

API Definition example

### YAML

```
DefinitionUri:
  Bucket: mybucket-name
  Key: mykey-name
  Version: 121212
```

## AWS::Serverless::LayerVersion

Creates a Lambda LayerVersion that contains library or runtime code needed by a Lambda Function.

**Important Note:** Since the release of the [UpdateReplacePolicy](#) resource attribute in AWS CloudFormation, [AWS::Lambda::LayerVersion](#) (recommended) offers the same benefits as [AWS::Serverless::LayerVersion](#) (p. 111).

When a Serverless LayerVersion is transformed, SAM also transforms the logical id of the resource so that old LayerVersions are not automatically deleted by CloudFormation when the resource is updated.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

### YAML

```
Type: AWS::Serverless::LayerVersion
Properties:
  CompatibleRuntimes: List
  ContentUri: String | LayerContent (p. 113)
  Description: String
  LayerName: String
  LicenseInfo: String
  RetentionPolicy: String
```

## Properties

### CompatibleRuntimes

List of runtimes compatible with this LayerVersion.

*Type:* List

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [CompatibleRuntimes](#) property of an `AWS::Lambda::LayerVersion`.

### ContentUri

AWS S3 Uri, local file path, or LayerContent object of the layer code.

If an AWS S3 Uri or LayerContent object is provided, The AWS S3 object referenced must be a valid ZIP archive that contains the contents of an [AWS Lambda layer](#).

If a local file path is provided, the template must go through the workflow that includes the `sam deploy` or `sam package` command, in order for the content to be transformed properly.

**Note:** Either CodeUri or InlineCode is required.

Type: String | [LayerContent](#) (p. 113)

Required: Yes

*CloudFormation Compatibility:* This property is similar to the [Content](#) property of an `AWS::Serverless::LayerVersion`. The nested Amazon S3 properties are named differently.

#### Description

Description of this layer.

Type: String

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [Description](#) property of an `AWS::Lambda::LayerVersion`.

#### LayerName

The name or Amazon Resource Name (ARN) of the layer.

Type: String

Required: No

Default: Resource logical id

*CloudFormation Compatibility:* This property is similar to the [LayerName](#) property of an `AWS::Lambda::LayerVersion`. If you don't specify a name, the logical id of the resource will be used as the name.

#### LicenseInfo

Information about the license for this LayerVersion.

Type: String

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [LicenseInfo](#) property of an `AWS::Lambda::LayerVersion`.

#### RetentionPolicy

Specifies whether old versions of your LayerVersion are retained or deleted after an update.

Supported values: Retain and Delete.

Type: String

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

*Additional Notes:* When you specify `Retain`, AWS SAM adds a [Resource Attribute](#) of `DeletionPolicy: Retain` to the transformed `AWS::Lambda::LayerVersion` resource.

## Return Values

### Ref

When the logical ID of this resource is provided to the `Ref` intrinsic function, it returns the resource ARN of the underlying `Lambda LayerVersion`.

For more information about using the `Ref` function, see [Ref](#).

## Examples

### LayerVersionExample

Example of a `LayerVersion`

#### YAML

```
Properties:
  LayerName: MyLayer
  Description: Layer description
  ContentUri: 's3://my-bucket/my-layer.zip'
  CompatibleRuntimes:
    - nodejs6.10
    - nodejs8.10
  LicenseInfo: 'Available under the MIT-0 license.'
  RetentionPolicy: Retain
```

## LayerContent

A ZIP archive that contains the contents of an [AWS Lambda layer](#).

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Bucket: String
Key: String
Version: String
```

## Properties

### Bucket

The Amazon S3 bucket of the layer archive.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [S3Bucket](#) property of the `AWS::Lambda::LayerVersion` Content data type.

#### Key

The Amazon S3 key of the layer archive.

*Type:* String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [S3Key](#) property of the `AWS::Lambda::LayerVersion` Content data type.

#### Version

For versioned objects, the version of the layer archive object to use.

*Type:* String

*Required:* No

*CloudFormation Compatibility:* This property is passed directly to the [S3ObjectVersion](#) property of the `AWS::Lambda::LayerVersion` Content data type.

## Examples

### LayerContent

Layer Content example

#### YAML

```
LayerContent:
  Bucket: mybucket-name
  Key: mykey-name
  Version: 121212
```

## AWS::Serverless::SimpleTable

Creates a DynamoDB table with a single attribute primary key. It is useful when data only needs to be accessed via a primary key.

To use the more advanced functionality of DynamoDB, use an [AWS::DynamoDB::Table](#) resource instead.

## Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Type: AWS::Serverless::SimpleTable
Properties:
  PrimaryKey: PrimaryKeyObject (p. 116)
  ProvisionedThroughput: ProvisionedThroughput
  SSESpecification: SSESpecification
  TableName: String
  Tags: Map
```

## Properties

### PrimaryKey

Attribute name and type to be used as the table's primary key. If not provided, the primary key will be a `String` with a value of `id`.

**Note:** The value of this property cannot be modified after this resource is created.

Type: [PrimaryKeyObject](#) (p. 116)

Required: No

*CloudFormation Compatibility:* This property is unique to AWS SAM and does not have an AWS CloudFormation equivalent.

### ProvisionedThroughput

Read and write throughput provisioning information.

If `ProvisionedThroughput` is not specified `BillingMode` will be specified as `PAY_PER_REQUEST`.

Type: [ProvisionedThroughput](#)

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [ProvisionedThroughput](#) property of an `AWS::DynamoDB::Table`.

### SSESpecification

Specifies the settings to enable server-side encryption.

Type: [SSESpecification](#)

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [SSESpecification](#) property of an `AWS::DynamoDB::Table`.

### TableName

Name for the DynamoDB Table.

Type: `String`

Required: No

*CloudFormation Compatibility:* This property is passed directly to the [TableName](#) property of an `AWS::DynamoDB::Table`.

### Tags

A map (string to string) that specifies the tags to be added to this `SimpleTable`. Keys and values are limited to alphanumeric characters. Keys can be 1 to 127 Unicode characters in length and cannot be prefixed with `aws:`. Values can be 1 to 255 Unicode characters in length.

Type: `Map`

Required: No



*CloudFormation Compatibility:* This property is similar to the [Tags](#) property of an `AWS::DynamoDB::Table`. The Tags property in SAM consists of Key:Value pairs; in CloudFormation it consists of a list of Tag objects.

## Return Values

### Ref

When the logical ID of this resource is provided to the Ref intrinsic function, it returns the resource name of the underlying DynamoDB table.

For more information about using the Ref function, see [Ref](#).

## Examples

### SimpleTableExample

Example of a SimpleTable

#### YAML

```
Properties:
  TableName: my-table
  PrimaryKey:
    Name: id
    Type: String
  ProvisionedThroughput:
    ReadCapacityUnits: 5
    WriteCapacityUnits: 5
  Tags:
    Department: Engineering
    AppType: Serverless
  SSESpecification:
    SSEEnabled: true
```

## PrimaryKeyObject

The object describing the properties of a primary key.

### Syntax

To declare this entity in your AWS SAM template, use the following syntax:

#### YAML

```
Name: String
Type: String
```

### Properties

#### Name

Attribute name of the primary key.

Type: String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [AttributeName](#) property of the `AWS::DynamoDB::Table AttributeDefinition` data type.

*Additional Notes:* This property is also passed to the [AttributeName](#) property of an `AWS::DynamoDB::Table KeySchema` data type.

Type

The data type for the primary key.

Supported values: String, Number, Binary.

Type: String

*Required:* Yes

*CloudFormation Compatibility:* This property is passed directly to the [AttributeType](#) property of the `AWS::DynamoDB::Table AttributeDefinition` data type.

## Examples

### PrimaryKey

Primary key example.

### YAML

```
Properties:
  PrimaryKey:
    Name: MyPrimaryKey
    Type: String
```

For reference information for all the AWS resource and property types that are supported by AWS CloudFormation and AWS SAM, see [AWS Resource and Property Types Reference](#) in the *AWS CloudFormation User Guide*.

# Resource Attributes

Resource attributes are attributes that you can add to a resource to control additional behaviors and relationships. For more information about resource attributes, see [Resource Attribute Reference](#) in the *AWS CloudFormation User Guide*.

AWS SAM resources support a subset of resource attributes that are supported by AWS CloudFormation resources. To see which AWS SAM resources support which resource attributes, see the following table.

Resource Type	CreationPolicy	DeletionPolicy	DependsOn	Metadata	UpdatePolicy	ReplacePolicy
<a href="#">AWS::Serverless::Api</a> (p. 28)	No	No	Yes	No	No	No
<a href="#">AWS::Serverless::Application</a> (p. 53)	No	No	Yes	No	No	No
<a href="#">AWS::Serverless::Function</a> (p. 56)	No	No	Yes	No	No	No
<a href="#">AWS::Serverless::HttpApi</a> (p. 102)	No	No	Yes	No	No	No

<a href="#">AWS::Serverless::LayerVersion</a> (p. 111)	No	Yes	Yes	No	No	No
<a href="#">AWS::Serverless::SimpleTable</a> (p. 114)	No	No	Yes	No	No	No

## Intrinsic Functions

Intrinsic functions are built-in functions that enable you to assign values to properties that are only available at runtime. For more information about intrinsic functions, see [Intrinsic Function Reference](#) in the *AWS CloudFormation User Guide*.

## API Gateway Extensions

API Gateway extensions are extensions to the OpenAPI specification that support the AWS-specific authorization and API Gateway-specific API integrations. For more information about API Gateway extensions, see [API Gateway Extensions to OpenAPI](#).

AWS SAM supports a subset of API Gateway extensions. To see which API Gateway extensions are supported by AWS SAM, see the following table.

API Gateway Extension	Supported by AWS SAM
<a href="#">x-amazon-apigateway-any-method</a> Object	Yes
<a href="#">x-amazon-apigateway-api-key-source</a> Property	No
<a href="#">x-amazon-apigateway-auth</a> Object	Yes
<a href="#">x-amazon-apigateway-authorizer</a> Object	Yes
<a href="#">x-amazon-apigateway-authtype</a> Property	Yes
<a href="#">x-amazon-apigateway-binary-media-types</a> Property	Yes
<a href="#">x-amazon-apigateway-documentation</a> Object	No
<a href="#">x-amazon-apigateway-endpoint-configuration</a> Object	No
<a href="#">x-amazon-apigateway-gateway-responses</a> Object	Yes
<a href="#">x-amazon-apigateway-gateway-responses.gatewayResponse</a> Object	Yes
<a href="#">x-amazon-apigateway-gateway-responses.responseParameters</a> Object	Yes
<a href="#">x-amazon-apigateway-gateway-responses.responseTemplates</a> Object	Yes
<a href="#">x-amazon-apigateway-integration</a> Object	Yes
<a href="#">x-amazon-apigateway-integration.requestTemplates</a> Object	Yes
<a href="#">x-amazon-apigateway-integration.requestParameters</a> Object	No
<a href="#">x-amazon-apigateway-integration.responses</a> Object	Yes
<a href="#">x-amazon-apigateway-integration.response</a> Object	Yes

<a href="#">x-amazon-apigateway-integration.responseTemplates</a> Object	Yes
<a href="#">x-amazon-apigateway-integration.responseParameters</a> Object	Yes
<a href="#">x-amazon-apigateway-request-validator</a> Property	No
<a href="#">x-amazon-apigateway-request-validators</a> Object	No
<a href="#">x-amazon-apigateway-request-validators.requestValidator</a> Object	No

# Authoring Serverless Applications

When you author a serverless application using AWS SAM, you construct an AWS SAM template to declare and configure the components of your application.

This section contains topics about how to validate your AWS SAM template, and how to build your application with dependencies. It also contains topics for how to use AWS SAM for certain use cases like working with Lambda layers, using nested applications, and controlling access to API Gateway APIs.

## Topics

- [Validating AWS SAM Template Files \(p. 120\)](#)
- [Building Applications with Dependencies \(p. 120\)](#)
- [Working with Layers \(p. 121\)](#)
- [Using Nested Applications \(p. 122\)](#)
- [Controlling Access to API Gateway APIs \(p. 125\)](#)

## Validating AWS SAM Template Files

Validate your templates with `sam validate` ([p. 176](#)). Currently, this command validates that the template provided is valid JSON / YAML. As with most AWS SAM CLI commands, it looks for a `template.[yaml|yml]` file in your current working directory by default. You can specify a different template file/location with the `-t` or `--template` option.

Example:

```
sam validate
<path-to-file>/template.yml is a valid SAM Template
```

### Note

The `sam validate` command requires AWS credentials to be configured. For more information, see [Configuration and Credential Files](#).

## Building Applications with Dependencies

You can use the `sam build` ([p. 163](#)) command to compile dependencies for Lambda functions written in Python. For example, if you write code that uses Python packages, such as a graphics library for image processing, you need to create a deployment package that works on the Amazon Linux AMI. The `sam build` command allows you to easily create deployment artifacts that target Lambda's execution environment, so that the functions you build locally run in a similar environment in the AWS Cloud.

The `sam build` command iterates through the functions in your application, looks for a manifest file (such as `requirements.txt`) that contain the dependencies, and automatically creates deployment artifacts that you can deploy to Lambda using the `sam package` and `sam deploy` commands.

If your Lambda function depends on packages that have natively compiled programs, you can use the `--use-container` flag. The `--use-container` flag compiles your functions in a Lambda-like environment locally, so they are in the right format when you deploy them to the AWS Cloud.

Examples:

```
# Build a deployment package
sam build

# Run the build process inside an AWS Lambda-like Docker container
sam build --use-container

# Build and run your functions locally
sam build && sam local invoke

# Build and package for deployment
sam build && sam package --s3-bucket <bucketname>

# For more options
sam build --help
```

## Working with Layers

The AWS SAM CLI supports applications that include layers. For more information about layers, see [Lambda Layers](#).

The following is an example AWS SAM template with a Lambda function that includes a layer:

```
ServerlessFunction:
  Type: AWS::Serverless::Function
  Properties:
    CodeUri: .
    Handler: my_handler
    Runtime: Python3.7
    Layers:
      - <LayerVersion ARN>
```

For more information about including layers in your application, see either [AWS::Serverless::Function](#) in the AWS SAM GitHub repository, or [AWS::Lambda::Function](#) in the *AWS CloudFormation User Guide*.

When you invoke your function using one of the `sam local` CLI subcommands, the layers package of your function is downloaded and cached on your local host. See the following chart for default cache directory locations. After the package is cached, the AWS SAM CLI overlays the layers onto a Docker image that's used to invoke your function. The AWS SAM CLI generates the names of the images it builds, as well as the `LayerVersions` that are held in the cache. You can find more details about the schema in the following sections.

To inspect the overlaid layers, execute the following command to start a bash session in the image that you want to inspect:

```
docker run -it --entrypoint=/bin/bash samcli/lambda:<Tag following the schema outlined in
  Docker Image Tag Schema> -i
```

### Layer Caching Directory name schema

Given a `LayerVersionArn` that's defined in your template, the AWS SAM CLI extracts the `LayerName` and `Version` from the ARN. It creates a directory to place the layer contents in named `LayerName-Version-<first 10 characters of sha256 of ARN>`.

Example:

```
ARN = arn:aws:lambda:us-west-2:111111111111:layer:myLayer:1
```

```
Directory name = myLayer-1-926eeb5ff1
```

### Docker Images tag schema

To compute the unique layers hash, combine all unique layer names with a delimiter of '-', take the SHA256 hash, and then take the first 25 characters.

Example:

```
ServerlessFunction:
  Type: AWS::Serverless::Function
  Properties:
    CodeUri: .
    Handler: my_handler
    Runtime: Python3.7
    Layers:
      - arn:aws:lambda:us-west-2:111111111111:layer:myLayer:1
      - arn:aws:lambda:us-west-2:111111111111:layer:mySecondLayer:1
```

Unique names are computed the same as the Layer Caching Directory name schema:

```
arn:aws:lambda:us-west-2:111111111111:layer:myLayer:1 = myLayer-1-926eeb5ff1
arn:aws:lambda:us-west-2:111111111111:layer:mySecondLayer:1 = mySecondLayer-1-6bc1022bdf
```

To compute the unique layers hash, combine all unique layer names with a delimiter of '-', take the sha256 hash, and then take the first 25 characters:

```
myLayer-1-926eeb5ff1-mySecondLayer-1-6bc1022bdf = 2dd7ac5ffb30d515926aef
```

Then combine this value with the function's runtime, with a delimiter of '-':

```
python3.7-2dd7ac5ffb30d515926aefffd
```

### Default Cache Directory Locations

OS	Location
Windows 7	C:\Users\<user>\AppData\Roaming\AWS SAM
Windows 8	C:\Users\<user>\AppData\Roaming\AWS SAM
Windows 10	C:\Users\<user>\AppData\Roaming\AWS SAM
macOS	~/aws-sam/layers-pkg
Unix	~/aws-sam/layers-pkg

## Using Nested Applications

A serverless application can include one or more **nested applications**. You can deploy a nested application as a stand-alone artifact or as a component of a larger application.

As serverless architectures grow, common patterns emerge in which the same components are defined in multiple application templates. You can now separate out common patterns as dedicated applications,

and then nest them as part of new or existing application templates. With nested applications, you can stay more focused on the business logic that's unique to your application.

To define a nested application in your serverless application, use the `AWS::Serverless::Application` resource type.

You can define nested applications from the following two sources:

- An **AWS Serverless Application Repository application** – You can define nested applications by using applications that are available to your account in the AWS Serverless Application Repository. These can be *private* applications in your account, applications that are *privately shared* with your account, or applications that are *publicly shared* in the AWS Serverless Application Repository. For more information about the different deployment permissions levels, see [Application Deployment Permissions](#) and [Publishing Applications](#) in the *AWS Serverless Application Repository Developer Guide*.
- A **local application** – You can define nested applications by using applications that are stored on your local file system.

See the following sections for details on how to use AWS SAM to define both of these types of nested applications in your serverless application.

**Note**

The maximum number of applications that can be nested in a serverless application is 200.  
The maximum number of parameters a nested application can have is 60.

## Defining a Nested Application from the AWS Serverless Application Repository

You can define nested applications by using applications that are available in the AWS Serverless Application Repository. You can also store and distribute applications that contain nested applications using the AWS Serverless Application Repository. To review details of a nested application in the AWS Serverless Application Repository, you can use the AWS SDK, the AWS CLI, or the Lambda console.

To define an application that's hosted in the AWS Serverless Application Repository in your serverless application's AWS SAM template, use the **Copy as SAM Resource** button on the detail page of every AWS Serverless Application Repository application. To do this, follow these steps:

1. Make sure that you're signed in to the AWS Management Console.
2. Find the application that you want to nest in the AWS Serverless Application Repository by using the steps in the [Browsing, Searching, and Deploying Applications](#) section of the *AWS Serverless Application Repository Developer Guide*.
3. Choose the **Copy as SAM Resource** button. The SAM template section for the application that you're viewing is now in your clipboard.
4. Paste the SAM template section into the `Resources` : section of the SAM template file for the application that you want to nest in this application.

The following is an example SAM template section for a nested application that's hosted in the AWS Serverless Application Repository:

```
Transform: AWS::Serverless-2016-10-31

Resources:
  applicationaliasname:
    Type: AWS::Serverless::Application
    Properties:
      Location:
```



```
ApplicationId: arn:aws:serverlessrepo:us-  
east-1:123456789012:applications/application-alias-name  
SemanticVersion: 1.0.0  
Parameters:  
  # Optional parameter that can have default value overridden  
  # ParameterName1: 15 # Uncomment to override default value  
  # Required parameter that needs value to be provided  
  ParameterName2: YOUR_VALUE
```

If there are no required parameter settings, you can omit the `Parameters:` section of the template.

### Important

Applications that contain nested applications hosted in the AWS Serverless Application Repository inherit the nested applications' sharing restrictions.

For example, suppose an application is publicly shared, but it contains a nested application that's only privately shared with the AWS account that created the parent application. In this case, if your AWS account doesn't have permission to deploy the nested application, you aren't able to deploy the parent application. For more information about permissions to deploy applications, see [Application Deployment Permissions](#) and [Publishing Applications](#) in the *AWS Serverless Application Repository Developer Guide*.

## Defining a Nested Application from the Local File System

You can define nested applications by using applications that are stored on your local file system. You do this by specifying the path to the AWS SAM template file that's stored on your local file system.

The following is an example SAM template section for a nested local application:

```
Transform: AWS::Serverless-2016-10-31  
  
Resources:  
  applicationaliasname:  
    Type: AWS::Serverless::Application  
    Properties:  
      Location: ../my-other-app/template.yaml  
      Parameters:  
        # Optional parameter that can have default value overridden  
        # ParameterName1: 15 # Uncomment to override default value  
        # Required parameter that needs value to be provided  
        ParameterName2: YOUR_VALUE
```

If there are no parameter settings, you can omit the `Parameters:` section of the template.

## Deploying Nested Applications

You can deploy your nested application by using the AWS SAM CLI command `sam deploy`. For more details, see [Deploying Serverless Applications \(p. 145\)](#).

### Note

When you deploy an application that contains nested applications, you must acknowledge that. You do this by passing `CAPABILITY_AUTO_EXPAND` to the [CreateCloudFormationChangeSet API](#), git status or using the `aws serverlessrepo create-cloud-formation-change-set` AWS CLI command.

For more information about acknowledging nested applications, see [Acknowledging IAM Roles, Resource Policies, and Nested Applications when Deploying Applications](#) in the *AWS Serverless Application Repository Developer Guide*.

## Controlling Access to API Gateway APIs

You can use AWS SAM to control who can access your API Gateway APIs by enabling authorization within your AWS SAM template.

AWS SAM supports several mechanisms for controlling access to your API Gateway APIs:

- **Lambda authorizers.** A Lambda authorizer (formerly known as a *custom authorizer*) is a Lambda function that you provide to control access to your API. When your API is called, this Lambda function is invoked with a request context or an authorization token that is provided by the client application. The Lambda function returns a policy document that specifies the operations that the caller is authorized to perform, if any. For more information about Lambda authorizers, see [Use API Gateway Lambda Authorizers](#) in the *API Gateway Developer Guide*. For examples of Lambda authorizers, see [Example: Defining Lambda Token Authorizers \(p. 126\)](#) and [Example: Defining Lambda Request Authorizers \(p. 126\)](#) later in this topic.
- **Amazon Cognito user pools.** Amazon Cognito user pools are user directories in Amazon Cognito. A client of your API must first sign a user in to the user pool, and obtain an identity or access token for the user. Then your API is called with one of the returned tokens. The API call succeeds only if the required token is valid. For more information about Amazon Cognito user pools, see [Control Access to REST API Using Amazon Cognito User Pools as Authorizer](#) in the *API Gateway Developer Guide*. For an example of Amazon Cognito user pools, see [Example: Defining Amazon Cognito User Pools \(p. 127\)](#) later in this topic.
- **IAM permissions.** You can control who can invoke your API using [IAM permissions](#). Users calling your API must be authenticated with IAM credentials. Calls to your API only succeed if there is an IAM policy attached to the IAM user that represents the API caller, an IAM group that contains the user, or an IAM role that is assumed by the user. For more information about IAM permissions, see [Control Access to an API with IAM Permissions](#) in the *API Gateway Developer Guide*. For an example of IAM permissions, see [Example: Defining IAM Permissions \(p. 128\)](#) later in this topic.
- **API keys.** API keys are alphanumeric string values that you distribute to application developer customers to grant access to your API. For more information about API keys, see [Create and Use Usage Plans with API Keys](#) in the *API Gateway Developer Guide*. For an example of API keys, see [Example: Defining API Keys \(p. 129\)](#) later in this topic.
- **Resource policies.** Resource policies are JSON policy documents that you can attach to an API Gateway API to control whether a specified principal (typically an IAM user or role) can invoke the API. For more information about resource policies, see [Control Access to an API with Amazon API Gateway Resource Policies](#) in the *API Gateway Developer Guide*. For an example of resource policies, see [Example: Defining Resource Policies \(p. 129\)](#) later in this topic.

In addition, you can use AWS SAM to customize the content of some API Gateway error responses. For more information about customizing API Gateway error responses, see [Set Up Gateway Responses to Customize Error Responses](#). For an example of customized responses, see [Example: Defining Customized Responses \(p. 130\)](#) later in this topic.

## Choosing a Mechanism to Control Access

The mechanism that you choose to control access to your API Gateway APIs depends on a few factors. For example, if you have a greenfield project that doesn't have either authorization or access control set

up yet, then Amazon Cognito user pools might be your best option. This is because when you set up user pools, you also set up both authentication and access control automatically.

However, if your application already has authentication set up, then using Lambda authorizers might be the best option. This is because you can call your existing authentication service and return a policy document based on the response. Also, if your application requires custom authentication or access control logic that user pools don't support, then Lambda authorizers might be your best option.

After you've decided which mechanism to use, see the corresponding section in this topic to see how to use AWS SAM to configure your application to use that mechanism.

## Example: Defining Lambda Token Authorizers

You can control access to your APIs by defining a Lambda Token authorizer within your AWS SAM template. To do this, you use the [API Auth Object](#) data type.

The following is an example AWS SAM template section for a Lambda Token authorizer:

```
Resources:
  MyApi:
    Type: AWS::Serverless::Api
    Properties:
      StageName: Prod
      Auth:
        DefaultAuthorizer: MyLambdaTokenAuthorizer
      Authorizers:
        MyLambdaTokenAuthorizer:
          FunctionArn: !GetAtt MyAuthFunction.Arn

  MyFunction:
    Type: AWS::Serverless::Function
    Properties:
      CodeUri: ./src
      Handler: index.handler
      Runtime: nodejs8.10
      Events:
        GetRoot:
          Type: Api
          Properties:
            RestApiId: !Ref MyApi
            Path: /
            Method: get

  MyAuthFunction:
    Type: AWS::Serverless::Function
    Properties:
      CodeUri: ./src
      Handler: authorizer.handler
      Runtime: nodejs8.10
```

For more information about API Gateway Lambda authorizers, see [Use API Gateway Lambda Authorizers](#) in the *API Gateway Developer Guide*.

For a full sample application that includes a Lambda Token authorizer, see [API Gateway + Lambda TOKEN Authorizer Example](#).

## Example: Defining Lambda Request Authorizers

You can control access to your APIs by defining a Lambda Request authorizer within your AWS SAM template. To do this, you use the [API Auth Object](#) data type.

The following is an example AWS SAM template section for a Lambda Request authorizer:

```
Resources:
  MyApi:
    Type: AWS::Serverless::Api
    Properties:
      StageName: Prod
      Auth:
        DefaultAuthorizer: MyLambdaRequestAuthorizer
        Authorizers:
          MyLambdaRequestAuthorizer:
            FunctionPayloadType: REQUEST
            FunctionArn: !GetAtt MyAuthFunction.Arn
            Identity:
              QueryStrings:
                - auth

  MyFunction:
    Type: AWS::Serverless::Function
    Properties:
      CodeUri: ./src
      Handler: index.handler
      Runtime: nodejs8.10
      Events:
        GetRoot:
          Type: Api
          Properties:
            RestApiId: !Ref MyApi
            Path: /
            Method: get

  MyAuthFunction:
    Type: AWS::Serverless::Function
    Properties:
      CodeUri: ./src
      Handler: authorizer.handler
      Runtime: nodejs8.10
```

For more information about API Gateway Lambda authorizers, see [Use API Gateway Lambda Authorizers](#) in the *API Gateway Developer Guide*.

For a full sample application that includes a Lambda Request authorizer, see [API Gateway + Lambda REQUEST Authorizer Example](#).

## Example: Defining Amazon Cognito User Pools

You can control access to your APIs by defining Amazon Cognito user pools within your AWS SAM template. To do this, you use the [API Auth Object](#) data type.

The following is an example AWS SAM template section for a user pool:

```
Resources:
  MyApi:
    Type: AWS::Serverless::Api
    Properties:
      StageName: Prod
      Cors: "*"
      Auth:
        DefaultAuthorizer: MyCognitoAuthorizer
        Authorizers:
          MyCognitoAuthorizer:
            UserPoolArn: !GetAtt MyCognitoUserPool.Arn
```

```
MyFunction:
  Type: AWS::Serverless::Function
  Properties:
    CodeUri: ./src
    Handler: lambda.handler
    Runtime: nodejs8.10
    Events:
      Root:
        Type: Api
        Properties:
          RestApiId: !Ref MyApi
          Path: /
          Method: GET

MyCognitoUserPool:
  Type: AWS::Cognito::UserPool
  Properties:
    UserPoolName: !Ref CognitoUserPoolName
    Policies:
      PasswordPolicy:
        MinimumLength: 8
    UsernameAttributes:
      - email
    Schema:
      - AttributeDataType: String
        Name: email
        Required: false

MyCognitoUserPoolClient:
  Type: AWS::Cognito::UserPoolClient
  Properties:
    UserPoolId: !Ref MyCognitoUserPool
    ClientName: !Ref CognitoUserPoolClientName
    GenerateSecret: false
```

For more information about Amazon Cognito user pools, see [Control Access to a REST API Using Amazon Cognito User Pools as Authorizer](#) in the *API Gateway Developer Guide*.

For a full sample application that includes a user pool as an authorizer, see [API Gateway + Cognito Auth + Cognito Hosted Auth Example](#).

## Example: Defining IAM Permissions

You can control access to your APIs by defining IAM permissions within your AWS SAM template. To do this, you use the [API Auth Object](#) data type.

The following is an example AWS SAM template section for IAM permissions:

```
Resources:
  MyApi:
    Type: AWS::Serverless::Api
    Properties:
      StageName: Prod
      Auth:
        DefaultAuthorizer: AWS_IAM

  MyFunction:
    Type: AWS::Serverless::Function
    Properties:
      CodeUri: .
      Handler: index.handler
      Runtime: nodejs8.10
```

```
Events:
  GetRoot:
    Type: Api
    Properties:
      RestApiId: !Ref MyApi
      Path: /
      Method: get
```

For more information about IAM permissions, see [Control Access to an API Using IAM Permissions](#) in the *API Gateway Developer Guide*.

For a full sample application that includes a user pool as an authorizer, see [API Gateway + IAM Permissions Example](#).

## Example: Defining API Keys

You can control access to your APIs by requiring API keys within your AWS SAM template. To do this, you use the [API Auth Object](#) data type.

The following is an example AWS SAM template section for API keys:

```
Resources:
  MyApi:
    Type: AWS::Serverless::Api
    Properties:
      StageName: Prod
      Auth:
        ApiKeyRequired: true # sets for all methods

  MyFunction:
    Type: AWS::Serverless::Function
    Properties:
      CodeUri: .
      Handler: index.handler
      Runtime: nodejs8.10
      Events:
        ApiKey:
          Type: Api
          Properties:
            RestApiId: !Ref MyApi
            Path: /
            Method: get
            Auth:
              ApiKeyRequired: true
```

For more information about API keys, see [Create and Use Usage Plans with API Keys](#) in the *API Gateway Developer Guide*.

## Example: Defining Resource Policies

You can control access to your APIs by attaching a resource policy within your AWS SAM template. To do this, you use the [API Auth Object](#) data type.

The following is an example AWS SAM template section for resource policies:

```
Resources:
  ExplicitApi:
    Type: AWS::Serverless::Api
    Properties:
      StageName: Prod
      EndpointConfiguration: PRIVATE
```

```
Auth:
  ResourcePolicy:
    CustomStatements: {
      Effect: 'Allow',
      Action: 'execute-api:Invoke',
      Resource: ['execute-api:/*/*/*'],
      Principal: '*'
    }
MinimalFunction:
  Type: 'AWS::Serverless::Function'
  Properties:
    CodeUri: s3://sam-demo-bucket/hello.zip
    Handler: hello.handler
    Runtime: python2.7
    Events:
      AddItem:
        Type: Api
        Properties:
          RestApiId:
            Ref: ExplicitApi
          Path: /add
          Method: post
```

For more information about resource policies, see [Control Access to an API with Amazon API Gateway Resource Policies](#) in the *API Gateway Developer Guide*.

## Example: Defining Customized Responses

You can customize some API Gateway error responses by defining response headers within your AWS SAM template. To do this, you use the [Gateway Response Object](#) data type.

The following is an example AWS SAM template section for API Gateway responses:

```
Resources:
  MyApi:
    Type: AWS::Serverless::Api
    Properties:
      StageName: Prod
      GatewayResponses:
        DEFAULT_4xx:
          ResponseParameters:
            Headers:
              Access-Control-Expose-Headers: "'WWW-Authenticate'"
              Access-Control-Allow-Origin: "'*'"

  GetFunction:
    Type: AWS::Serverless::Function
    Properties:
      Handler: index.get
      Runtime: nodejs6.10
      InlineCode: module.exports = async () => throw new Error('Check out the response headers!')
      Events:
        GetResource:
          Type: Api
          Properties:
            Path: /error
            Method: get
            RestApiId: !Ref MyApi
```

For more information about customizing API Gateway messages, see [Set Up Gateway Responses to Customize Error Responses](#) in the *API Gateway Developer Guide*.

For a full sample application that includes a customized error response, see [API Gateway + GatewayResponse Example](#).



# Testing and Debugging Serverless Applications

With the AWS SAM command line interface (CLI), you can locally test and "step-through" debug your serverless applications before uploading your application to the AWS Cloud. You can verify whether your application is behaving as expected, debug what's wrong, and fix any issues, before going through the steps of packaging and deploying your application.

When you locally invoke a Lambda function in debug mode within the AWS SAM CLI, you can then attach a debugger to it. With the debugger, you can step through your code line by line, see the values of various variables, and fix issues the same way you would for any other application.

## Topics

- [Invoking Functions Locally \(p. 132\)](#)
- [Running API Gateway Locally \(p. 134\)](#)
- [Running Automated Tests \(p. 136\)](#)
- [Generating Sample Event Payloads \(p. 138\)](#)
- [Step-Through Debugging Lambda Functions Locally \(p. 138\)](#)
- [Passing Additional Runtime Debug Arguments \(p. 144\)](#)

## Invoking Functions Locally

You can invoke your function locally by using the `sam local invoke` (p. 169) command and providing its function logical ID and an event file. Alternatively, `sam local invoke` also accepts `stdin` as an event.

### Note

The `sam local invoke` command described in this section corresponds to the AWS CLI command `aws lambda invoke`. You can use either version of this command to invoke a Lambda function that you've uploaded to the AWS Cloud.

You must execute `sam local invoke` in the project directory containing the function you want to invoke.

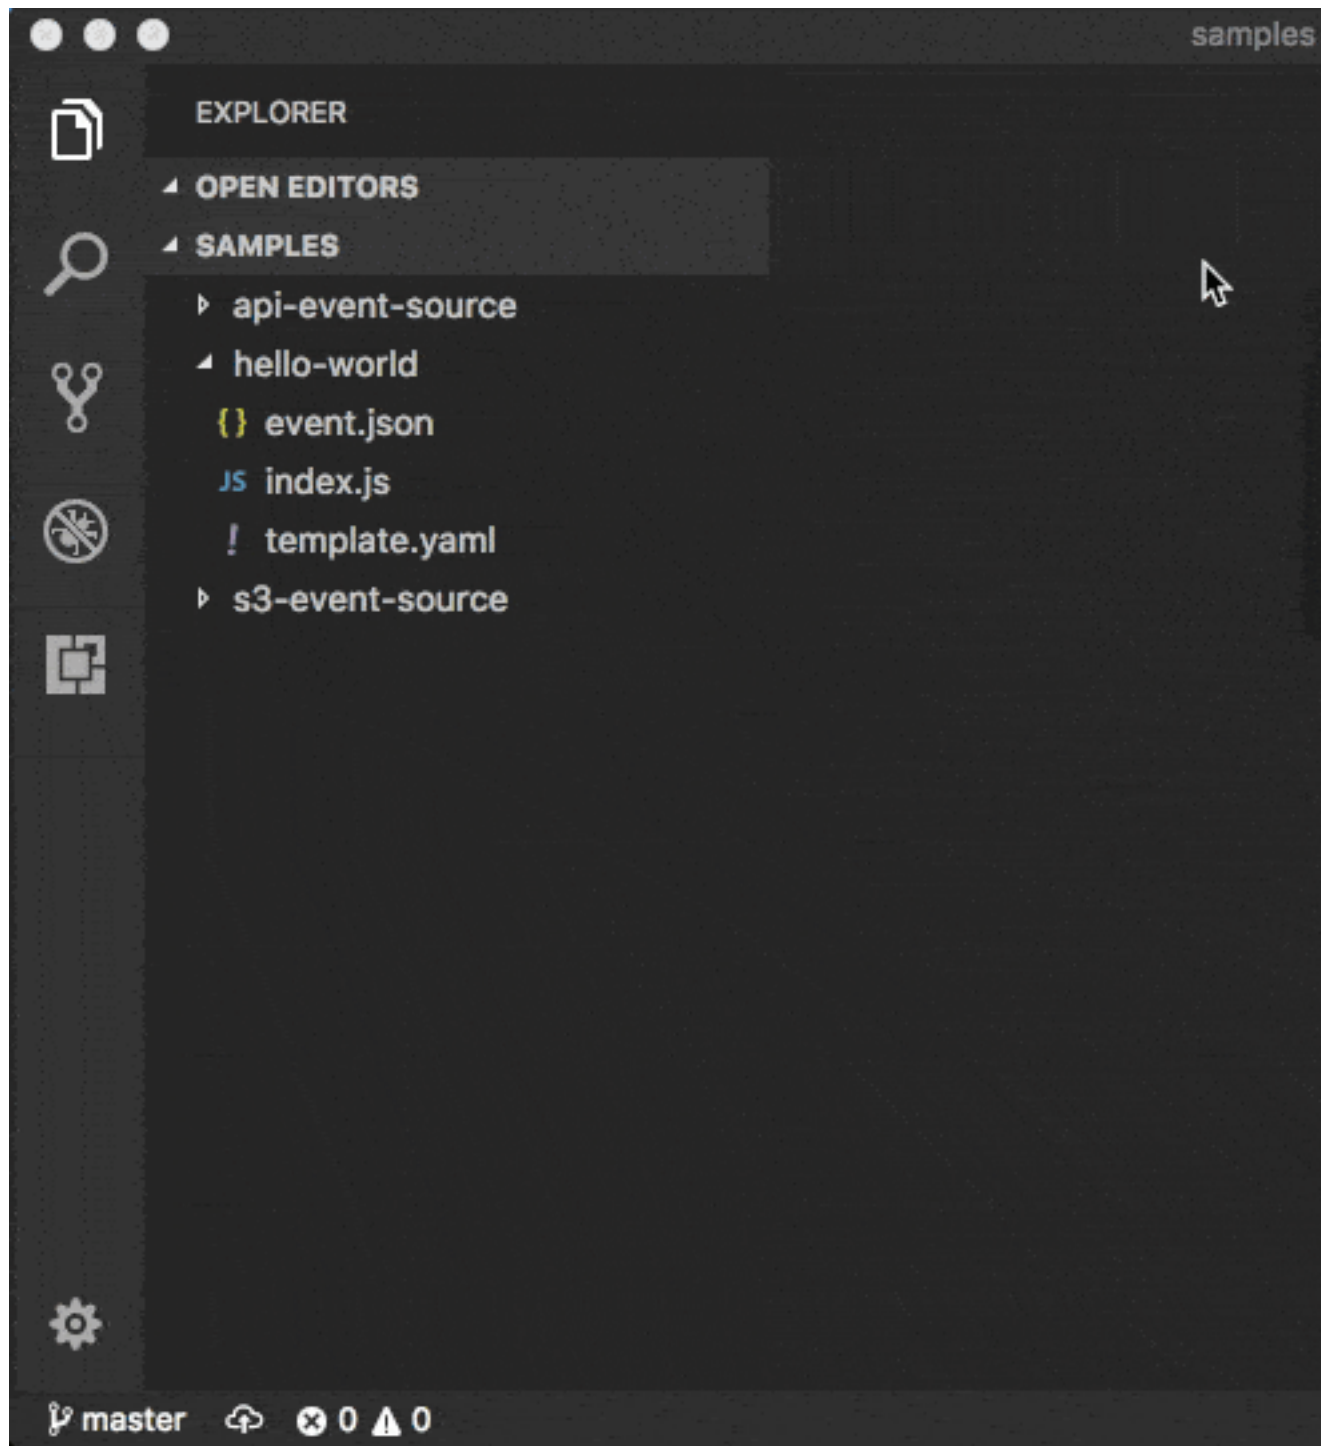
Examples:

```
# Invoking function with event file
$ sam local invoke "Ratings" -e event.json

# Invoking function with event via stdin
$ echo '{"message": "Hey, are you there?" }' | sam local invoke "Ratings"

# For more options
$ sam local invoke --help
```

This animation shows invoking a Lambda function locally using Microsoft Visual Studio Code:



## Environment Variable File

You can use the `--env-vars` argument with the `invoke` or `start-api` commands. You do this to provide a JSON file that contains values to override the environment variables that are already defined in your function template. Structure the file as follows:

```
{
  "MyFunction1": {
    "TABLE_NAME": "localtable",
    "BUCKET_NAME": "testBucket"
  },
  "MyFunction2": {
    "TABLE_NAME": "localtable",
    "STAGE": "dev"
  }
}
```

For example, if you save this content in a file named `env.json`, then the following command uses this file to override the included environment variables:

```
sam local invoke --env-vars env.json
```

## Layers

If your application includes layers, see [Working with Layers \(p. 121\)](#) for more information about how to debug layers issues on your local host.

## Running API Gateway Locally

Use the `sam local start-api` (p. 170) command to start a local instance of API Gateway that you will use to test HTTP request/response functionality. This functionality features hot reloading to enable you to quickly develop and iterate over your functions.

### Note

"Hot reloading" is when only the files that changed are refreshed without losing the state of the application. In contrast, "live reloading" is when the entire application is refreshed, such that the state of the application is lost.


You must execute `sam local invoke` in the project directory containing the function you want to invoke.

Example:

```
sam local start-api
```

AWS SAM automatically finds any functions within your AWS SAM template that have `Api` event sources defined. Then, it mounts them at the defined HTTP paths.

This animation shows running API Gateway locally using Microsoft Visual Studio Code:



In the following example, the Ratings function mounts `ratings.py:handler()` at `/ratings` for GET requests:

```
Ratings:
  Type: AWS::Serverless::Function
  Properties:
    Handler: ratings.handler
```

```
Runtime: python3.6
Events:
  Api:
    Type: Api
    Properties:
      Path: /ratings
      Method: get
```

By default, AWS SAM uses [Proxy Integration](#) and expects the response from your Lambda function to include one or more of the following: `statusCode`, `headers`, or `body`.

For example:

```
// Example of a Proxy Integration response
exports.handler = (event, context, callback) => {
  callback(null, {
    statusCode: 200,
    headers: { "x-custom-header" : "my custom header value" },
    body: "hello world"
  });
}
```

For examples in other AWS Lambda languages, see [Proxy Integration](#).

#### *Environment Variable File*

You can use the `--env-vars` argument with the `invoke` or `start-api` commands to provide a JSON file that contains values to override the environment variables already defined in your function template. Structure the file as follows:

```
{
  "MyFunction1": {
    "TABLE_NAME": "localtable",
    "BUCKET_NAME": "testBucket"
  },
  "MyFunction2": {
    "TABLE_NAME": "localtable",
    "STAGE": "dev"
  },
}
```

For example, if you save this content in a file named `env.json`, then the following command uses this file to override the included environment variables:

```
sam local start-api --env-vars env.json
```

## Layers

If your application includes layers, see [Working with Layers \(p. 121\)](#) for more information about how to debug layers issues on your local host.

## Running Automated Tests

You can use the `sam local invoke` command to manually test your code by running Lambda functions locally. With the AWS SAM CLI, you can easily author automated integration tests by first running tests against local Lambda functions before deploying to the AWS Cloud.

The `sam local start-lambda` command starts a local endpoint that emulates the AWS Lambda invoke endpoint. You can invoke it from your automated tests. Because this endpoint emulates the AWS Lambda invoke endpoint, you can write tests once, and then run them (without any modifications) against the local Lambda function, or against a deployed Lambda function. You can also run the same tests against a deployed AWS SAM stack in your CI/CD pipeline.

This is how the process works:

1. Start the local Lambda endpoint.

Start the local Lambda endpoint by running the following command in the directory that contains your AWS SAM template:

```
sam local start-lambda
```

This command starts a local endpoint at `http://127.0.0.1:3001` that emulates AWS Lambda. You can run your automated tests against this local Lambda endpoint. When you invoke this endpoint using the AWS CLI or SDK, it locally executes the Lambda function that's specified in the request, and returns a response.

2. Run an integration test against the local Lambda endpoint.

In your integration test, you can use the AWS SDK to invoke your Lambda function with test data, wait for response, and verify that the response is what you expect. To run the integration test locally, you should configure the AWS SDK to send a Lambda Invoke API call to invoke the local Lambda endpoint that you started in previous step.

The following is a Python example (the AWS SDKs for other languages have similar configurations):

```
import boto3
import botocore

# Set "running_locally" flag if you are running the integration test locally
running_locally = True

if running_locally:
    # Create Lambda SDK client to connect to appropriate Lambda endpoint
    lambda_client = boto3.client('lambda',
        region_name="us-west-2",
        endpoint_url="http://127.0.0.1:3001",
        use_ssl=False,
        verify=False,
        config=botocore.client.Config(
            signature_version=botocore.UNSIGNED,
            read_timeout=1,
            retries={'max_attempts': 0},
        )
    )
else:
    lambda_client = boto3.client('lambda')

# Invoke your Lambda function as you normally usually do. The function will run
# locally if it is configured to do so
response = lambda_client.invoke(FunctionName="HelloWorldFunction")

# Verify the response
assert response == "Hello World"
```

You can use this code to test deployed Lambda functions by setting `running_locally` to `False`. This sets up the AWS SDK to connect to AWS Lambda in the AWS Cloud.

## Generating Sample Event Payloads

To make local development and testing of Lambda functions easier, you can generate and customize event payloads for a number of AWS services like API Gateway, AWS CloudFormation, Amazon S3, and so on.

For the full list of services that you can generate sample event payloads for, use this command:

```
sam local generate-event --help
```

For the list of options you can use for a particular service, use this command:

```
sam local generate-event [SERVICE] --help
```

Examples:

```
#Generates the event from S3 when a new object is created
sam local generate-event s3 put

# Generates the event from S3 when an object is deleted
sam local generate-event s3 delete
```

## Step-Through Debugging Lambda Functions Locally

You can use AWS SAM with a number of AWS toolkits to test and debug your serverless applications locally.

For example, you can perform step-through debugging of your Lambda functions. Step-through debugging makes it easier to understand what the code is doing. It tightens the feedback loop by making it possible for you to find and troubleshoot issues that you might run into in the cloud.

### Using AWS Toolkits

AWS toolkits are plugins that provide you with the ability to perform many common debugging tasks, like setting breakpoints, executing code line by line, and inspecting the values of variables. Toolkits make it easier for you to develop, debug, and deploy serverless applications that are built using AWS. They provide an experience for building, testing, debugging, deploying, and invoking Lambda functions that's integrated into the integrated development environment (IDE).

For more information about AWS toolkits that you can use with AWS SAM, see the following:

- [AWS Toolkit for JetBrains](#)
- [AWS Toolkit for PyCharm](#)

- [AWS Toolkit for IntelliJ](#)
- [AWS Toolkit for Visual Studio Code](#)

## Running AWS SAM Locally

The commands `sam local invoke` and `sam local start-api` both support local step-through debugging of your Lambda functions. To run AWS SAM locally with step-through debugging support enabled, specify `--debug-port` or `-d` on the command line. For example:

```
# Invoke a function locally in debug mode on port 5858
sam local invoke -d 5858 <function logical id>

# Start local API Gateway in debug mode on port 5858
sam local start-api -d 5858
```

### Note

If you're using `sam local start-api`, the local API Gateway instance exposes all of your Lambda functions. However, because you can specify a single debug port, you can only debug one function at a time. You need to call your API before the AWS SAM CLI binds to the port, which allows the debugger to connect.

### Topics


The following topics provide examples of how to set up your environment to test and debug your serverless applications locally.

- [Step-Through Debugging Node.js Functions Locally \(p. 139\)](#)
- [Step-Through Debugging Python Functions Locally \(p. 141\)](#)
- [Step-Through Debugging Golang Functions Locally \(p. 143\)](#)

## Step-Through Debugging Node.js Functions Locally

The following is an example that shows how to debug a Node.js function with Microsoft Visual Studio Code:





To set up Microsoft Visual Studio Code for step-through debugging Node.js functions with the AWS SAM CLI, use the following launch configuration. Before you do this, set the directory where the `template.yaml` file is located as the workspace root in Microsoft Visual Studio Code:

```
{
  "version": "0.2.0",
  "configurations": [
```

```
{
  "name": "Attach to SAM CLI",
  "type": "node",
  "request": "attach",
  "address": "localhost",
  "port": 5858,
  // From the sam init example, it would be "${workspaceRoot}/hello-world"
  "localRoot": "${workspaceRoot}/{directory of node app}",
  "remoteRoot": "/var/task",
  "protocol": "inspector",
  "stopOnEntry": false
}
]
```

**Note**

The `localRoot` is set based on what the `CodeUri` points at in the `template.yaml` file. If there are nested directories within the `CodeUri`, that needs to be reflected in the `localRoot`.

**Note**

Node.js versions earlier than 7 (for example, Node.js 4.3 and Node.js 6.10) use the `legacy` protocol, while Node.js versions including and later than 7 (for example, Node.js 8.10) use the `inspector` protocol. Be sure to specify the corresponding protocol in the `protocol` entry of your launch configuration. This was tested with Microsoft Visual Studio Code versions 1.26, 1.27, and 1.28 for the `legacy` and `inspector` protocols.

## Step-Through Debugging Python Functions Locally

Python step-through debugging requires you to enable remote debugging in your Lambda function code. This is a two-step process:

1. Install the [ptvsd library](#) and enable it within your code.
2. Configure your IDE to connect to the debugger that you configured for your function.

Because this might be your first time using the AWS SAM CLI, start with a boilerplate Python application, and install both the application's dependencies and `ptvsd`:

```
sam init --runtime python3.6 --name python-debugging
cd python-debugging/

# Install dependencies of our boilerplate app
pip install -r hello_world/requirements.txt -t hello_world/build/

# Install ptvsd library for step through debugging
pip install ptvsd -t hello_world/build/

cp hello_world/app.py hello_world/build/
```

## Ptvsd Configuration

Next, you need to enable `ptvsd` within your code. To do this, open `hello_world/build/app.py`, and add the following `ptvsd` specifics:

```
import ptvsd

# Enable ptvsd on 0.0.0.0 address and on port 5890 that we'll connect later with our IDE
ptvsd.enable_attach(address=('0.0.0.0', 5890), redirect_output=True)
ptvsd.wait_for_attach()
```

Use `0.0.0.0` instead of `localhost` for listening across all network interfaces. 5890 is the debugging port that you want to use.

## Microsoft Visual Studio Code

Now that you have the dependencies and `ptvsd` enabled within your code, you can configure Microsoft Visual Studio Code debugging. Assuming that you're still in the application folder and have the code command in your path, open Microsoft Visual Studio Code by using this command:

```
code .
```

### Note

If you don't have code in your path, open a new instance of Microsoft Visual Studio Code from the `python-debugging/` folder that you created earlier.

To set up Microsoft Visual Studio Code for debugging with the AWS SAM CLI, use the following launch configuration:

```
{
  "version": "0.2.0",
  "configurations": [
    {
      "name": "SAM CLI Python Hello World",
      "type": "python",
      "request": "attach",
      "port": 5890,
      "host": "localhost",
      "pathMappings": [
        {
          "localRoot": "${workspaceFolder}/hello_world/build",
          "remoteRoot": "/var/task"
        }
      ]
    }
  ]
}
```

For Microsoft Visual Studio Code, the property `localRoot` under the `pathMappings` key is important. There are two reasons that help explain this setup:

- **localRoot:** This path is to be mounted in the Docker container, and needs to have both the application and dependencies at the root level.
- **workspaceFolder:** This path is the absolute path where the Microsoft Visual Studio Code instance was opened.

If you opened Microsoft Visual Studio Code in a different location other than `python-debugging/`, you need to replace it with the absolute path where `python-debugging/` is located.

After the Microsoft Visual Studio Code debugger configuration is complete, make sure to add a breakpoint anywhere you want to in `hello_world/build/app.py`, and then proceed as follows:

1. Run the AWS SAM CLI to invoke your function.
2. Send a request to the URL to invoke the function and initialize `ptvsd` code execution.
3. Start the debugger within Microsoft Visual Studio Code.

```
# Remember to hit the URL before starting the debugger in Microsoft Visual Studio Code
```

```
sam local start-api -d 5890

# OR

# Change HelloWorldFunction to reflect the logical name found in template.yaml
sam local generate-event apigateway aws-proxy | sam local invoke HelloWorldFunction -d 5890
```

## Step-Through Debugging Golang Functions Locally

Golang function step-through debugging is slightly different when compared to Node.js, Java, and Python. We require [Delve](#) as the debugger, and wrap your function with it at runtime. The debugger is run in headless mode, listening on the debug port.

When you're debugging, you must compile your function in debug mode:

```
GOARCH=amd64 GOOS=linux go build -gcflags='-N -l' -o <output path> <path to code directory>
```

### Delve Debugger

You must compile [Delve](#) to run in the container and provide its local path with the `--debugger-path` argument.

Build [Delve](#) locally as follows:

```
GOARCH=amd64 GOOS=linux go build -o <delve folder path>/dlv github.com/go-delve/delve/cmd/dlv
```

### Delve Debugger Path

The output path needs to end in `/dlv`. The Docker container expects the `dlv` binary file to be in the `<delve folder path>`. If it's not, a mounting issue occurs.

#### Note

The `--debugger-path` is the path to the directory that contains the `dlv` binary file that's compiled from the previous code.

#### Example:

Invoke AWS SAM similar to the following:

```
sam local start-api -d 5986 --debugger-path <delve folder path>
```

### Delve Debugger API Version

To run the [Delve](#) debugger with an API version of your choice, specify the desired API version using an [additional debug argument \(p. 144\)](#) `-delveAPI`.

#### Note

For IDEs such as GoLand, Microsoft Visual Studio Code, etc., it is important to run [Delve](#) in API version 2 mode.

#### Example

Invoke AWS SAM with the [Delve](#) debugger in API version 2 mode:

```
sam local start-api -d 5986 --debugger-path <delve folder path> --debug-args "-delveAPI=2"
```

## Example

The following is an example launch configuration for Microsoft Visual Studio Code to attach to a debug session.

```
{
  "version": "0.2.0",
  "configurations": [
    {
      "name": "Connect to Lambda container",
      "type": "go",
      "request": "launch",
      "mode": "remote",
      "remotePath": "",
      "port": <debug port>,
      "host": "127.0.0.1",
      "program": "${workspaceRoot}",
      "env": {},
      "args": [],
    },
  ]
}
```

## Passing Additional Runtime Debug Arguments

To pass additional runtime arguments when you're debugging your function, use the environment variable `DEBUGGER_ARGS`. This passes a string of arguments directly into the run command that the AWS SAM CLI uses to start your function.

For example, if you want to load a debugger like iKpdb at the runtime of your Python function, you could pass the following as `DEBUGGER_ARGS`: `-m ikpdb --ikpdb-port=5858 --ikpdb-working-directory=/var/task/ --ikpdb-client-working-directory=/myApp --ikpdb-address=0.0.0.0`. This would load iKpdb at runtime with the other arguments you've specified.

In this case, your full AWS SAM CLI command would be:

```
DEBUGGER_ARGS="-m ikpdb --ikpdb-port=5858 --ikpdb-working-directory=/var/task/ --ikpdb-client-working-directory=/myApp --ikpdb-address=0.0.0.0" echo {} | sam local invoke -d 5858 myFunction
```

You can pass debugger arguments to the functions of all runtimes.

# Deploying Serverless Applications

AWS SAM uses AWS CloudFormation as the underlying deployment mechanism. For more information, see [What Is AWS CloudFormation?](#).

You can deploy your application by using AWS SAM command line interface (CLI) commands. You can also use other AWS services that integrate with AWS SAM to automate your deployments.

## Packaging and Deploying Using the AWS SAM CLI

After you develop and test your serverless application locally, you can deploy your application by using the `sam package` and `sam deploy` commands.

### Note

Both the `sam package` and `sam deploy` commands described in this section are identical to their AWS CLI equivalent commands `aws cloudformation package` and `aws cloudformation deploy`, respectively.

The `sam package` command zips your code artifacts, uploads them to Amazon S3, and produces a packaged AWS SAM template file that's ready to be used. The `sam deploy` command uses this file to deploy your application. For example, the following command generates a `packaged.yaml` file:

```
# Package SAM template
sam package --template-file sam.yaml --s3-bucket mybucket --output-template-file
packaged.yaml
```

The following `sam deploy` command takes the packaged AWS SAM template file that was created earlier, and deploys your serverless application:

```
# Deploy packaged SAM template
sam deploy --template-file ./packaged.yaml --stack-name mystack --capabilities
CAPABILITY_IAM
```

### Note

To deploy an application that contains one or more nested applications, you must include the `CAPABILITY_AUTO_EXPAND` capability in the `sam deploy` command.

## Publishing Serverless Applications

The AWS Serverless Application Repository is a service that hosts serverless applications that are built using AWS SAM. If you want to share serverless applications with others, you can publish them in the AWS Serverless Application Repository. You can also search, browse, and deploy serverless applications that have been published by others. For more information, see [What Is the AWS Serverless Application Repository?](#).

## Automating Deployments

You can use AWS SAM with a number of other AWS services to automate the deployment process of your serverless application.

- **CodeBuild:** You use CodeBuild to build, locally test, and package your serverless application. For more information, see [What Is CodeBuild?](#).

- **CodeDeploy:** You use [CodeDeploy](#) to gradually deploy updates to your serverless applications. For more information on how to do this, see [Deploying Serverless Applications Gradually \(p. 146\)](#).
- **CodePipeline:** You use CodePipeline to model, visualize, and automate the steps that are required to release your serverless application. For more information, see [What Is CodePipeline?](#).

### Topics

- [Deploying Serverless Applications Gradually \(p. 146\)](#)

## Deploying Serverless Applications Gradually

If you use AWS SAM to create your serverless application, it comes built-in with [CodeDeploy](#) to help ensure safe Lambda deployments. With just a few lines of configuration, AWS SAM does the following for you:

- Deploys new versions of your Lambda function, and automatically creates aliases that point to the new version.
- Gradually shifts customer traffic to the new version until you're satisfied that it's working as expected, or you roll back the update.
- Defines pre-traffic and post-traffic test functions to verify that the newly deployed code is configured correctly and your application operates as expected.
- Rolls back the deployment if CloudWatch alarms are triggered.

### Note

If you enable gradual deployments through your AWS SAM template, a CodeDeploy resource is automatically created for you. You can view the CodeDeploy resource directly through the AWS Management Console.

### Example

The following example demonstrates a simple version of using CodeDeploy to gradually shift customers to your newly deployed version:

```
Resources:
  MyLambdaFunction:
    Type: AWS::Serverless::Function
    Properties:
      Handler: index.handler
      Runtime: nodejs4.3
      CodeUri: s3://bucket/code.zip

      AutoPublishAlias: live

  DeploymentPreference:
    Type: Canary10Percent10Minutes
    Alarms:
      # A list of alarms that you want to monitor
      - !Ref AliasErrorMetricGreaterThanZeroAlarm
      - !Ref LatestVersionErrorMetricGreaterThanZeroAlarm
    Hooks:
      # Validation Lambda functions that are run before & after traffic shifting
      PreTraffic: !Ref PreTrafficLambdaFunction
      PostTraffic: !Ref PostTrafficLambdaFunction
```

These revisions to the AWS SAM template do the following:

- **AutoPublishAlias:** By adding this property and specifying an alias name, AWS SAM:
  - Detects when new code is being deployed, based on changes to the Lambda function's Amazon S3 URI.
  - Creates and publishes an updated version of that function with the latest code.
  - Creates an alias with a name that you provide (unless an alias already exists), and points to the updated version of the Lambda function. Function invocations should use the alias qualifier to take advantage of this. If you aren't familiar with Lambda function versioning and aliases, see [AWS Lambda Function Versioning and Aliases](#).
- **Deployment Preference Type:** In the previous example, 10 percent of your customer traffic is immediately shifted to your new version. After 10 minutes, all traffic is shifted to the new version. However, if your pre-hook/post-hook tests fail, or if a CloudWatch alarm is triggered, CodeDeploy rolls back your deployment. The following table outlines other traffic-shifting options that are available beyond the one used earlier. Note the following:
  - **Canary:** Traffic is shifted in two increments. You can choose from predefined canary options. The options specify the percentage of traffic that's shifted to your updated Lambda function version in the first increment, and the interval, in minutes, before the remaining traffic is shifted in the second increment.
  - **Linear:** Traffic is shifted in equal increments with an equal number of minutes between each increment. You can choose from predefined linear options that specify the percentage of traffic that's shifted in each increment and the number of minutes between each increment.
  - **All-at-once:** All traffic is shifted from the original Lambda function to the updated Lambda function version at once.

Deployment Preference Type
Canary10Percent30Minutes
Canary10Percent5Minutes
Canary10Percent10Minutes
Canary10Percent15Minutes
Linear10PercentEvery10Minutes
Linear10PercentEvery1Minute
Linear10PercentEvery2Minutes
Linear10PercentEvery3Minutes
AllAtOnce

- **Alarms:** These are CloudWatch alarms that are triggered by any errors raised by the deployment. They automatically roll back your deployment. An example is if the updated code you're deploying is creating errors within the application. Another example is if any [AWS Lambda](#) or custom CloudWatch metrics that you specified have breached the alarm threshold.
- **Hooks:** These are pre-traffic and post-traffic test functions that run sanity checks before traffic shifting starts to the new version, and after traffic shifting completes.
  - **PreTraffic:** Before traffic shifting starts, CodeDeploy invokes the pre-traffic hook Lambda function. This Lambda function must call back to CodeDeploy and indicate success or failure. If the function fails, it aborts and reports a failure back to AWS CloudFormation. If the function succeeds, CodeDeploy proceeds to traffic shifting.
  - **PostTraffic:** After traffic shifting completes, CodeDeploy invokes the post-traffic hook Lambda function. This is similar to the pre-traffic hook, where the function must call back to CodeDeploy to report a success or failure. Use post-traffic hooks to run integration tests or other validation actions.



For more information, see [SAM Reference to Safe Deployments](#).

# Monitoring Serverless Applications

After you deploy your serverless application to the AWS Cloud, you need to verify that it's operating properly on an ongoing basis.

## Topics

- [Working with Logs \(p. 149\)](#)

## Working with Logs

To simplify troubleshooting, the AWS SAM CLI has a command called `sam logs` ([p. 173](#)). This command lets you fetch logs generated by your Lambda function from the command line.

### Note

The `sam logs` command works for all AWS Lambda functions, not just the ones you deploy using AWS SAM.

## Fetching Logs by AWS CloudFormation Stack

When your function is a part of an AWS CloudFormation stack, you can fetch logs by using the function's logical ID:

```
sam logs -n HelloWorldFunction --stack-name mystack
```

## Fetching Logs by Lambda Function Name

Or, you can fetch logs by using the function's name:

```
sam logs -n mystack-HelloWorldFunction-1FJ8PD
```

## Tailing Logs

Add the `--tail` option to wait for new logs and see them as they arrive. This is helpful during deployment or when you're troubleshooting a production issue.

```
sam logs -n HelloWorldFunction --stack-name mystack --tail
```

## Viewing Logs for a Specific Time Range

You can view logs for a specific time range by using the `-s` and `-e` options:

```
sam logs -n HelloWorldFunction --stack-name mystack -s '10min ago' -e '2min ago'
```

## Filtering Logs

Use the `--filter` option to quickly find logs that match terms, phrases, or values in your log events:

```
sam logs -n HelloWorldFunction --stack-name mystack --filter "error"
```

In the output, the AWS SAM CLI underlines all occurrences of the word "error" so you can easily locate the filter keyword within the log output.

## Error Highlighting

When your Lambda function crashes or times out, the AWS SAM CLI highlights the timeout message in red. This helps you easily locate specific executions that are timing out within a giant stream of log output.

## JSON Pretty Printing

If your log messages print JSON strings, the AWS SAM CLI automatically pretty prints the JSON to help you visually parse and understand the JSON.

# Publishing Serverless Applications Using the AWS SAM CLI

You can use the AWS SAM CLI to publish your application to the AWS Serverless Application Repository to make it available for others to find and deploy. To make an AWS SAM application public, you must create it, and all of the other Amazon or AWS resources it uses, in us-east-1 or us-east-2.

The application that you want to publish must be one that you've defined using AWS SAM. You also need to have tested it locally and/or in the AWS Cloud. The application's deployment package and AWS SAM template are the inputs to the following procedure steps.

The following instructions either create a new application, create a new version of an existing application, or update the metadata of an existing application. This depends on whether the application already exists in the AWS Serverless Application Repository, and whether any application metadata is changing. For more information about application metadata that's used to publish applications, see [AWS SAM Template Metadata Section Properties \(p. 154\)](#).

## Prerequisites

Before you publish an application to the AWS Serverless Application Repository, you need the following:

- A valid AWS account with an IAM user that has administrator permissions. See [Set Up an AWS Account](#).
- Version 1.16.77 or later of the AWS CLI installed. See [Installing the AWS Command Line Interface](#). If you have the AWS CLI installed, you can get the version by running the following command:

```
aws --version
```

- The AWS SAM CLI (command line interface) installed. See [Installing the AWS SAM CLI](#). You can determine whether the AWS SAM CLI is installed by running the following command:

```
sam --version
```

- A valid AWS Serverless Application Model (AWS SAM) template.
- Your application code and dependencies referenced by the AWS SAM template.
- A semantic version for your application (required to share your application publicly). This value can be as simple as 1.0.
- A URL that points to your application's source code.
- A README.md file. This file should describe how customers can use your application, and how to configure it before deploying it in their own AWS accounts.
- A LICENSE.txt file (required to share your application publicly).
- A valid Amazon S3 bucket policy that grants the service read permissions for artifacts uploaded to Amazon S3 when you packaged your application. To do this, follow these steps:
  1. Open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
  2. Choose the Amazon S3 bucket that you used to package your application.
  3. Choose the **Permissions** tab.
  4. Choose the **Bucket Policy** button.

5. Paste the following policy statement into the **Bucket policy editor**. Make sure to substitute your bucket name in the Resource property value.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "serverlessrepo.amazonaws.com"
      },
      "Action": "s3:GetObject",
      "Resource": "arn:aws:s3:::<your-bucket-name>/*"
    }
  ]
}
```

6. Choose the **Save** button.

## Publishing a New Application

### Step 1: Add a Metadata Section to the AWS SAM Template

First add a Metadata section to your AWS SAM template. Provide the application information to be published to the AWS Serverless Application Repository.

The following is an example Metadata section:

```
Metadata:
  AWS::ServerlessRepo::Application:
    Name: my-app
    Description: hello world
    Author: user1
    SpdxLicenseId: Apache-2.0
    LicenseUrl: LICENSE.txt
    ReadmeUrl: README.md
    Labels: ['tests']
    HomePageUrl: https://github.com/user1/my-app-project
    SemanticVersion: 0.0.1
    SourceCodeUrl: https://github.com/user1/my-app-project

Resources:
  HelloWorldFunction:
    Type: AWS::Lambda::Function
    Properties:
      ...
      CodeUri: source-code1
      ...
```

For more information about the properties of the Metadata section in the AWS SAM template, see [AWS SAM Template Metadata Section Properties](#) (p. 154).

### Step 2: Package the Application

Execute the following AWS SAM CLI command:

```
sam package \  
  --template-file template.yaml \  
  --output-template-file packaged.yaml \  
  --s3-bucket <your-bucket-name>
```

The command uploads the application artifacts to Amazon S3 and outputs a new template file called `packaged.yaml`. You use this file in the next step to publish the application to the AWS Serverless Application Repository. The `packaged.yaml` template file is similar to the original template file (`template.yaml`), but has a key difference—the `CodeUri`, `LicenseUrl`, and `ReadmeUrl` properties point to the Amazon S3 bucket and objects that contain the respective artifacts.

The following snippet from an example `packaged.yaml` template file shows the `CodeUri` property:

```
MySampleFunction:  
  Type: AWS::Serverless::Function  
  Properties:  
    CodeUri: s3://bucketname/fbd77a3647a4f47a352fcObjectGUID  
  
...
```

## Step 3: Publish the Application

Execute the following AWS SAM CLI command:

```
sam publish \  
  --template packaged.yaml \  
  --region us-east-1
```

The output of the `sam publish` command includes a link to the AWS Serverless Application Repository directly to your application. You can also go to the AWS Serverless Application Repository landing page directly and search for your application.

Your application is set to private by default, so it isn't visible to other AWS accounts. In order to share your application with others, you must either make it public or grant permission to a specific list of AWS accounts. For information on sharing your application by using the AWS CLI, see [Using Resource-based Policies for the AWS Serverless Application Repository](#). For information on sharing your application using the console, see [Sharing an Application Through the Console](#).

## Publishing a New Version of an Existing Application

After you've published an application to the AWS Serverless Application Repository, you might want to publish a new version of it. For example, you might have changed your Lambda function code or added a new component to your application architecture.

To update an application that you've previously published, you publish the application using the same process as above. You provide the same application name that you originally published it with, but with a new `SemanticVersion` value. You also provide the application name and `SemanticVersion` number in the `Metadata` section of the AWS SAM template file.

For example, if you published an application with the name `SampleApp` and `SemanticVersion 1.0.0`, to update that application, the AWS SAM template must have application name `SampleApp`, and the `SemanticVersion` can be `1.0.1` (or anything different from `1.0.0`).

## Additional Topics

- [AWS SAM Template Metadata Section Properties \(p. 154\)](#)

## AWS SAM Template Metadata Section Properties

`AWS::ServerlessRepo::Application` is a metadata key that you can use to specify application information that you want published to the AWS Serverless Application Repository.

### Note

AWS CloudFormation [intrinsic functions](#) aren't supported by the `AWS::ServerlessRepo::Application` metadata key.

## Properties

This table provides information about the properties of the `Metadata` section of the AWS SAM template. This section is required to publish applications to the AWS Serverless Application Repository using the AWS SAM CLI.

Property	Type	Required	Description
Name	String	TRUE	The name of the application.  Minimum length=1. Maximum length=140.  Pattern: "[a-zA-Z0-9\\-]+";
Description	String	TRUE	The description of the application.  Minimum length=1. Maximum length=256.
Author	String	TRUE	The name of the author publishing the application.  Minimum length=1. Maximum length=127.  Pattern "[a-z0-9]([a-z0-9] -?!-)*[a-z0-9]?\$";
SpdxLicenseId	String	FALSE	A valid license identifier. To view the list of valid license identifiers, see <a href="#">SPDX License List</a> on the <i>Software Package Data Exchange (SPDX)</i> website.
LicenseUrl	String	FALSE	The reference to a local license file, or an Amazon S3 link to a license file, that matches the <code>spdxLicenseId</code> value of your application.  An AWS SAM template file that hasn't been packaged using the <code>sam package</code> command can have a reference to a local file for this property. However, for an application to be published using the <code>sam publish</code> command, this property must be a reference to an Amazon S3 bucket.  Maximum size: 5 MB.  You must provide a value for this property in order to make your application public. Note that you cannot

Property	Type	Required	Description
			update this property after your application has been published. So, to add a license to an application, you must either delete it first, or publish a new application with a different name.
ReadmeUrl	String	FALSE	<p>The reference to a local readme file or an Amazon S3 link to the readme file that contains a more detailed description of the application and how it works.</p> <p>An AWS SAM template file that hasn't been packaged using the <code>sam package</code> command can have a reference to a local file for this property. However, to be published using the <code>sam publish</code> command, this property must be a reference to an Amazon S3 bucket.</p> <p>Maximum size: 5 MB.</p>
Labels	String	FALSE	<p>The labels that improve discovery of applications in search results.</p> <p>Minimum length=1. Maximum length=127. Maximum number of labels: 10.</p> <p>Pattern: <code>^[a-zA-Z0-9+\\-\\.\\@]+\$</code>;</p>
HomePageUrl	String	FALSE	A URL with more information about the application—for example, the location of your GitHub repository for the application.
SemanticVersion	String	FALSE	<p>The semantic version of the application. For the Semantic Versioning specification, see the <a href="#">Semantic Versioning</a> website.</p> <p>You must provide a value for this property in order to make your application public.</p>
SourceCodeUrl	String	FALSE	A link to a public repository for the source code of your application.

## Use Cases

This section lists the use cases for publishing applications, along with the `Metadata` properties that are processed for that use case. Properties that are *not* listed for a given use case are ignored.

- **Creating a new application** – A new application is created if there is no application in the AWS Serverless Application Repository with a matching name for an account.
  - Name
  - SpdxLicenseId
  - LicenseUrl
  - Description
  - Author
  - ReadmeUrl
  - Labels
  - HomePageUrl



- `SourceCodeUrl`
  - `SemanticVersion`
  - The content of the AWS SAM template (for example, any event sources, resources, and Lambda function code)
- 
- **Creating an application version** – An application version is created if there is already an application in the AWS Serverless Application Repository with a matching name for an account *and* the `SemanticVersion` is changing.
    - `Description`
    - `Author`
    - `ReadmeUrl`
    - `Labels`
    - `HomePageUrl`
    - `SourceCodeUrl`
    - `SemanticVersion`
    - The content of the AWS SAM template (for example, any event sources, resources, and Lambda function code)
- 
- **Updating an application** – An application is updated if there is already an application in the AWS Serverless Application Repository with a matching name for an account *and* the `SemanticVersion` is *not* changing.
    - `Description`
    - `Author`
    - `ReadmeUrl`
    - `Labels`
    - `HomePageUrl`

## Example

The following is an example `Metadata` section:

```
Metadata:
  AWS::ServerlessRepo::Application:
    Name: my-app
    Description: hello world
    Author: user1
    SpdxLicenseId: Apache-2.0
    LicenseUrl: LICENSE.txt
    ReadmeUrl: README.md
    Labels: ['tests']
    HomePageUrl: https://github.com/user1/my-app-project
    SemanticVersion: 0.0.1
    SourceCodeUrl: https://github.com/user1/my-app-project
```

# Example Serverless Applications

The following examples show you how to download, test, and deploy a number of additional serverless applications—including how to configure event sources and AWS resources.

## Topics

- [Process DynamoDB Events \(p. 157\)](#)
- [Process Amazon S3 Events \(p. 159\)](#)

## Process DynamoDB Events

With this example application, you build on what you learned in the overview and the Quick Start guide, and install another example application. This application consists of a Lambda function that's invoked by a DynamoDB table event source. The Lambda function is very simple—it logs data that was passed in through the event source message.

This exercise shows you how to mimic event source messages that are passed to Lambda functions when they're invoked.

## Before You Begin

Make sure that you've completed the required setup in the [Installing the AWS SAM CLI \(p. 3\)](#).

## Step 1: Initialize the Application

In this section, you download the application package, which consists of an AWS SAM template and application code.

### To initialize the application

1. Run the following command at an AWS SAM CLI command prompt.

```
sam init \  
--location gh:aws-samples/cookiecutter-aws-sam-dynamodb-python \  
--no-input
```

2. Review the contents of the directory that the command created (`dynamodb_event_reader/`):
  - `template.yaml` – Defines two AWS resources that the Read DynamoDB application needs: a Lambda function and a DynamoDB table. The template also defines mapping between the two resources.
  - `read_dynamodb_event/` directory – Contains the DynamoDB application code.

## Step 2: Test the Application Locally

For local testing, use the AWS SAM CLI to generate a sample DynamoDB event and invoke the Lambda function:

```
sam local generate-event dynamodb update | sam local invoke ReadDynamoDBEvent
```

The `generate-event` command creates a test event source message like the messages that are created when all components are deployed to the AWS Cloud. This event source message is piped to the Lambda function `ReadDynamoDBEvent`.

Verify that the expected messages are printed to the console, based on the source code in `app.py`.

## Step 3: Package the Application

After testing your application locally, you use the AWS SAM CLI to create a deployment package, which you use to deploy the application to the AWS Cloud.

### To create a Lambda deployment package

1. Create an S3 bucket in the location where you want to save the packaged code. If you want to use an existing S3 bucket, skip this step.

```
aws s3 mb s3://bucketname
```

2. Create the deployment package by running the following package CLI command at the command prompt.

```
sam package \  
  --template-file template.yaml \  
  --output-template-file packaged.yaml \  
  --s3-bucket bucketname
```

You specify the new template file, `packaged.yaml`, when you deploy the application in the next step.

## Step 4: Deploy the Application

Now that you've created the deployment package, you use it to deploy the application to the AWS Cloud. You then test the application.

### To deploy the serverless application to the AWS Cloud

- In the AWS SAM CLI, use the `deploy` CLI command to deploy all of the resources that you defined in the template.

```
sam deploy \  
  --template-file packaged.yaml \  
  --stack-name sam-app \  
  --capabilities CAPABILITY_IAM \  
  --region us-east-1
```

In the command, the `--capabilities` parameter allows AWS CloudFormation to create an IAM role.

AWS CloudFormation creates the AWS resources that are defined in the template. You can access the names of these resources in the AWS CloudFormation console.

### To test the serverless application in the AWS Cloud

1. Open the DynamoDB console.
2. Insert a record into the table that you just created.

3. Go to the **Metrics** tab of the table, and choose **View all CloudWatch metrics**. In the CloudWatch console, choose **Logs** to be able to view the log output.

## Process Amazon S3 Events

With this example application, you build on what you learned in the previous examples, and install a more complex application. This application consists of a Lambda function that's invoked by an Amazon S3 object upload event source. This exercise shows you how to access AWS resources and make AWS service calls through a Lambda function.

This sample serverless application processes object-creation events in Amazon S3. For each image that's uploaded to a bucket, Amazon S3 detects the object-created event and invokes a Lambda function. The Lambda function invokes Amazon Rekognition to detect text that's in the image. It then stores the results returned by Amazon Rekognition in a DynamoDB table.

### Note

With this example application, you perform steps in a slightly different order than in previous examples. The reason for this is that this example requires that AWS resources are created and IAM permissions are configured *before* you can test the Lambda function locally. We're going to leverage AWS CloudFormation to create the resources and configure the permissions for you. Otherwise, you would need to do this manually before you can test the Lambda function locally. Because this example is more complicated, be sure that you're familiar with installing the previous example applications before executing this one.

## Before You Begin

Make sure that you've completed the required setup in the [Installing the AWS SAM CLI \(p. 3\)](#).

## Step 1: Initialize the Application

In this section, you download the sample application, which consists of an AWS SAM template and application code.

### To initialize the application

1. Run the following command at an AWS SAM CLI command prompt.

```
sam init \  
--location https://github.com/aws-samples/cookiecutter-aws-sam-s3-rekognition-dynamodb-  
python \  
--no-input
```

2. Review the contents of the directory that the command created (`aws_sam_ocr/`):
  - `template.yaml` – Defines three AWS resources that the Amazon S3 application needs: a Lambda function, an Amazon S3 bucket, and a DynamoDB table. The template also defines the mappings and permissions between these resources.
  - `src/` directory – Contains the Amazon S3 application code.
  - `SampleEvent.json` – The sample event source, which is used for local testing.

## Step 2: Package the Application

Before you can test this application locally, you must use the AWS SAM CLI to create a deployment package, which you use to deploy the application to the AWS Cloud. This deployment creates the necessary AWS resources and permissions that are required to test the application locally.

### To create a Lambda deployment package

1. Create an S3 bucket in the location where you want to save the packaged code. If you want to use an existing S3 bucket, skip this step.

```
aws s3 mb s3://bucketname
```

2. Create the deployment package by running the following package CLI command at the command prompt.

```
sam package \  
  --template-file template.yaml \  
  --output-template-file packaged.yaml \  
  --s3-bucket bucketname
```

You specify the new template file, `packaged.yaml`, when you deploy the application in the next step.

## Step 3: Deploy the Application

Now that you've created the deployment package, you use it to deploy the application to the AWS Cloud. You then test the application by invoking it in the AWS Cloud.

### To deploy the serverless application to the AWS Cloud

- In the AWS SAM CLI, use the `deploy` command to deploy all of the resources that you defined in the template.

```
sam deploy \  
  --template-file packaged.yaml \  
  --stack-name aws-sam-ocr \  
  --capabilities CAPABILITY_IAM \  
  --region us-east-1
```

In the command, the `--capabilities` parameter allows AWS CloudFormation to create an IAM role.

AWS CloudFormation creates the AWS resources that are defined in the template. You can access the names of these resources in the AWS CloudFormation console.

### To test the serverless application in the AWS Cloud

1. Upload an image to the Amazon S3 bucket that you created for this sample application.
2. Open the DynamoDB console and find the table that was created. See the table for results returned by Amazon Rekognition.
3. Verify that the DynamoDB table contains new records that contain text that Amazon Rekognition found in the uploaded image.

## Step 4: Test the Application Locally

Before you can test the application locally, you must first retrieve the names of the AWS resources that were created by AWS CloudFormation.

- Retrieve the Amazon S3 key name and bucket name from AWS CloudFormation. Modify the `SampleEvent.json` file by replacing the values for the object key, bucket name, and bucket ARN.
- Retrieve the DynamoDB table name. This name is used for the following `sam local invoke` command.

Use the AWS SAM CLI to generate a sample Amazon S3 event and invoke the Lambda function:

```
TABLE_NAME=Table name obtained from AWS CloudFormation console sam local invoke --event  
SampleEvent.json
```

The `TABLE_NAME=` portion sets the DynamoDB table name. The `--event` parameter specifies the file that contains the test event message to pass to the Lambda function.

You can now verify that the expected DynamoDB records were created, based on the results returned by Amazon Rekognition.

# AWS SAM Reference

## AWS SAM Specification

The AWS SAM specification is an open-source specification under the Apache 2.0 license. The current version of the AWS SAM specification is available in the [AWS Serverless Application Model \(AWS SAM\) Specification \(p. 22\)](#).

AWS SAM templates are an extension of AWS CloudFormation templates. For the full reference for AWS CloudFormation templates, see [AWS CloudFormation Template Reference](#).

## AWS SAM CLI Command Reference

The **AWS SAM CLI** is a command line tool that operates on an AWS SAM template and application code. With the AWS SAM CLI, you can invoke Lambda functions locally, create a deployment package for your serverless application, deploy your serverless application to the AWS Cloud, and so on.

You can use the AWS SAM CLI commands to develop, test, and deploy your serverless applications to the AWS Cloud. The following are some examples of AWS SAM CLI commands:

- `sam init` – If you're a first-time AWS SAM CLI user, you can run the `sam init` command without any parameters to create a Hello World application. The command generates a preconfigured AWS SAM template and example application code in the language that you choose.
- `sam local invoke` and `sam local start-api` – Use these commands to test your application code locally, before deploying it to the AWS Cloud.
- `sam logs` – Use this command to fetch logs generated by your Lambda function. This can help you with testing and debugging your application after you've deployed it to the AWS Cloud.
- `sam package` – Use this command to bundle your application code and dependencies into a "deployment package". The deployment package is needed to upload your application to the AWS Cloud.
- `sam deploy` – Use this command to deploy your serverless application to the AWS Cloud. It creates the AWS resources and sets permissions and other configurations that are defined in the AWS SAM template.

## AWS SAM Policy Templates

AWS SAM allows you to choose from a list of policy templates to scope the permissions of your Lambda functions to the resources that are used by your application.

### Topics

- [AWS Serverless Application Model \(AWS SAM\) Specification \(p. 22\)](#)
- [AWS SAM CLI Command Reference \(p. 163\)](#)
- [AWS SAM Policy Templates \(p. 178\)](#)
- [Telemetry in the AWS SAM CLI \(p. 210\)](#)

# AWS SAM CLI Command Reference

This section is the reference for the AWS SAM CLI commands.

## Topics

- [sam build](#) (p. 163)
- [sam deploy](#) (p. 164)
- [sam init](#) (p. 166)
- [sam local generate-event](#) (p. 168)
- [sam local invoke](#) (p. 169)
- [sam local start-api](#) (p. 170)
- [sam local start-lambda](#) (p. 172)
- [sam logs](#) (p. 173)
- [sam package](#) (p. 174)
- [sam publish](#) (p. 175)
- [sam validate](#) (p. 176)

## sam build

Use this command to build your Lambda source code and generate deployment artifacts that target Lambda's execution environment. By doing this, the functions that you build locally run in a similar environment in the AWS Cloud.

The `sam build` command iterates through the functions in your application, looks for a manifest file (such as `requirements.txt`) that contains the dependencies, and automatically creates deployment artifacts that you can deploy to Lambda using the `sam package` and `sam deploy` commands. You can also use `sam build` in combination with other commands like `sam local invoke` to test your application locally.

To use this command, update your AWS SAM template to specify the path to your function's source code in the resource's `Code` or `CodeUri` property.

### Usage:

```
sam build [OPTIONS]
```

### Examples:

```
To build on your workstation, run this command in folder containing  
SAM template. Built artifacts will be written to .aws-sam/build folder  
$ sam build
```

```
To build inside a AWS Lambda like Docker container  
$ sam build --use-container
```

```
To build & run your functions locally  
$ sam build && sam local invoke
```

```
To build and package for deployment  
$ sam build && sam package --s3-bucket <bucketname>
```

### Options:



Option	Description
-b, --build-dir DIRECTORY	The path to a folder where the built artifacts are stored. This directory and all of its content will be removed with this option.
-s, --base-dir DIRECTORY	Resolves relative paths to the function's source code with respect to this folder. Use this if the AWS SAM template and your source code aren't in the same enclosing folder. By default, relative paths are resolved with respect to the template's location.
-u, --use-container	If your functions depend on packages that have natively compiled dependencies, use this flag to build your function inside an AWS Lambda-like Docker container.
-m, --manifest PATH	The path to a custom dependency manifest (ex: package.json) to use instead of the default one.
-t, --template PATH	The AWS SAM template file [default: template.[yaml yml]].
--parameter-overrides	Optional. A string that contains AWS CloudFormation parameter overrides, encoded as key-value pairs. Use the same format as the AWS CLI—for example, 'ParameterKey=KeyPairName, ParameterValue=MyKey ParameterKey=InstanceType, ParameterValue=t1.micro'.
--skip-pull-image	Specifies whether the command should skip pulling down the latest Docker image for Lambda runtime.
--docker-network TEXT	Specifies the name or id of an existing Docker network to Lambda Docker containers should connect to, along with the default bridge network. If not specified, the Lambda containers will only connect to the default bridge Docker network.
--profile TEXT	Selects a specific profile from your credential file to get AWS credentials.
--region TEXT	Sets the AWS Region of the service (for example, us-east-1).
--debug	Turns on debug logging to print debug message generated by the AWS SAM CLI.
--help	Shows this message and exits.

## sam deploy

Deploys an AWS SAM application.

This command now comes with a guided interactive mode, which you can enable by specifying the `--guided` parameter. The interactive mode walks you through the parameters required for deployment, provides default options, and saves these options in a configuration file in your project folder. You can execute subsequent deployments of your application by simply executing `sam deploy` and the needed parameters will be retrieved from the AWS SAM CLI configuration file.

Deploying Lambda functions through AWS CloudFormation requires an Amazon S3 bucket for the Lambda deployment package. AWS SAM CLI now creates and manages this Amazon S3 bucket for you.

### Usage:

```
sam deploy [OPTIONS] [ARGS]...
```

**Options:**

Option	Description
-g, --guided	Specify this flag to allow AWS SAM to guide you through the deployment using guided prompts.  For more information about settings that optionally get stored when specifying this parameter, see <a href="#">AWS SAM CLI Config (p. 176)</a>
--template-file PATH	The path where your AWS SAM template is located. Default: template.[yaml yml].
--stack-name TEXT	The name of the AWS CloudFormation stack you're deploying to. If you specify an existing stack, the command updates the stack. If you specify a new stack, the command creates it. Required.
--s3-bucket TEXT	The name of the Amazon S3 bucket where artifacts that are referenced in your template. This is required the deployments of templates sized greater than 51,200 bytes.
--s3-prefix TEXT	Prefix added to the artifacts name that are uploaded to the Amazon S3 bucket. The prefix name is a path name (folder name) for the Amazon S3 bucket.
--capabilities LIST	A list of capabilities that you must specify to allow AWS CloudFormation to create certain stacks. Some stack templates might include resources that can affect permissions in your AWS account, for example, by creating new AWS Identity and Access Management (IAM) users. For those stacks, you must explicitly acknowledge their capabilities by specifying this parameter. The only valid values are <code>CAPABILITY_IAM</code> and <code>CAPABILITY_NAMED_IAM</code> . If you have IAM resources, you can specify either capability. If you have IAM resources with custom names, you must specify <code>CAPABILITY_NAMED_IAM</code> . If you don't specify this parameter, this action returns an <code>InsufficientCapabilities</code> error.
--region TEXT	The AWS Region to deploy to (for example, us-east-1).
--profile TEXT	Select a specific profile from your credential file to get AWS credentials.
--kms-key-id TEXT	The ID of an AWS KMS key used to encrypt artifacts that are at rest in the Amazon S3 bucket.
--force-upload	Override existing files in the Amazon S3 bucket. Specify this flag to upload artifacts even if they match existing artifacts in the Amazon S3 bucket.
--no-execute-changeset	Indicates whether to execute the change set. Specify this flag if you want to view your stack changes before executing the change set. This command creates an AWS CloudFormation change set and then exits without executing the change set. If you want to execute the changeset, the stack changes can be made by running the same command without this flag.
--role-arn TEXT	The Amazon Resource Name (ARN) of an AWS Identity and Access Management (IAM) role that AWS CloudFormation assumes when executing the change set.

Option	Description
<code>--fail-on-empty-changeset</code>   <code>--no-fail-on-empty-changeset</code>	Specify whether to return a non-zero exit code if there are no changes to be made to the stack. The default behavior is to return a non-zero exit code.
<code>--confirm-changeset</code>	Prompt to confirm before deploying the computed changeset.
<code>--use-json</code>	Output JSON for the AWS CloudFormation template. YAML is used by default.
<code>--metadata</code>	A map of metadata to attach to all artifacts that are referenced in your template. Optional.
<code>--notification-arns LIST</code>	Amazon Simple Notification Service topic Amazon Resource Names (ARNs) that AWS CloudFormation associates with the stack.
<code>--tags</code>	A list of tags to associate with the stack that is created or updated. AWS CloudFormation also propagates these tags to resources in the stack if the resource supports it.
<code>--parameter-overrides</code>	A string that contains AWS CloudFormation parameter overrides encoded as key=value pairs. Use the same format as the AWS CLI. For example, <code>ParameterKey=KeyPairName,ParameterValue=MyKey</code> <code>ParameterKey=InstanceType,ParameterValue=t1.micro</code> .
<code>--debug</code>	Turns on debug logging.
<code>--help</code>	Shows this message and exits.

## sam init

Initializes a serverless application with an AWS SAM template. The template provides a folder structure for your Lambda functions, and is connected to an event source such as APIs, S3 buckets, or DynamoDB tables. This application includes everything you need to get started and to eventually extend it into a production-scale application.

### Usage:

```
sam init [OPTIONS]
```

### Note

With AWS SAM version 0.30.0 or later, you can initialize your application using one of two modes: 1) Interactive workflow, or 2) Providing all required parameters.

- *Interactive Workflow:* Through the interactive initialize workflow you can input either 1) your project name, preferred runtime, and template file, or 2) the location of a custom template.
- *Providing Parameters:* Provide all required parameters.

If you provide a subset of required parameters, you will be prompted for the additional required information.

### Examples:

```
Initializes a new SAM project with required parameters passed as parameters
```

```
sam init --runtime python3.7 --dependency-manager pip --app-template hello-world --name
sam-app

Initializes a new SAM project using custom template in a Git/Mercurial repository

# gh being expanded to github url
sam init --location gh:aws-samples/cookiecutter-aws-sam-python

sam init --location git+ssh://git@github.com/aws-samples/cookiecutter-aws-sam-python.git

sam init --location hg+ssh://hg@bitbucket.org/repo/template-name

# Initializes a new SAM project using custom template in a Zipfile
sam init --location /path/to/template.zip

sam init --location https://example.com/path/to/template.zip

# Initializes a new SAM project using custom template in a local path
sam init --location /path/to/template/folder
```

#### Options:

Option	Description
--no-interactive	Disable interactive prompting for init parameters, and fail if any required values are missing.
-l, --location TEXT	<p>The template or application location (Git, Mercurial, HTTP/HTTPS, ZIP, path).</p> <p>This parameter is required if --no-interactive is specified and --runtime, --name, and --app-template are not provided.</p> <p>For Git repositories, you must use location of the root of the repository.</p>
-r, --runtime [python2.7   nodejs6.10   ruby2.5   java8   python3.7   nodejs8.10   dotnetcore2.0   nodejs10.x   dotnetcore2.1   dotnetcore1.0   python3.6   go1.x]	<p>The Lambda runtime of your application.</p> <p>This parameter is required if --no-interactive is specified and --location is not provided.</p>
-d, --dependency-manager [gradle   mod   maven   bundler   npm   cli-package   pip]	Dependency manager of your Lambda runtime
-o, --output-dir PATH	The location where the initialized application is output.
-n, --name TEXT	<p>The name of your project to be generated as a folder.</p> <p>This parameter is required if --no-interactive is specified and --location is not provided.</p>
--app-template TEXT	Identifier of the managed application template you want to use. If not sure, call 'sam init' without options for an interactive workflow.

Option	Description
	This parameter is required if <code>--no-interactive</code> is specified and <code>--location</code> is not provided.  This parameter is only available in SAM CLI version 0.30.0 or later. Specifying this parameter with an earlier version will result in an error.
<code>--no-input</code>	Disables Cookiecutter prompting and accept default values that are defined in the template configuration.
<code>--debug</code>	Turns on debug logging.
<code>-h, --help</code>	Shows this message and exits.

## sam local generate-event

Generates sample payloads from different event sources, such as Amazon S3, Amazon API Gateway, and Amazon SNS. These payloads contain the information that the event sources send to your Lambda functions.

### Usage:

```
sam local generate-event [OPTIONS] COMMAND [ARGS]...
```

### Examples:

```
Generate the event that S3 sends to your Lambda function when a new object is uploaded
sam local generate-event s3 [put/delete]

# You can even customize the event by adding parameter flags. To find which flags apply to
your command,
run:

sam local generate-event s3 [put/delete] --help

# Then you can add in those flags that you wish to customize using

sam local generate-event s3 [put/delete] --bucket <bucket> --key <key>

# After you generate a sample event, you can use it to test your Lambda function locally
sam local generate-event s3 [put/delete] --bucket <bucket> --key <key> | sam local invoke
<function logical id>
```

### Options:

Option	Description
<code>--help</code>	Shows this message and exits.

### Commands:

- alexa-skills-kit
- alexa-smart-home
- apigateway
- batch

- cloudformation
- cloudfront
- cloudwatch
- codecommit
- codepipeline
- cognito
- config
- dynamodb
- kinesis
- lex
- rekognition
- s3
- ses
- sns
- sqs
- stepfunctions

## sam local invoke

Invokes a local Lambda function once and quits after invocation completes.

This is useful for developing serverless functions that handle asynchronous events (such as Amazon S3 or Amazon Kinesis events). It can also be useful if you want to compose a script of test cases. The event body can be passed in either by `stdin` (default), or by using the `--event` parameter. The runtime output (logs etc) is output to `stderr`, and the Lambda function result is output to `stdout`.

### Usage:

```
sam local invoke [OPTIONS] [FUNCTION_IDENTIFIER]
```

### Options:

Option	Description
<code>-e, --event PATH</code>	The JSON file that contains event data that's passed to the Lambda function when it's invoked. If you don't specify this option, no event is assumed. To input JSON from <code>stdin</code> you must pass in the value <code>'-'</code> .
<code>--no-event</code>	Invokes the function with an empty event.
<code>-t, --template PATH</code>	The AWS SAM template file [default: <code>template.[yaml yml]</code> ].
<code>-n, --env-vars PATH</code>	The JSON file that contains values for the Lambda function's environment variables. For more information about environment variables files, see <a href="#">Environment Variable File (p. 133)</a> .
<code>--parameter-overrides</code>	Optional. A string that contains AWS CloudFormation parameter overrides encoded as key-value pairs. Use the same format as the AWS CLI—for example, <code>'ParameterKey=KeyPairName, ParameterValue=MyKey ParameterKey=InstanceType,ParameterValue=t1.micro'</code> .
<code>-d, --debug-port TEXT</code>	When specified, starts the Lambda function container in debug mode and exposes this port on the local host.

Option	Description
--debugger-path TEXT	The host path to a debugger that will be mounted into the Lambda container.
--debug-args TEXT	Additional arguments to be passed to the debugger.
-v, --docker-volume-basedir TEXT	The location of the base directory where the AWS SAM file exists. If Docker is running on a remote machine, you must mount the path where the AWS SAM file exists on the Docker machine, and modify this value to match the remote machine.
--docker-network TEXT	The name or ID of an existing Docker network that Lambda Docker containers should connect to, along with the default bridge network. If this isn't specified, the Lambda containers only connect to the default bridge Docker network.
-l, --log-file TEXT	The log file to send runtime logs to.
--layer-cache-basedir DIRECTORY	Specifies the location basedir where the layers your template uses are downloaded to.
--skip-pull-image	Specifies whether the CLI should skip pulling down the latest Docker image for the Lambda runtime.
--force-image-build	Specifies whether the CLI should rebuild the image used for invoking functions with layers.
--profile TEXT	The AWS credentials profile to use.
--region TEXT	Sets the AWS Region of the service (for example, us-east-1).
--debug	Turns on debug logging.
--help	Shows this message and exits.

## sam local start-api

Allows you to run your serverless application locally for quick development and testing. When you run this command in a directory that contains your serverless functions and your AWS SAM template, it creates a local HTTP server that hosts all of your functions.

When it's accessed (through a browser, CLI, and so on), it starts a Docker container locally to invoke the function. It reads the CodeUri property of the AWS::Serverless::Function resource to find the path in your file system that contains the Lambda function code. This could be the project's root directory for interpreted languages like Node.js and Python, or a build directory that stores your compiled artifacts or a Java Archive (JAR) file.

If you're using an interpreted language, local changes are available immediately in the Docker container on every invoke. For more compiled languages or projects that require complex packing support, we recommend that you run your own building solution, and point AWS SAM to the directory or file that contains the build artifacts.

### Usage:

```
sam local start-api [OPTIONS]
```

### Options:

Option	Description
--host TEXT	The local hostname or IP address to bind to (default: '127.0.0.1').
-p, --port INTEGER	The local port number to listen on (default: '3000').
-s, --static-dir TEXT	Any static asset (for example, CSS/JavaScript/HTML) files located in this directory are presented at /.
-t, --template PATH	The AWS SAM template file [default: template.[yaml yml]].
-n, --env-vars PATH	The JSON file that contains values for the Lambda function's environment variables.
--parameter-overrides	Optional. A string that contains AWS CloudFormation parameter overrides encoded as key-value pairs. Use the same format as the AWS CLI—for example, 'ParameterKey=KeyPairName, ParameterValue=MyKey ParameterKey=InstanceType,ParameterValue=t1.micro'.
-d, --debug-port TEXT	When specified, starts the Lambda function container in debug mode and exposes this port on the local host.
--debugger-path TEXT	The host path to a debugger that will be mounted into the Lambda container.
--debug-args TEXT	Additional arguments to be passed to the debugger.
-v, --docker-volume-basedir TEXT	The location of the base directory where the AWS SAM file exists. If Docker is running on a remote machine, you must mount the path where the AWS SAM file exists on the Docker machine, and modify this value to match the remote machine.
--docker-network TEXT	The name or ID of an existing Docker network that the Lambda Docker containers should connect to, along with the default bridge network. If this isn't specified, the Lambda containers only connect to the default bridge Docker network.
-l, --log-file TEXT	The log file to send runtime logs to.
--layer-cache-basedir DIRECTORY	Specifies the location basedir where the Layers your template uses are downloaded to.
--skip-pull-image	Specifies whether the CLI should skip pulling down the latest Docker image for the Lambda runtime.
--force-image-build	Specifies whether CLI should rebuild the image used for invoking functions with layers.
--profile TEXT	The AWS credentials profile to use.
--region TEXT	Sets the AWS Region of the service (for example, us-east-1).
--debug	Turns on debug logging.
--help	Shows this message and exits.



## sam local start-lambda

Enables you to programmatically invoke your Lambda function locally by using the AWS CLI or SDKs. This command starts a local endpoint that emulates AWS Lambda. You can run your automated tests against this local Lambda endpoint. When you send an invoke to this endpoint using the AWS CLI or SDK, it locally executes the Lambda function that's specified in the request.

### Usage:

```
sam local start-lambda [OPTIONS]
```

### Examples:

```
# SETUP
# -----
# Start the local Lambda endpoint by running this command in the directory that contains
# your AWS SAM template.

sam local start-lambda

# USING AWS CLI
# -----
# Then, you can invoke your Lambda function locally using the AWS CLI

aws lambda invoke --function-name "HelloWorldFunction" --endpoint-url
"http://127.0.0.1:3001" --no-verify-ssl out.txt

# USING AWS SDK
# -----
# You can also use the AWS SDK in your automated tests to invoke your functions
# programmatically.
# Here is a Python example:
#
#     self.lambda_client = boto3.client('lambda',
#                                       endpoint_url="http://127.0.0.1:3001",
#                                       use_ssl=False,
#                                       verify=False,
#                                       config=Config(signature_version=UNSIGNED,
#                                                     read_timeout=0,
#                                                     retries={'max_attempts': 0}))
#     self.lambda_client.invoke(FunctionName="HelloWorldFunction")
```

### Options:

Option	Description
--host TEXT	The local hostname or IP address to bind to (default: '127.0.0.1').
-p, --port INTEGER	The local port number to listen on (default: '3001').
-t, --template PATH	The AWS SAM template file [default: template.[yaml yml]].
-n, --env-vars PATH	The JSON file that contains values for the Lambda function's environment variables.
--parameter-overrides	Optional. A string that contains AWS CloudFormation parameter overrides encoded as key-value pairs. Use the same format as the AWS CLI—for example, 'ParameterKey=KeyPairName, ParameterValue=MyKey ParameterKey=InstanceType,ParameterValue=t1.micro'.

Option	Description
-d, --debug-port TEXT	When specified, starts the Lambda function container in debug mode, and exposes this port on the local host.
--debugger-path TEXT	The host path to a debugger to be mounted into the Lambda container.
--debug-args TEXT	Additional arguments to be passed to the debugger.
-v, --docker-volume-basedir TEXT	The location of the base directory where the AWS SAM file exists. If Docker is running on a remote machine, you must mount the path where the AWS SAM file exists on the Docker machine, and modify this value to match the remote machine.
--docker-network TEXT	The name or ID of an existing Docker network that Lambda Docker containers should connect to, along with the default bridge network. If this is specified, the Lambda containers only connect to the default bridge Docker network.
-l, --log-file TEXT	The log file to send runtime logs to.
--layer-cache-basedir DIRECTORY	Specifies the location basedir where the layers your template uses are downloaded to.
--skip-pull-image	Specifies whether the CLI should skip pulling down the latest Docker image for the Lambda runtime.
--force-image-build	Specify whether the CLI should rebuild the image used for invoking functions with layers.
--profile TEXT	The AWS credentials profile to use.
--region TEXT	Sets the AWS Region of the service (for example, us-east-1).
--debug	Turns on debug logging.
--help	Shows this message and exits.

## sam logs

Fetches logs that are generated by your Lambda function.

When your functions are a part of an AWS CloudFormation stack, you can fetch logs by using the function's logical ID when you specify the stack name.

### Usage:

```
sam logs [OPTIONS]
```

### Examples:

```
sam logs -n HelloWorldFunction --stack-name mystack

# Or, you can fetch logs using the function's name.
sam logs -n mystack-HelloWorldFunction-1FJ8PD36GML2Q

# You can view logs for a specific time range using the -s (--start-time) and -e (--end-time) options
sam logs -n HelloWorldFunction --stack-name mystack -s '10min ago' -e '2min ago'
```

```
# You can also add the --tail option to wait for new logs and see them as they arrive.
sam logs -n HelloWorldFunction --stack-name mystack --tail

# Use the --filter option to quickly find logs that match terms, phrases or values in your
log events.
sam logs -n HelloWorldFunction --stack-name mystack --filter "error"
```

#### Options:

Option	Description
-n, --name TEXT	The name of your Lambda function. If this function is part of an AWS CloudFormation stack, this can be the logical ID of the function resource in the AWS CloudFormation/AWS SAM template. [required]
--stack-name TEXT	The name of the AWS CloudFormation stack that the function is a part of.
--filter TEXT	Lets you specify an expression to quickly find logs that match terms, phrases, or values in your log events. This can be a simple keyword (for example, "error") or a pattern that's supported by Amazon CloudWatch Logs. For the syntax, see the <a href="#">Amazon CloudWatch Logs documentation</a> .
-s, --start-time TEXT	Fetches logs starting at this time. The time can be relative values like '5mins ago', 'yesterday', or a formatted timestamp like '2018-01-01 10:10:10'. It defaults to '10mins ago'.
-e, --end-time TEXT	Fetches logs up to this time. The time can be relative values like '5mins ago', 'tomorrow', or a formatted timestamp like '2018-01-01 10:10:10'.
-t, --tail	Tails the log output. This ignores the end time argument and continues to fetch logs as they become available.
--profile TEXT	Selects a specific profile from your credential file to get AWS credentials.
--region TEXT	Sets the AWS Region of the service (for example, us-east-1).
--debug	Turns on debug logging to print debug messages that are generated by the AWS SAM CLI.
--help	Shows this message and exits.

## sam package

Packages an AWS SAM application. It creates a ZIP file of your code and dependencies, and uploads it to Amazon S3. It then returns a copy of your AWS SAM template, replacing references to local artifacts with the Amazon S3 location where the command uploaded the artifacts.

#### Note

[sam deploy \(p. 164\)](#) now implicitly performs the functionality of `sam package`. You can use the [sam deploy \(p. 164\)](#) command directly to package and deploy your application.

#### Usage:

```
sam package [OPTIONS] [ARGS]...
```

#### Options:

Option	Description
--template-file PATH	The path where your AWS SAM template is located. Default: template.[yaml yml].
--s3-bucket TEXT	The name of the S3 bucket where this command uploads the artifacts that are referenced in your template. Required.
--s3-prefix TEXT	Prefix added to the artifacts name that are uploaded to the Amazon S3 bucket. The prefix name is a path name (folder name) for the Amazon S3 bucket.
--kms-key-id TEXT	The ID of an AWS KMS key used to encrypt artifacts that are at rest in the Amazon S3 bucket.
--output-template-file PATH	The path to the file where the command writes the packaged template. If you don't specify a path, the command writes the template to the standard output.
--use-json	Output JSON for the AWS CloudFormation template. YAML is used by default.
--force-upload	Override existing files in the Amazon S3 bucket. Specify this flag to upload artifacts even if they match existing artifacts in the Amazon S3 bucket.
--metadata	A map of metadata to attach to all artifacts that are referenced in your template. Optional.
--profile TEXT	Select a specific profile from your credential file to get AWS credentials.
--region TEXT	Sets the AWS Region of the service (for example, us-east-1).
--debug	Turns on debug logging.
--help	Shows this message and exits.

### Note

If the AWS SAM template contains a `Metadata` section for `ServerlessRepo`, and the `LicenseUrl` or `ReadmeUrl` properties contain references to local files, you must update AWS CLI to version 1.16.77 or later. For more information about the `Metadata` section of AWS SAM templates and publishing applications with AWS SAM CLI, see [Publishing Serverless Applications Using the AWS SAM CLI \(p. 151\)](#).

## sam publish

Publish an AWS SAM application to the AWS Serverless Application Repository. This command takes a packaged AWS SAM template and publishes the application to the specified region.

This command expects the AWS SAM template to include a `Metadata` containing application metadata required for publishing. Furthermore, these properties must include references to Amazon S3 buckets for `LicenseUrl` and `ReadmeUrl` values, and not references to local files. For more details about the `Metadata` section of the AWS SAM template, see [Publishing Serverless Applications Using the AWS SAM CLI \(p. 151\)](#).

This command creates the application as private by default, so you must share the application before other AWS accounts are allowed to view and deploy the application. For more information on sharing applications see [Using Resource-Based Policies for the AWS Serverless Application Repository](#).

**Usage:**

```
sam publish [OPTIONS]
```

**Examples:**

```
# To publish an application  
sam publish --template packaged.yaml --region us-east-1
```

**Options:**

Option	Description
-t, --template PATH	AWS SAM template file [default: template.[yaml yml]].
--semantic-version TEXT	Optional. The semantic version of the application provided by this parameter overrides SemanticVersion in the Metadata section of the template file. <a href="https://semver.org/">https://semver.org/</a>
--profile TEXT	Select a specific profile from your credential file to get AWS credentials.
--region TEXT	Sets the AWS Region of the service (for example, us-east-1).
--debug	Turns on debug logging.
--help	Shows this message and exits.

## sam validate

Validates an AWS SAM template.

**Usage:**

```
sam validate [OPTIONS]
```

**Options:**

Option	Description
-t, --template PATH	The AWS SAM template file [default: template.[yaml yml]].
--profile TEXT	Selects a specific profile from your credential file to get AWS credentials.
--region TEXT	Sets the AWS Region of the service (for example, us-east-1).
--debug	Turns on debug logging.
--help	Shows this message and exits.

## AWS SAM CLI Config

The AWS SAM CLI now supports a project-level configuration file that stores default parameters for AWS SAM commands. This configuration file stores the parameters to use for command executions. If the

required parameters are available in the configuration file, you can run commands without providing all parameters each execution.

For example, when executing the `sam deploy --guided` command, AWS SAM CLI automatically adds the required parameters into the configuration file. You can subsequently execute `sam deploy` with no parameters, and the values will be retrieved from the configuration file.

## Config Basics

The configuration file for a project should be created as `samconfig.toml` in your project's root directory. If you run a command like `sam deploy --guided` this file will be created for you. These files are written in TOML format: <https://github.com/toml-lang/toml>.

## Example

Here is an example AWS SAM CLI config file that is created when using the `sam deploy --guided` command:

```
version=0.1
[default.deploy.parameters]
region = "us-west-2"
stack_name = "my-app-stack"
capabilities = "CAPABILITY_IAM"
s3_bucket = "my-source-bucket"
```

## Options

### region

Set the region you want to deploy to. If you omit region, then the AWS SAM CLI will try to resolve the region through AWS profiles and then environment variables.

### stack\_name

The name of the AWS CloudFormation stack you're deploying to.

### capabilities

A list of capabilities that you must specify to allow AWS CloudFormation to create certain stacks. Some stack templates might include resources that can affect permissions in your AWS account, for example, by creating new AWS Identity and Access Management (IAM) users. For those stacks, you must explicitly acknowledge their capabilities by specifying this parameter.

The only valid values are `CAPABILITY_IAM` and `CAPABILITY_NAMED_IAM`. If you have IAM resources, you can specify either capability. If you have IAM resources with custom names, you must specify `CAPABILITY_NAMED_IAM`. If you don't specify this parameter, this action returns an `InsufficientCapabilities` error.

Allowed values: `CAPABILITY_IAM`, `CAPABILITY_NAMED_IAM`

Default value: `CAPABILITY_IAM`

### s3\_bucket

The S3 bucket name used for deployment artifacts for the `sam deploy` command. You may use your own Amazon S3 bucket by providing your own value here.

# AWS SAM Policy Templates

AWS SAM allows you to choose from a list of policy templates to scope the permissions of your Lambda functions to the resources that are used by your application.

AWS SAM applications in the AWS Serverless Application Repository that use policy templates don't require any special customer acknowledgments to deploy the application from the AWS Serverless Application Repository.

If you want to request a new policy template to be added, do the following:

1. Submit a pull request against the `policy_templates.json` source file in the `develop` branch of the AWS SAM GitHub project. You can find the source file in [policy\\_templates.json](#) on the GitHub website.
2. Submit an issue in the AWS SAM GitHub project that includes the reasons for your pull request and a link to the request. Use this link to submit a new issue: [AWS Serverless Application Model: Issues](#).

## Examples

There are two AWS SAM template examples in this section: one with a policy template that includes placeholder values, and one that doesn't include placeholder values.

### Example 1: Policy Template with Placeholder Values

The following example shows that the [SQSPollerPolicy \(p. 182\)](#) policy template expects a `QueueName` as a resource. The AWS SAM template retrieves the name of the "MyQueue" Amazon SQS queue, which you can create in the same application or requested as a parameter to the application.

```
MyFunction:
  Type: 'AWS::Serverless::Function'
  Properties:
    CodeUri: ${codeuri}
    Handler: hello.handler
    Runtime: python2.7
    Policies:
      - SQSPollerPolicy:
        QueueName:
          !GetAtt MyQueue.QueueName
```

### Example 2: Policy Template with No Placeholder Values

The following example contains the [CloudWatchPutMetricPolicy \(p. 183\)](#) policy template, which has no placeholder values.

```
MyFunction:
  Type: 'AWS::Serverless::Function'
  Properties:
    CodeUri: ${codeuri}
    Handler: hello.handler
    Runtime: python2.7
    Policies:
      - CloudWatchPutMetricPolicy: {}
```

## Policy Template Table

The following is a table of the available policy templates.

Policy Template	Description		
<a href="#">SQSPollerPolicy (p. 182)</a>	Gives permission to poll an Amazon SQS Queue.		
<a href="#">LambdaInvokePolicy (p. 182)</a>	Gives permission to invoke a Lambda function, alias, or version.		
<a href="#">CloudWatchDescribeAlarmHistoryPolicy (p. 183)</a>	Gives permission to describe CloudWatch alarm history.		
<a href="#">CloudWatchPutMetricDataPolicy (p. 183)</a>	Gives permission to put metrics to CloudWatch.		
<a href="#">EC2DescribePolicy (p. 183)</a>	Gives permission to describe Amazon EC2 instances.		
<a href="#">DynamoDBCrudPolicy (p. 184)</a>	Gives create, read, update, and delete permissions to a DynamoDB table.		
<a href="#">DynamoDBReadPolicy (p. 184)</a>	Gives read-only permission to a DynamoDB table.		
<a href="#">DynamoDBReconfigurePolicy (p. 185)</a>	Gives permission to reconfigure a DynamoDB table.		
<a href="#">SESSendBouncePolicy (p. 185)</a>	Gives SendBounce permission to an Amazon SES identity.		
<a href="#">ElasticsearchHttpPostPolicy (p. 186)</a>	Gives POST permission to Amazon Elasticsearch Service.		
<a href="#">S3ReadPolicy (p. 187)</a>	Gives read-only permission to objects in an Amazon S3 bucket.		
<a href="#">S3CrudPolicy (p. 187)</a>	Gives create, read, update, and delete permission to objects in an Amazon S3 bucket.		
<a href="#">AMIDescribePolicy (p. 188)</a>	Gives permission to describe Amazon Machine Images (AMIs).		
<a href="#">CloudFormationDescribeStackPolicy (p. 188)</a>	Gives permission to describe AWS CloudFormation stacks.		
<a href="#">RekognitionDetectFacesPolicy (p. 189)</a>	Gives permission to detect faces, labels, and text.		
<a href="#">RekognitionNoDataComparePolicy (p. 189)</a>	Gives permission to compare and detect faces and labels.		
<a href="#">RekognitionReadFacesPolicy (p. 190)</a>	Gives permission to list and search faces.		
<a href="#">RekognitionWriteOccurrencePolicy (p. 190)</a>	Gives permission to create collection and index faces.		



Policy Template	Description		
<a href="#">SQSSendMessagePolicy (p. 190)</a>	Gives permission to send message to an Amazon SQS queue.		
<a href="#">SNSPublishMessagePolicy (p. 190)</a>	Gives permission to publish a message to an Amazon SNS topic.		
<a href="#">VPCAccessPolicy (p. 191)</a>	Gives access to create, delete, describe, and detach Elastic Network Interfaces.		
<a href="#">DynamoDBStreamFullPolicy (p. 191)</a>	Gives permission to describe and read DynamoDB streams and records.		
<a href="#">KinesisStreamReadOnlyPolicy (p. 192)</a>	Gives permission to list and read an Amazon Kinesis stream.		
<a href="#">SESCrudPolicy (p. 192)</a>	Gives permission to send email and verify identity.		
<a href="#">SNSCrudPolicy (p. 192)</a>	Gives permission to create, publish, and subscribe to Amazon SNS topics.		
<a href="#">KinesisCrudPolicy (p. 193)</a>	Gives permission to create, publish, and delete an Amazon Kinesis stream.		
<a href="#">KMSTDecryptPolicy (p. 194)</a>	Gives permission to decrypt with an AWS KMS key.		
<a href="#">KMSEncryptPolicy (p. 195)</a>	Gives permission to encrypt with an AWS KMS key.		
<a href="#">PollyFullAccessPolicy (p. 195)</a>	Gives full access permission to Amazon Polly lexicon resources.		
<a href="#">S3FullAccessPolicy (p. 196)</a>	Gives full access permission to objects in an Amazon S3 bucket.		
<a href="#">CodePipelineLambdaFullPolicy (p. 197)</a>	Gives permission for a Lambda function invoked by CodePipeline to report the status of the job.		
<a href="#">ServerlessRepoReadOnlyPolicy (p. 197)</a>	Gives permission to create and list applications in the AWS Serverless Application Repository service.		
<a href="#">EC2CopyImagePolicy (p. 197)</a>	Gives permission to copy Amazon EC2 images.		
<a href="#">AWSSecretsManagerRotatePolicy (p. 198)</a>	Gives permission to rotate a secret in AWS Secrets Manager.		
<a href="#">AWSSecretsManagerGetPolicy (p. 199)</a>	Gives permission to get secret value for the specified AWS Secrets Manager secret.		
<a href="#">CodePipelineReadPolicy (p. 199)</a>	Gives read permission to get details about a CodePipeline pipeline.		
<a href="#">CloudWatchDashboardsPolicy (p. 199)</a>	Gives permissions to put metrics to operate on CloudWatch dashboards.		

Policy Template	Description		
<a href="#">RekognitionFacesModelPolicy</a> (p. 200)	Gives permission to add, delete, and search faces in a collection.		
<a href="#">RekognitionFacesPolicy</a> (p. 200)	Gives permission to compare and detect faces and labels.		
<a href="#">RekognitionLabelsPolicy</a> (p. 200)	Gives permission to detect object and moderation labels.		
<a href="#">DynamoDBBackupPolicy</a> (p. 201)	Gives read and write permission to DynamoDB on-demand backups for a table.		
<a href="#">DynamoDBRestorePolicy</a> (p. 201)	Gives permission to restore a DynamoDB table from backup.		
<a href="#">ComprehendBasicPolicy</a> (p. 200)	Gives permission for detecting entities, key phrases, languages, and sentiments.		
<a href="#">MobileAnalyticsWriteOnlyPolicy</a> (p. 202)	Gives write-only permission to put event data for all application resources.		
<a href="#">PinpointEndpointApplicationPolicy</a> (p. 201)	Gives permission to get and update endpoints for an Amazon Pinpoint application.		
<a href="#">FirehoseWritePolicy</a> (p. 202)	Gives permission to write to a Kinesis Data Firehose delivery stream.		
<a href="#">FirehoseCrudPolicy</a> (p. 202)	Gives permission to create, write, update, and delete a Kinesis Data Firehose delivery stream.		
<a href="#">EKSDescribePolicy</a> (p. 204)	Gives permission to describe or list Amazon EKS clusters.		
<a href="#">CostExplorerReadOnlyPolicy</a> (p. 204)	Gives read-only permission to the read-only Cost Explorer APIs for billing history.		
<a href="#">OrganizationsListAccountsPolicy</a> (p. 205)	Gives read-only permission to list child account names and IDs.		
<a href="#">SESBulkTemplatedEmailPolicy</a> (p. 205)	Gives permission to send email, templated email, templated bulk emails and verify identity.		
<a href="#">SESEmailTemplatePolicy</a> (p. 205)	Gives permission to create, get, list, update and delete Amazon SES email templates.		
<a href="#">FilterLogEventsPolicy</a> (p. 206)	Gives permission to filter log events from a specified log group.		
<a href="#">SSMParameterReadOnlyPolicy</a> (p. 206)	Gives permission to access a parameter to load secrets in this account.		
<a href="#">StepFunctionsExecutionPolicy</a> (p. 207)	Gives permission to start a Step Functions state machine execution.		

Policy Template	Description		
<a href="#">CodeCommitCrudPolicy</a>	Gives permissions to create/read/update/delete objects within a specific CodeCommit repository.		
<a href="#">CodeCommitReadOnlyPolicy</a>	Gives permissions to read objects within a specific CodeCommit repository.		
<a href="#">AthenaQueryPolicy</a>	Gives permissions to execute Athena queries.		

## Policy Template List

The following are the available policy templates, along with the permissions that are applied to each one. AWS SAM automatically populates the placeholder items (such as AWS Region and account ID) with the appropriate information.

### SQSPollerPolicy

Gives permission to poll an Amazon SQS Queue.

```

"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "sqs:ChangeMessageVisibility",
      "sqs:ChangeMessageVisibilityBatch",
      "sqs:DeleteMessage",
      "sqs:DeleteMessageBatch",
      "sqs:GetQueueAttributes",
      "sqs:ReceiveMessage"
    ],
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:sqs:${AWS::Region}:${AWS::AccountId}:${queueName}",
        {
          "queueName": {
            "Ref": "QueueName"
          }
        }
      ]
    }
  }
]

```

### LambdaInvokePolicy

Gives permission to invoke a Lambda function, alias, or version.

```

"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "lambda:InvokeFunction"
    ],
  }
]

```

```
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:lambda:${AWS::Region}:${AWS::AccountId}:function:${functionName}*",
        {
          "functionName": {
            "Ref": "FunctionName"
          }
        }
      ]
    }
  }
}
```

## CloudWatchDescribeAlarmHistoryPolicy

Gives permission to describe CloudWatch alarm history.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "cloudwatch:DescribeAlarmHistory"
    ],
    "Resource": "*"
  }
]
```

## CloudWatchPutMetricPolicy

Gives permission to put metrics to CloudWatch.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "cloudwatch:PutMetricData"
    ],
    "Resource": "*"
  }
]
```

## EC2DescribePolicy

Gives permission to describe Amazon EC2 instances.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "ec2:DescribeRegions",
      "ec2:DescribeInstances"
    ],
    "Resource": "*"
  }
]
```

```
}
]
```

## DynamoDBCrudPolicy

Gives create, read, update, and delete permissions to a DynamoDB table.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "dynamodb:GetItem",
      "dynamodb>DeleteItem",
      "dynamodb:PutItem",
      "dynamodb:Scan",
      "dynamodb:Query",
      "dynamodb:UpdateItem",
      "dynamodb:BatchWriteItem",
      "dynamodb:BatchGetItem",
      "dynamodb:DescribeTable",
      "dynamodb:ConditionCheckItem"
    ],
    "Resource": [
      {
        "Fn::Sub": [
          "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}",
          {
            "tableName": {
              "Ref": "TableName"
            }
          }
        ]
      },
      {
        "Fn::Sub": [
          "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}/index/*",
          {
            "tableName": {
              "Ref": "TableName"
            }
          }
        ]
      }
    ]
  }
]
```

## DynamoDBReadPolicy

Gives read-only permission to a DynamoDB table.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "dynamodb:GetItem",
```

```
        "dynamodb:Scan",
        "dynamodb:Query",
        "dynamodb:BatchGetItem",
        "dynamodb:DescribeTable"
    ],
    "Resource": [
        {
            "Fn::Sub": [
                "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}",
                {
                    "tableName": {
                        "Ref": "TableName"
                    }
                }
            ]
        },
        {
            "Fn::Sub": [
                "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}/index/*",
                {
                    "tableName": {
                        "Ref": "TableName"
                    }
                }
            ]
        }
    ]
}
```

## DynamoDBReconfigurePolicy

Gives permission to reconfigure a DynamoDB table.

```
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "dynamodb:UpdateTable"
            ],
            "Resource": {
                "Fn::Sub": [
                    "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}",
                    {
                        "tableName": {
                            "Ref": "TableName"
                        }
                    }
                ]
            }
        }
    ]
}
```

## SESSendBouncePolicy

Gives SendBounce permission to an Amazon SES identity.

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "ses:SendBounce"
        ],
        "Resource": {
          "Fn::Sub": [
            "arn:${AWS::Partition}:ses:${AWS::Region}:${AWS::AccountId}:identity/${identityName}",
            {
              "identityName": {
                "Ref": "IdentityName"
              }
            }
          ]
        }
      }
    ]
  }
```

## ElasticsearchHttpPostPolicy

Gives POST and PUT permission to Amazon Elasticsearch Service.

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "es:ESHttpPost",
          "es:ESHttpPut"
        ],
        "Resource": {
          "Fn::Sub": [
            "arn:${AWS::Partition}:es:${AWS::Region}:${AWS::AccountId}:domain/${domainName}/*",
            {
              "domainName": {
                "Ref": "DomainName"
              }
            }
          ]
        }
      }
    ]
  }
```

## S3ReadPolicy

Gives read-only permission to objects in an Amazon S3 bucket.

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "s3:GetObject",
          "s3:ListBucket",
          "s3:GetBucketLocation",

```

```

        "s3:GetObjectVersion",
        "s3:GetLifecycleConfiguration"
    ],
    "Resource": [
        {
            "Fn::Sub": [
                "arn:${AWS::Partition}:s3:::${bucketName}",
                {
                    "bucketName": {
                        "Ref": "BucketName"
                    }
                }
            ]
        },
        {
            "Fn::Sub": [
                "arn:${AWS::Partition}:s3:::${bucketName}/*",
                {
                    "bucketName": {
                        "Ref": "BucketName"
                    }
                }
            ]
        }
    ]
}
]

```

## S3CrudPolicy

Gives create, read, update, and delete permission to objects in an Amazon S3 bucket.

```

"Statement": [
    {
        "Effect": "Allow",
        "Action": [
            "s3:GetObject",
            "s3:ListBucket",
            "s3:GetBucketLocation",
            "s3:GetObjectVersion",
            "s3:PutObject",
            "s3:PutObjectAcl",
            "s3:GetLifecycleConfiguration",
            "s3:PutLifecycleConfiguration",
            "s3:DeleteObject"
        ],
        "Resource": [
            {
                "Fn::Sub": [
                    "arn:${AWS::Partition}:s3:::${bucketName}",
                    {
                        "bucketName": {
                            "Ref": "BucketName"
                        }
                    }
                ]
            },
            {
                "Fn::Sub": [
                    "arn:${AWS::Partition}:s3:::${bucketName}/*",
                    {
                        "bucketName": {

```



```
        "Ref": "BucketName"
      }
    ]
  }
}
```

## AMIDescribePolicy

Gives permission to describe Amazon Machine Images (AMIs).

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "ec2:DescribeImages"
        ],
        "Resource": {
          "Fn::Sub": "arn:${AWS::Partition}:ec2:${AWS::Region}:${AWS::AccountId}:image/*"
        }
      }
    ]
```

## CloudFormationDescribeStacksPolicy

Gives permission to describe AWS CloudFormation stacks.

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "cloudformation:DescribeStacks"
        ],
        "Resource": {
          "Fn::Sub": "arn:${AWS::Partition}:cloudformation:${AWS::Region}:${AWS::AccountId}:stack/*"
        }
      }
    ]
```

## RekognitionDetectOnlyPolicy

Gives permission to detect faces, labels, and text.

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "rekognition:DetectFaces",
          "rekognition:DetectLabels",
```

```
        "rekognition:DetectModerationLabels",
        "rekognition:DetectText"
    ],
    "Resource": "*"
}
]
```

## RekognitionNoDataAccessPolicy

Gives permission to compare and detect faces and labels.

```
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "rekognition:CompareFaces",
                "rekognition:DetectFaces",
                "rekognition:DetectLabels",
                "rekognition:DetectModerationLabels"
            ],
            "Resource": {
                "Fn::Sub": [
                    "arn:${AWS::Partition}:rekognition:${AWS::Region}:",
                    "${AWS::AccountId}:collection/${collectionId}",
                    {
                        "collectionId": {
                            "Ref": "CollectionId"
                        }
                    }
                ]
            }
        }
    ]
}
```

## RekognitionReadPolicy

Gives permission to list and search faces.

```
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "rekognition:ListCollections",
                "rekognition:ListFaces",
                "rekognition:SearchFaces",
                "rekognition:SearchFacesByImage"
            ],
            "Resource": {
                "Fn::Sub": [
                    "arn:${AWS::Partition}:rekognition:${AWS::Region}:",
                    "${AWS::AccountId}:collection/${collectionId}",
                    {
                        "collectionId": {
                            "Ref": "CollectionId"
                        }
                    }
                ]
            }
        }
    ]
}
```

```
    }  
  }  
]
```

## RekognitionWriteOnlyAccessPolicy

Gives permission to create collection and index faces.

```
    "Statement": [  
      {  
        "Effect": "Allow",  
        "Action": [  
          "rekognition:CreateCollection",  
          "rekognition:IndexFaces"  
        ],  
        "Resource": {  
          "Fn::Sub": [  
            "arn:${AWS::Partition}:rekognition:${AWS::Region}:  
${AWS::AccountId}:collection/${collectionId}",  
            {  
              "collectionId": {  
                "Ref": "CollectionId"  
              }  
            }  
          ]  
        }  
      }  
    ]  
  }  
]
```

## SQSSendMessagePolicy

Gives permission to send message to an Amazon SQS queue.

```
    "Statement": [  
      {  
        "Effect": "Allow",  
        "Action": [  
          "sqs:SendMessage*"  
        ],  
        "Resource": {  
          "Fn::Sub": [  
            "arn:${AWS::Partition}:sqs:${AWS::Region}:${AWS::AccountId}:${queueName}",  
            {  
              "queueName": {  
                "Ref": "QueueName"  
              }  
            }  
          ]  
        }  
      }  
    ]  
  }  
]
```

## SNSPublishMessagePolicy

Gives permission to publish a message to an Amazon SNS topic.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "sns:Publish"
    ],
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:sns:${AWS::Region}:${AWS::AccountId}:${topicName}",
        {
          "topicName": {
            "Ref": "TopicName"
          }
        }
      ]
    }
  }
]
```

## VPCAccessPolicy

Gives access to create, delete, describe, and detach Elastic Network Interfaces.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "ec2:CreateNetworkInterface",
      "ec2:DeleteNetworkInterface",
      "ec2:DescribeNetworkInterfaces",
      "ec2:DetachNetworkInterface"
    ],
    "Resource": "*"
  }
]
```

## DynamoDBStreamReadPolicy

Gives permission to describe and read DynamoDB streams and records.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "dynamodb:DescribeStream",
      "dynamodb:GetRecords",
      "dynamodb:GetShardIterator"
    ],
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}/stream/${streamName}",
        {
          "tableName": {
            "Ref": "TableName"
          }
        }
      ]
    }
  }
]
```

```

        },
        "streamName": {
            "Ref": "StreamName"
        }
    ]
},
{
    "Effect": "Allow",
    "Action": [
        "dynamodb:ListStreams"
    ],
    "Resource": {
        "Fn::Sub": [
            "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}/stream/*",
            {
                "tableName": {
                    "Ref": "TableName"
                }
            }
        ]
    }
}
]

```

## KinesisStreamReadPolicy

Gives permission to list and read an Amazon Kinesis stream.

```

    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "kinesis:ListStreams",
                "kinesis:DescribeLimits"
            ],
            "Resource": {
                "Fn::Sub": "arn:${AWS::Partition}:kinesis:${AWS::Region}:
${AWS::AccountId}:stream/*"
            }
        },
        {
            "Effect": "Allow",
            "Action": [
                "kinesis:DescribeStream",
                "kinesis:DescribeStreamSummary",
                "kinesis:GetRecords",
                "kinesis:GetShardIterator"
            ],
            "Resource": {
                "Fn::Sub": [
                    "arn:${AWS::Partition}:kinesis:${AWS::Region}:${AWS::AccountId}:stream/
${streamName}",
                    {
                        "streamName": {
                            "Ref": "StreamName"
                        }
                    }
                ]
            }
        }
    ]
}

```

```
    }  
  ]  
}
```

## SESCrudPolicy

Gives permission to send email and verify identity.

```
    "Statement": [  
      {  
        "Effect": "Allow",  
        "Action": [  
          "ses:GetIdentityVerificationAttributes",  
          "ses:SendEmail",  
          "ses:VerifyEmailIdentity"  
        ],  
        "Resource": {  
          "Fn::Sub": [  
            "arn:${AWS::Partition}:ses:${AWS::Region}:${AWS::AccountId}:identity/  
${identityName}",  
            {  
              "identityName": {  
                "Ref": "IdentityName"  
              }  
            }  
          ]  
        }  
      }  
    ]  
  }  
}
```

## SNSCrudPolicy

Gives permission to create, publish, and subscribe to Amazon SNS topics.

```
    "Statement": [  
      {  
        "Effect": "Allow",  
        "Action": [  
          "sns:ListSubscriptionsByTopic",  
          "sns:CreateTopic",  
          "sns:SetTopicAttributes",  
          "sns:Subscribe",  
          "sns:Publish"  
        ],  
        "Resource": {  
          "Fn::Sub": [  
            "arn:${AWS::Partition}:sns:${AWS::Region}:${AWS::AccountId}:${topicName}*",  
            {  
              "topicName": {  
                "Ref": "TopicName"  
              }  
            }  
          ]  
        }  
      }  
    ]  
  }  
}
```

## KinesisCrudPolicy

Gives permission to create, publish, and delete an Amazon Kinesis stream.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "kinesis:AddTagsToStream",
      "kinesis:CreateStream",
      "kinesis:DecreaseStreamRetentionPeriod",
      "kinesis>DeleteStream",
      "kinesis:DescribeStream",
      "kinesis:DescribeStreamSummary",
      "kinesis:GetShardIterator",
      "kinesis:IncreaseStreamRetentionPeriod",
      "kinesis:ListTagsForStream",
      "kinesis:MergeShards",
      "kinesis:PutRecord",
      "kinesis:PutRecords",
      "kinesis:SplitShard",
      "kinesis:RemoveTagsFromStream"
    ],
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:kinesis:${AWS::Region}:${AWS::AccountId}:stream/${streamName}",
        {
          "streamName": {
            "Ref": "StreamName"
          }
        }
      ]
    }
  }
]
```

## KMSDecryptPolicy

Gives permission to decrypt with an AWS KMS key.

```
"Statement": [
  {
    "Action": "kms:Decrypt",
    "Effect": "Allow",
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:kms:${AWS::Region}:${AWS::AccountId}:key/${keyId}",
        {
          "keyId": {
            "Ref": "KeyId"
          }
        }
      ]
    }
  }
]
```

## KMSEncryptPolicy

Gives permission to encrypt with an AWS KMS key.

```
"Statement": [
  {
    "Action": "kms:Encrypt",
    "Effect": "Allow",
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:kms:${AWS::Region}:${AWS::AccountId}:key/${keyId}",
        {
          "keyId": {
            "Ref": "KeyId"
          }
        }
      ]
    }
  }
]
```

## PollyFullAccessPolicy

Gives full access permission to Amazon Polly lexicon resources.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "polly:GetLexicon",
      "polly:DeleteLexicon"
    ],
    "Resource": [
      {
        "Fn::Sub": [
          "arn:${AWS::Partition}:polly:${AWS::Region}:${AWS::AccountId}:lexicon/
${lexiconName}",
          {
            "lexiconName": {
              "Ref": "LexiconName"
            }
          }
        ]
      }
    ]
  },
  {
    "Effect": "Allow",
    "Action": [
      "polly:DescribeVoices",
      "polly:ListLexicons",
      "polly:PutLexicon",
      "polly:SynthesizeSpeech"
    ],
    "Resource": [
      {
        "Fn::Sub": "arn:${AWS::Partition}:polly:${AWS::Region}:
${AWS::AccountId}:lexicon/*"
      }
    ]
  }
]
```



```
    ]
  }
]
```

## S3FullAccessPolicy

Gives full access permission to objects in an Amazon S3 bucket.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "s3:GetObject",
      "s3:GetObjectAcl",
      "s3:GetObjectVersion",
      "s3:PutObject",
      "s3:PutObjectAcl",
      "s3:DeleteObject",
      "s3:DeleteObjectTagging",
      "s3:DeleteObjectVersionTagging",
      "s3:GetObjectTagging",
      "s3:GetObjectVersionTagging",
      "s3:PutObjectTagging",
      "s3:PutObjectVersionTagging"
    ],
    "Resource": [
      {
        "Fn::Sub": [
          "arn:${AWS::Partition}:s3:::${bucketName}/*",
          {
            "bucketName": {
              "Ref": "BucketName"
            }
          }
        ]
      }
    ]
  },
  {
    "Effect": "Allow",
    "Action": [
      "s3:ListBucket",
      "s3:GetBucketLocation",
      "s3:GetLifecycleConfiguration",
      "s3:PutLifecycleConfiguration"
    ],
    "Resource": [
      {
        "Fn::Sub": [
          "arn:${AWS::Partition}:s3:::${bucketName}",
          {
            "bucketName": {
              "Ref": "BucketName"
            }
          }
        ]
      }
    ]
  }
]
```

## CodePipelineLambdaExecutionPolicy

Gives permission for a Lambda function invoked by CodePipeline to report the status of the job.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "codepipeline:PutJobSuccessResult",
      "codepipeline:PutJobFailureResult"
    ],
    "Resource": "*"
  }
]
```

## ServerlessRepoReadWriteAccessPolicy

Gives permission to create and list applications in the AWS Serverless Application Repository service.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "serverlessrepo:CreateApplication",
      "serverlessrepo:CreateApplicationVersion",
      "serverlessrepo:GetApplication",
      "serverlessrepo:ListApplications",
      "serverlessrepo:ListApplicationVersions"
    ],
    "Resource": [
      {
        "Fn::Sub": "arn:${AWS::Partition}:serverlessrepo:${AWS::Region}:
${AWS::AccountId}:applications/*"
      }
    ]
  }
]
```

## EC2CopyImagePolicy

Gives permission to copy Amazon EC2 images.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "ec2:CopyImage"
    ],
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:ec2:${AWS::Region}:${AWS::AccountId}:image/
${imageId}",
        {
          "imageId": {
```

```
        "Ref": "ImageId"
      }
    }
  ]
}
```

## AWSecretsManagerRotationPolicy

Gives permission to rotate a secret in AWS Secrets Manager.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "secretsmanager:DescribeSecret",
      "secretsmanager:GetSecretValue",
      "secretsmanager:PutSecretValue",
      "secretsmanager:UpdateSecretVersionStage"
    ],
    "Resource": {
      "Fn::Sub": "arn:${AWS::Partition}:secretsmanager:${AWS::Region}:
${AWS::AccountId}:secret:*"
    },
    "Condition": {
      "StringEquals": {
        "secretsmanager:resource/AllowRotationLambdaArn": {
          "Fn::Sub": [
            "arn:${AWS::Partition}:lambda:${AWS::Region}:
${AWS::AccountId}:function:${functionName}",
            {
              "functionName": {
                "Ref": "FunctionName"
              }
            }
          ]
        }
      }
    }
  },
  {
    "Effect": "Allow",
    "Action": [
      "secretsmanager:GetRandomPassword"
    ],
    "Resource": "*"
  }
]
```

## AWSecretsManagerGetSecretValuePolicy

Gives permission to GetSecretValue for the specified AWS Secrets Manager secret.

```
"Statement": [
  {
    "Effect": "Allow",
```

```
    "Action": [
      "secretsmanager:GetSecretValue"
    ],
    "Resource": {
      "Fn::Sub": [
        "${secretArn}",
        {
          "secretArn": {
            "Ref": "SecretArn"
          }
        }
      ]
    }
  }
}
```

## CodePipelineReadOnlyPolicy

Gives read permission to get details about a CodePipeline pipeline.

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "codepipeline:ListPipelineExecutions"
        ],
        "Resource": {
          "Fn::Sub": [
            "arn:${AWS::Partition}:codepipeline:${AWS::Region}:${AWS::AccountId}:",
            "${pipelineName}",
            {
              "pipelineName": {
                "Ref": "PipelineName"
              }
            }
          ]
        }
      }
    ]
}
```

## CloudWatchDashboardPolicy

Gives permissions to put metrics to operate on CloudWatch dashboards.

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "cloudwatch:GetDashboard",
          "cloudwatch:ListDashboards",
          "cloudwatch:PutDashboard",
          "cloudwatch:ListMetrics"
        ],
        "Resource": "*"
      }
    ]
}
```

## RekognitionFacesManagementPolicy

Gives permission to add, delete, and search faces in a collection.

```
"Statement": [{
  "Effect": "Allow",
  "Action": [
    "rekognition:IndexFaces",
    "rekognition:DeleteFaces",
    "rekognition:SearchFaces",
    "rekognition:SearchFacesByImage",
    "rekognition:ListFaces"
  ],
  "Resource": {
    "Fn::Sub": [
      "arn:${AWS::Partition}:rekognition:${AWS::Region}:
${AWS::AccountId}:collection/${collectionId}",
      {
        "collectionId": {
          "Ref": "CollectionId"
        }
      }
    ]
  }
}]
```

## RekognitionFacesPolicy

Gives permission to compare and detect faces and labels.

```
"Statement": [{
  "Effect": "Allow",
  "Action": [
    "rekognition:CompareFaces",
    "rekognition:DetectFaces"
  ],
  "Resource": "*"
}]
```

## RekognitionLabelsPolicy

Gives permission to detect object and moderation labels.

```
"Statement": [{
  "Effect": "Allow",
  "Action": [
    "rekognition:DetectLabels",
    "rekognition:DetectModerationLabels"
  ],
  "Resource": "*"
}]
```

## DynamoDBBackupFullAccessPolicy

Gives read and write permission to DynamoDB on-demand backups for a table.

```
"Statement": [{
  "Effect": "Allow",
  "Action": [
    "dynamodb:CreateBackup",
    "dynamodb:DescribeContinuousBackups"
  ],
  "Resource": {
    "Fn::Sub": [
      "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}",
      {
        "tableName": {
          "Ref": "TableName"
        }
      }
    ]
  }
}, {
  "Effect": "Allow",
  "Action": [
    "dynamodb>DeleteBackup",
    "dynamodb:DescribeBackup",
    "dynamodb:ListBackups"
  ],
  "Resource": {
    "Fn::Sub": [
      "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}/backup/*",
      {
        "tableName": {
          "Ref": "TableName"
        }
      }
    ]
  }
}]
```

## DynamoDBRestoreFromBackupPolicy

Gives permission to restore a DynamoDB table from backup.

```
"Statement": [{
  "Effect": "Allow",
  "Action": [
    "dynamodb:RestoreTableFromBackup"
  ],
  "Resource": {
    "Fn::Sub": [
      "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}/backup/*",
      {
        "tableName": {
          "Ref": "TableName"
        }
      }
    ]
  }
}]
```

```
    }
  }
]
},
{
  "Effect": "Allow",
  "Action": [
    "dynamodb:PutItem",
    "dynamodb:UpdateItem",
    "dynamodb>DeleteItem",
    "dynamodb:GetItem",
    "dynamodb:Query",
    "dynamodb:Scan",
    "dynamodb:BatchWriteItem"
  ],
  "Resource": {
    "Fn::Sub": [
      "arn:${AWS::Partition}:dynamodb:${AWS::Region}:${AWS::AccountId}:table/
${tableName}",
      {
        "tableName": {
          "Ref": "TableName"
        }
      }
    ]
  }
}
]
```

## ComprehendBasicAccessPolicy

Gives permission for detecting entities, key phrases, languages, and sentiments.

```
"Statement": [{
  "Effect": "Allow",
  "Action": [
    "comprehend:BatchDetectKeyPhrases",
    "comprehend:DetectDominantLanguage",
    "comprehend:DetectEntities",
    "comprehend:BatchDetectEntities",
    "comprehend:DetectKeyPhrases",
    "comprehend:DetectSentiment",
    "comprehend:BatchDetectDominantLanguage",
    "comprehend:BatchDetectSentiment"
  ],
  "Resource": "*"
}]
```

## MobileAnalyticsWriteOnlyAccessPolicy

Gives write-only permission to put event data for all application resources.

```
"Statement": [
  {
    "Effect": "Allow",
```

```
    "Action": [
      "mobileanalytics:PutEvents"
    ],
    "Resource": "*"
  }
]
```

## PinpointEndpointAccessPolicy

Gives permission to get and update endpoints for an Amazon Pinpoint application.

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "mobiletargeting:GetEndpoint",
          "mobiletargeting:UpdateEndpoint",
          "mobiletargeting:UpdateEndpointsBatch"
        ],
        "Resource": {
          "Fn::Sub": [
            "arn:${AWS::Partition}:mobiletargeting:${AWS::Region}:${AWS::AccountId}:apps/${pinpointApplicationId}/endpoints/*",
            {
              "pinpointApplicationId": {
                "Ref": "PinpointApplicationId"
              }
            }
          ]
        }
      }
    ]
  }
]
```

## FirehoseWritePolicy

Gives permission to write to a Kinesis Data Firehose delivery stream.

```
    "Statement": [
      {
        "Effect": "Allow",
        "Action": [
          "firehose:PutRecord",
          "firehose:PutRecordBatch"
        ],
        "Resource": {
          "Fn::Sub": [
            "arn:${AWS::Partition}:firehose:${AWS::Region}:${AWS::AccountId}:deliverystream/${deliveryStreamName}",
            {
              "deliveryStreamName": {
                "Ref": "DeliveryStreamName"
              }
            }
          ]
        }
      }
    ]
  }
]
```



## FirehoseCrudPolicy

Gives permission to create, write, update, and delete a Kinesis Data Firehose delivery stream.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "firehose:CreateDeliveryStream",
      "firehose>DeleteDeliveryStream",
      "firehose:DescribeDeliveryStream",
      "firehose:PutRecord",
      "firehose:PutRecordBatch",
      "firehose:UpdateDestination"
    ],
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:firehose:${AWS::Region}:",
        "${AWS::AccountId}:deliverystream/${deliveryStreamName}",
        {
          "deliveryStreamName": {
            "Ref": "DeliveryStreamName"
          }
        }
      ]
    }
  }
]
```

## EKSDescribePolicy

Gives permission to describe or list Amazon EKS clusters.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "eks:DescribeCluster",
      "eks:ListClusters"
    ],
    "Resource": "*"
  }
]
```

## CostExplorerReadOnlyPolicy

Gives read-only permission to the read-only Cost Explorer APIs for billing history.

```
"Statement": [{
  "Effect": "Allow",
  "Action": [
    "ce:GetCostAndUsage",
```

```
        "ce:GetDimensionValues",
        "ce:GetReservationCoverage",
        "ce:GetReservationPurchaseRecommendation",
        "ce:GetReservationUtilization",
        "ce:GetTags"
    ],
    "Resource": "*"
  }]
}
```

## OrganizationsListAccountsPolicy

Gives read-only permission to list child account names and IDs.

```
  "Statement": [{
    "Effect": "Allow",
    "Action": [
      "organizations:ListAccounts"
    ],
    "Resource": "*"
  }]
}
```

## SESBulkTemplatedCrudPolicy

Gives permission to send email, templated email, templated bulk emails and verify identity.

```
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "ses:GetIdentityVerificationAttributes",
        "ses:SendEmail",
        "ses:SendRawEmail",
        "ses:SendTemplatedEmail",
        "ses:SendBulkTemplatedEmail",
        "ses:VerifyEmailIdentity"
      ],
      "Resource": {
        "Fn::Sub": [
          "arn:${AWS::Partition}:ses:${AWS::Region}:${AWS::AccountId}:identity/${identityName}",
          {
            "identityName": {
              "Ref": "IdentityName"
            }
          }
        ]
      }
    }
  ]
}
```

## SESEmailTemplateCrudPolicy

Gives permission to create, get, list, update and delete Amazon SES email templates.

```
"Statement": [{
  "Effect": "Allow",
  "Action": [
    "ses:CreateTemplate",
    "ses:GetTemplate",
    "ses:ListTemplates",
    "ses:UpdateTemplate",
    "ses>DeleteTemplate",
    "ses:TestRenderTemplate"
  ],
  "Resource": "*"
}]
```

## FilterLogEventsPolicy

Gives permission to filter log events from a specified log group.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "logs:FilterLogEvents"
    ],
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:logs:${AWS::Region}:${AWS::AccountId}:log-group:
${logGroupName}:log-stream:*",
        {
          "logGroupName": {
            "Ref": "LogGroupName"
          }
        }
      ]
    }
  }
]
```

## SSMParameterReadPolicy

Gives permission to access a parameter to load secrets in this account.

**Note**

If you are not using default key, you will also need `KMSDecryptPolicy`.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "ssm:DescribeParameters"
    ],
    "Resource": "*"
  },
  {
    "Effect": "Allow",
    "Action": [
      "ssm:GetParameters",
      "ssm:GetParameter",

```

```
        "ssm:GetParametersByPath"
      ],
      "Resource": {
        "Fn::Sub": [
          "arn:${AWS::Partition}:ssm:${AWS::Region}:${AWS::AccountId}:parameter/",
          "${parameterName}",
          {
            "parameterName": {
              "Ref": "ParameterName"
            }
          }
        ]
      }
    }
  ]
}
```

## StepFunctionsExecutionPolicy

Gives permission to start a Step Functions state machine execution.

```
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "states:StartExecution"
      ],
      "Resource": {
        "Fn::Sub": [
          "arn:${AWS::Partition}:states:${AWS::Region}:${AWS::AccountId}:stateMachine:${stateMachineName}",
          {
            "stateMachineName": {
              "Ref": "StateMachineName"
            }
          }
        ]
      }
    }
  ]
}
```

## CodeCommitCrudPolicy

Gives permissions to create/read/update/delete objects within a specific CodeCommit repository.

```
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "codecommit:GitPull",
        "codecommit:GitPush",
        "codecommit:CreateBranch",
        "codecommit>DeleteBranch",
        "codecommit:GetBranch",
        "codecommit:ListBranches",
        "codecommit:MergeBranchesByFastForward",
        "codecommit:MergeBranchesBySquash",
        "codecommit:MergeBranchesByThreeWay",
        "codecommit:UpdateDefaultBranch",

```

```

        "codecommit:BatchDescribeMergeConflicts",
        "codecommit:CreateUnreferencedMergeCommit",
        "codecommit:DescribeMergeConflicts",
        "codecommit:GetMergeCommit",
        "codecommit:GetMergeOptions",
        "codecommit:BatchGetPullRequests",
        "codecommit:CreatePullRequest",
        "codecommit:DescribePullRequestEvents",
        "codecommit:GetCommentsForPullRequest",
        "codecommit:GetCommitsFromMergeBase",
        "codecommit:GetMergeConflicts",
        "codecommit:GetPullRequest",
        "codecommit:ListPullRequests",
        "codecommit:MergePullRequestByFastForward",
        "codecommit:MergePullRequestBySquash",
        "codecommit:MergePullRequestByThreeWay",
        "codecommit:PostCommentForPullRequest",
        "codecommit:UpdatePullRequestDescription",
        "codecommit:UpdatePullRequestStatus",
        "codecommit:UpdatePullRequestTitle",
        "codecommit>DeleteFile",
        "codecommit:GetBlob",
        "codecommit:GetFile",
        "codecommit:GetFolder",
        "codecommit:PutFile",
        "codecommit>DeleteCommentContent",
        "codecommit:GetComment",
        "codecommit:GetCommentsForComparedCommit",
        "codecommit:PostCommentForComparedCommit",
        "codecommit:PostCommentReply",
        "codecommit:UpdateComment",
        "codecommit:BatchGetCommits",
        "codecommit>CreateCommit",
        "codecommit:GetCommit",
        "codecommit:GetCommitHistory",
        "codecommit:GetDifferences",
        "codecommit:GetObjectIdentifier",
        "codecommit:GetReferences",
        "codecommit:GetTree",
        "codecommit:GetRepository",
        "codecommit:UpdateRepositoryDescription",
        "codecommit:ListTagsForResource",
        "codecommit:TagResource",
        "codecommit:UntagResource",
        "codecommit:GetRepositoryTriggers",
        "codecommit:PutRepositoryTriggers",
        "codecommit:TestRepositoryTriggers",
        "codecommit:GetBranch",
        "codecommit:GetCommit",
        "codecommit:UploadArchive",
        "codecommit:GetUploadArchiveStatus",
        "codecommit:CancelUploadArchive"
    ],
    "Resource": {
        "Fn::Sub": [
            "arn:${AWS::Partition}:codecommit:${AWS::Region}:${AWS::AccountId}:",
            "${repositoryName}",
            {
                "repositoryName": {
                    "Ref": "RepositoryName"
                }
            }
        ]
    }
}
]

```

## CodeCommitReadPolicy

Gives permissions to read objects within a specific CodeCommit repository.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "codecommit:GitPull",
      "codecommit:GetBranch",
      "codecommit:ListBranches",
      "codecommit:BatchDescribeMergeConflicts",
      "codecommit:DescribeMergeConflicts",
      "codecommit:GetMergeCommit",
      "codecommit:GetMergeOptions",
      "codecommit:BatchGetPullRequests",
      "codecommit:DescribePullRequestEvents",
      "codecommit:GetCommentsForPullRequest",
      "codecommit:GetCommitsFromMergeBase",
      "codecommit:GetMergeConflicts",
      "codecommit:GetPullRequest",
      "codecommit:ListPullRequests",
      "codecommit:GetBlob",
      "codecommit:GetFile",
      "codecommit:GetFolder",
      "codecommit:GetComment",
      "codecommit:GetCommentsForComparedCommit",
      "codecommit:BatchGetCommits",
      "codecommit:GetCommit",
      "codecommit:GetCommitHistory",
      "codecommit:GetDifferences",
      "codecommit:GetObjectIdentifier",
      "codecommit:GetReferences",
      "codecommit:GetTree",
      "codecommit:GetRepository",
      "codecommit:ListTagsForResource",
      "codecommit:GetRepositoryTriggers",
      "codecommit:TestRepositoryTriggers",
      "codecommit:GetBranch",
      "codecommit:GetCommit",
      "codecommit:GetUploadArchiveStatus"
    ],
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:codecommit:${AWS::Region}:${AWS::AccountId}:",
        "${repositoryName}",
        {
          "repositoryName": {
            "Ref": "RepositoryName"
          }
        }
      ]
    }
  }
]
```

## AthenaQueryPolicy

Gives permissions to execute Athena queries.

```
"Statement": [
  {
    "Effect": "Allow",
    "Action": [
      "athena:ListWorkGroups",
      "athena:GetExecutionEngine",
      "athena:GetExecutionEngines",
      "athena:GetNamespace",
      "athena:GetCatalogs",
      "athena:GetNamespaces",
      "athena:GetTables",
      "athena:GetTable"
    ],
    "Resource": "*"
  },
  {
    "Effect": "Allow",
    "Action": [
      "athena:StartQueryExecution",
      "athena:GetQueryResults",
      "athena>DeleteNamedQuery",
      "athena:GetNamedQuery",
      "athena:ListQueryExecutions",
      "athena:StopQueryExecution",
      "athena:GetQueryResultsStream",
      "athena:ListNamedQueries",
      "athena:CreateNamedQuery",
      "athena:GetQueryExecution",
      "athena:BatchGetNamedQuery",
      "athena:BatchGetQueryExecution",
      "athena:GetWorkGroup"
    ],
    "Resource": {
      "Fn::Sub": [
        "arn:${AWS::Partition}:athena:${AWS::Region}:${AWS::AccountId}:workgroup/${workgroupName}",
        {
          "workgroupName": {
            "Ref": "WorkGroupName"
          }
        }
      ]
    }
  }
]
```

## Telemetry in the AWS SAM CLI

At AWS, we develop and launch services based on what we learn from interactions with customers. We use customer feedback to iterate on our product. Telemetry is additional information that helps us to better understand our customers' needs, diagnose issues, and deliver features that improve the customer experience.

The AWS SAM CLI collects telemetry, such as generic usage metrics, system and environment information, and errors. For details of the types of telemetry collected, see [Types of Information Collected \(p. 211\)](#).

The AWS SAM CLI does **not** collect personal information, such as usernames or email addresses. It also does not extract sensitive project-level information.

Customers control whether telemetry is enabled, and can change their settings at any point of time. If telemetry remains enabled, the AWS SAM CLI sends telemetry data in the background without requiring any additional customer interaction.

## Disabling Telemetry for a Session

In macOS and Linux operating systems, you can disable telemetry for a single session. To disable telemetry for your current session, run the following command to set the environment variable `SAM_CLI_TELEMETRY` to `false`. You must repeat the command for each new terminal or session.

```
export SAM_CLI_TELEMETRY=0
```

## Disabling Telemetry for Your Profile in All Sessions

Run the following commands to disable telemetry for all sessions when you're running the AWS SAM CLI on your operating system.

### To disable telemetry in Linux

1. Run:

```
echo "export SAM_CLI_TELEMETRY=0" >> ~/.profile
```

2. Run:

```
source ~/.profile
```

### To disable telemetry in macOS

1. Run:

```
echo "export SAM_CLI_TELEMETRY=0" >> ~/.profile
```

2. Run:

```
source ~/.profile
```

### To disable telemetry in Windows

1. Run:

```
setx SAM_CLI_TELEMETRY 0
```

2. Run:

```
refreshenv
```

## Types of Information Collected

- **Usage information** – The generic commands and subcommands that are run.



- **Errors and diagnostic information** – The status and duration of commands that are run, including exit codes, internal exception names, and failures when connecting to Docker.
- **System and environment information** – The Python version, operating system (Windows, Linux, or macOS), and environment in which the AWS SAM CLI is executed (for example, AWS CodeBuild, an AWS IDE toolkit, or a terminal).

## Learn More

The telemetry data that's collected adheres to the AWS data privacy policies. For more information, see the following:

- [AWS Service Terms](#)
- [Data Privacy](#)

# Document History for AWS Serverless Application Model

The following table describes the important changes in each release of the *AWS Serverless Application Model Developer Guide*. For notification about updates to this documentation, you can subscribe to an RSS feed by choosing the RSS button in the top menu panel.

- **Latest documentation update:** January 17, 2020

update-history-change	update-history-description	update-history-date
<a href="#">Setting up AWS credentials (p. 213)</a>	Added instructions for setting up AWS credentials, for readers who have not already set them to use with other tools like one of the AWS SDKs or the AWS CLI. For more information, see <a href="#">Setting Up AWS Credentials</a> .	January 17, 2020
<a href="#">AWS SAM Specification and AWS SAM CLI updates (p. 213)</a>	Migrated the AWS SAM Specification from GitHub. For more information see <a href="#">AWS SAM Specification</a> . Also updated the deployment workflow with changes to the <code>sam deploy</code> command.	November 25, 2019
<a href="#">New options for controlling access to API Gateway APIs and policy template updates (p. 213)</a>	Added new options for controlling access to API Gateway APIs: IAM permissions, API keys, and resource policies. For more information, see <a href="#">Controlling Access to API Gateway APIs</a> . Also updated two policy templates: <code>RekognitionFacesPolicy</code> and <code>ElasticsearchHttpPostPolicy</code> . For more information, see <a href="#">AWS SAM Policy Templates</a> .	August 29, 2019
<a href="#">Getting Started updates (p. 213)</a>	Updated Getting Started chapter with improved installation instructions for the AWS SAM CLI and the Hello World tutorial. For more information, see <a href="#">Getting Started</a> .	July 25, 2019
<a href="#">Controlling access to API Gateway APIs (p. 213)</a>	Added support for controlling access to API Gateway APIs. For more information, see <a href="#">Controlling Access to API Gateway APIs</a> .	March 21, 2019

<a href="#">Added sam publish to the AWS SAM CLI (p. 213)</a>	The new <code>sam publish</code> command in the AWS SAM CLI simplifies the process for publishing serverless applications in the AWS Serverless Application Repository. For more information, see <a href="#">Publishing Serverless Applications Using the AWS SAM CLI</a> .	December 21, 2018
<a href="#">Nested applications and layers support (p. 213)</a>	Added support for nested applications and layers. For more information, see <a href="#">Nested Applications</a> and <a href="#">Working with Layers</a> .	November 29, 2018
<a href="#">Added sam build to the AWS SAM CLI (p. 213)</a>	The new <code>sam build</code> command in the AWS SAM CLI simplifies the process for compiling serverless applications with dependencies so you can locally test and deploy these applications. For more information, see <a href="#">Building Applications with Dependencies</a> .	November 19, 2018
<a href="#">Added new installation options for the AWS SAM CLI (p. 213)</a>	Added Linuxbrew (Linux), MSI (Windows), and Homebrew (macOS) installation options for the AWS SAM CLI. For more information, see <a href="#">Installing the AWS SAM CLI</a> .	November 7, 2018
<a href="#">New guide (p. 213)</a>	This is the first release of the <i>AWS Serverless Application Model Developer Guide</i> .	October 17, 2018